



D 2026

RICOSLAM: LOCALIZATION AND MAPPING IN DYNAMIC ENVIRONMENTS

RICARDO BARBOSA SOUSA

TESE DE DOUTORAMENTO APRESENTADA
À FACULDADE DE ENGENHARIA DA UNIVERSIDADE DO PORTO EM
ENGENHARIA ELETROTÉCNICA E DE COMPUTADORES

RicoSLAM: Localization and Mapping in Dynamic Environments

Ricardo B. Sousa 



Doctoral Program in Electrical and Computer Engineering

Supervisor: Prof. António Paulo Moreira , PhD

Second Supervisor: Héber Miguel Sobreira , PhD

April 14, 2026

RicoSLAM: Localization and Mapping in Dynamic Environments

Ricardo B. Sousa 

Doctoral Program in Electrical and Computer Engineering

PhD Thesis submitted for the doctoral final oral examination . . . :

President: Doctor Luís Miguel Almeida (TBC), Full Professor,
Faculdade de Engenharia da Universidade do Porto, Portugal

Referee: Doctor Robert Deaves (TBC), Honorary Professor,
School of Engineering, University of Birmingham, United Kingdom

Referee: Doctor Carlos Cardeira (TBC), Associate Professor,
Instituto Superior Técnico, Universidade de Lisboa, Portugal

Referee: Doctor Miguel Riem de Oliveira (TBC), Assistant Professor,
Department of Mechanical Engineering, Universidade de Aveiro, Portugal

Referee: Doctor Paulo G. Costa, Associate Professor,
Faculdade de Engenharia da Universidade do Porto, Portugal

Advisor: Doctor António Paulo Moreira, Full Professor,
Faculdade de Engenharia da Universidade do Porto, Portugal

April 14, 2026

Resumo

TBC

Palavras-chave: TBC.

Abstract

TBC

Keywords: TBC.

Acknowledgments

First and foremost, I would like to thank my supervisors, Prof. António Paulo Moreira and Héber Miguel Sobreira, for their unwavering support throughout my PhD journey and for the opportunity to work with them. This journey would not have been possible without them, including the invitation by Prof. António Paulo upon completing my M.Sc. dissertation, starting to work with Héber at CRIIS – Centre for Robotics in Industry and Intelligent Systems from INESC TEC – Institute for Systems and Computer Engineering, Technology and Science, and then exploring and researching perception, Simultaneous Localization and Mapping (SLAM), calibration, control systems, and mobile robotics in general.

Next, a special thanks to my colleagues from CRIIS, including from the laboratories iiLab – Industry and Innovation Laboratory and TRIBElab – Laboratory of Robotics and IoT for Smart Precision Agriculture and Forestry. Although my research is related to the intralogistics industry and more related to iiLab’s scope and mission, I began my PhD journey working side by side with TRIBElab, as the lab is located on the campus of the Faculty of Engineering, University of Porto (FEUP). This experience was excellent in terms of the team environment and spirit, while also providing the opportunity to interact daily with my mates from the master’s degree, namely André and Daniel. Additionally, the TRIBElab leader, Filipe Santos, was one of the most important people who helped me decide that staying at INESC TEC would be the best choice to continue my journey and pursue a PhD degree, given the interaction with real-world applications and the opportunity to move beyond simulators and the theoretical world. So, a special thanks to him for that support, and to André and Daniel, two close companions and friends who supported me greatly throughout my PhD. After I moved to the iiLab facilities, it has been a fantastic journey, marked by the creation of new team partnerships and demonstrating that working side by side with colleagues from the same application field and different research areas can result in fascinating research experiences. A special thanks to my team leader Héber, also the co-supervisor of my PhD Thesis, Prof. Manuel Silva, for his guidance on the GreenAuto project directly linked with my PhD research and all the support given by him and the current and past CRIIS coordination team, to all my teammates of the Mobile Robotics Development Team (MRDT), and all my colleagues from iiLab that I interact daily with and has been a pleasure to work with them.

Another very special person in the PhD program was Giorgio Grisetti from Sapienza University of Rome. Upon gathering persons of interest for a PhD internship, Giorgio clearly distinguished himself by having everything I wanted to supervise me in that experience: theoretical background, both in filter and graph-based formulations for the SLAM problem, and joining the theory with a tremendous experience in implementing SLAM algorithms with real-world applications, including having developed countless open-source applications used by everyone in the world (GMapping, collaboration in the g2o framework, among others). I contacted him by email on July 18th, 2023, a Tuesday, after dinner, thinking that I would only receive an answer after the summer vacations or not at all. To my complete surprise, the very next day, at 8 AM (Rome local time), I received a meeting invitation from Giorgio scheduled for Friday of the same week. At that meeting, the internship was agreed upon, including the timeframe (common availability between January and June 2024 and a time period of three months), the scope, and, especially, how to prepare for the internship, by doing the entire Probabilistic Robotics course from Giorgio that had online video recordings from the COVID time (2020/21 and 2021/22 academic years). Indeed, the experience in Rome was life-changing in terms of my research, having a tremendous impact on my PhD Thesis, as evidenced by the extension of my stay from three to six months (from February until July 2023). Of course, a lot of that experience was thanks to Giorgio, for being the great person he is, not caring about societal status (in terms of academic background and professional

success) and being hands-on, interacting daily with his PhD students and with external ones like me. There are no words enough to thank him for that, a massive **Thank You** for the experience! Additionally, I would like to thank my colleagues from the Robotics Vision and Perception Group (RVP-Group) and the LabRoCoCo for their hospitality and engaging research discussions. Last but not least, a big thank you to INESC TEC and, especially, to the past CRIIS coordination team (Prof. António Paulo and Germano Veiga) for giving me the opportunity to undertake the PhD internship. Although I transitioned from the national Fundação para a Ciência e a Tecnologia (FCT) PhD scholarship to a researcher contract at INESC TEC in April 2023, the CRIIS coordination not only supported but also encouraged me to go to the internship. So, a big thank you to them and the teammates at CRIIS, who, of course, supported me throughout the journey.

In parallel with my PhD, I had the opportunity to be part of the 5dpo Robotics Team at FEUP. I want to thank all past and present members of the team, with special thanks to Prof. António Paulo, Cláudia Rocha, João Martins, José Pedro, José Sarmiento, and Maria Lopes. Thank you for all the time we spent together, for our participation in the Portuguese Robotics Open (FNR), for the frustration of the defeats that helped us grow as a team, and for the happiness, experience, knowledge, and pride that came with our victories. Thank you also for all the hard work, partnership, friendship, and the many engaging discussions, forming the backbone of an incredible team and experience, one in which I truly felt part of a team fighting toward common goals, working together without hidden interests, supporting one another to achieve our objectives while enjoying our time together and learning from each other.

Fortunately, I have always had a strong set of friendships behind me, supporting me since high school, my master's degree, and now throughout my PhD. A big thank you to my friends from high school and further: Bárbara and Ruben, Bruno and Juliana, Beatriz and Rocha, Hugo and Márcia, Mariana and Tommy, Pimentas, and Sâncio. They are to me like a family, always having the weekly coffee in Gaia to discuss economics, politics, a little bit of football (unfortunately, most of them support Porto, not my Benfica!) and gossip. Furthermore, another big thanks to all my friends from the masters at FEUP who are still present for the monthly craft beer at Cervejaria do Carmo or Cervejaria Gulden Draak Bierhuis in Porto and possibly dinner afterwards, depending on the mood, and the yearly Christmas dinner, including Bruno, Daniel and Clara, Hélder, Mariano, Mariana and Granja, Pedro, Ricardo and Isabel, Tiago, Tovar, and Xavi. It's quite fascinating that, even after completing our master's degrees at FEUP, we have still been meeting regularly since 2020. I will never forget the special moment when Bruno, Hélder, Mariano, and Tovar took a road trip from Germany to Rome, Italy, and back, during May 2024, while I was still on my PhD internship. Indeed, that weekend we spent together exploring Rome and enjoying good food and drink was the best of my whole stay in Italy, forever etched in my memory. Also, I would like to extend another special thanks to Daniel and Pedro. We started as classmates, study buddies, and now I consider you both as brothers. I thank you for all the advice, discussions, and sharing of experiences that we have been doing, and I hope to continue doing it throughout many years with you!

A special thank you to all the people who were important and had an impact on me throughout my PhD journey in ways not previously mentioned. First and foremost, to Daniela, who was and will continue to be, regardless of where life leads, a truly special person to me. She supported me profoundly on a personal level, and I am grateful for all the memories we shared, the advice we exchanged, and the many long and enjoyable walks we took over these years. Additionally, I would like to thank Helena and Joana for the valuable advice they provided at various points along the way, demonstrating that both physical and mental strength are essential to achieving success in life. Moreover, I would also like to extend my gratitude to professors António Paulo and João Tasso Sousa, Ana Rufino Ferreira, Catarina Pinto, and Héber for all the support they provided during the preparation for applying to the dual degree PhD program in Electrical and Computer Engineering from Carnegie Mellon University (CMU) at the end of 2020. Although the application was not successful, their help was essential in that endeavour, from providing recommendation letters to CV preparation, drafting motivation letters, and searching for potential supervisors on the CMU side. Their encouragement motivated me to pursue a PhD internship abroad, which led me to Rome, where I had the great privilege of working with Giorgio.

Last but not least, my deepest and most sincere thanks go to my family, my parents, José Fernando and Maria, my brother Rafael and his partner. Their support has been constant throughout my life. They have always been there for me, ready to help in any way I needed, and they gave me the freedom to explore, take risks, and

make mistakes that helped me grow. Work may come and go, as do friendships and professional connections, but what remains unwavering is their presence. For that, I am endlessly grateful.

Ricardo B. Sousa

*“I don’t stop when I’m tired.
I stop when I’m done.”*

David Goggins

Contents

Resumo	i
Abstract	iii
Acknowledgments	v
Abbreviations	xxi
1 Introduction	1
1.1 Context	1
1.2 Motivation	3
1.3 Problem Statement	4
1.4 Research Questions and Hypotheses	4
1.5 Scope and Objectives	6
1.6 Contributions	7
1.6.1 Main contributions	7
1.6.2 Complementary contributions	8
1.6.3 Public datasets	8
1.6.4 Open-source Software	9
1.7 Document Structure	9
2 Literature Review	11
2.1 Simultaneous Localization and Mapping (SLAM)	11
2.1.1 Problem introduction	11
2.1.2 Scan matching	13
2.1.3 Particle filter	13
2.1.4 Pose graph	15
2.2 Systematic Literature Review	17
2.2.1 Purpose of the study	17
2.2.2 Review methodology	19
2.2.3 Results overview	23
2.3 Long-Term SLAM in Dynamic Environments	25
2.3.1 Map representation	30
2.3.2 Map matching	31
2.3.3 Dynamic objects detection	33
2.3.4 Prediction modeling	36
2.4 Summary and Conclusions	38

3	Datasets	43
3.1	Mobile Platforms	43
3.1.1	5dpo three-wheeled omnidirectional robot	43
3.1.2	Modified Hangfa Discovery Q2 four-wheeled omnidirectional platform	43
3.1.3	IDEPA's tricycle autonomous robot	43
3.1.4	Avantgarde tricycle autonomous forklift	43
3.2	Faculdade de Engenharia da Universidade do Porto	43
3.3	IILABS 3D – iiLab Indoor LiDAR-based SLAM	43
3.4	Summary and Conclusions	43
4	Line Fitting-Based Corner-Like Detector for 2D Laser Scanners Data	45
4.1	Background	45
4.1.1	Line fitting intersection	46
4.1.2	Curvature-based detection	47
4.2	Corner-Like Feature Detector	47
4.2.1	Linear least-squares regression	48
4.2.2	Feature extraction	49
4.3	Experimental Results and Discussion	51
4.3.1	Feature extraction results	51
4.3.2	Features observability	52
4.3.3	Features association distance	52
4.3.4	Integration with a line-segment feature extraction algorithm	54
4.4	Summary and Conclusions	54
5	Point Cloud Registration with Distance Maps	57
5.1	Background	57
5.1.1	Point cloud registration	57
5.1.2	Least-squares minimization	60
5.2	Registration Methodology	68
5.2.1	Distance maps	68
5.2.2	Problem formulation	76
5.2.3	Point-to-point error formulation	78
5.2.4	Point-to-plane error formulation	78
5.3	Experimental Results and Discussion	81
5.3.1	Synthetic benchmarks	81
5.3.2	Real-world laser odometry evaluation	81
5.3.3	Computational performance	81
5.3.4	Discussion	81
5.4	Summary and Conclusions	81
6	Simultaneous Localization and Mapping (SLAM) in Dynamic Environments	83
6.1	Background	83
6.2	Online SLAM	83
6.3	Dynamics Removal	83
6.4	Offline Global Bundle Adjustment	83
6.5	Experimental Results and Discussion	83
6.6	Summary and Conclusions	83

7	Conclusions and Future Work	85
7.1	Conclusions and Contributions	85
7.2	Final Remarks and Future Work	85
7.2.1	Future Work	85
7.2.2	Involvement in the Scientific Community	85
	References	87
A	A Systematic Literature Review on Long-Term Localization and Mapping for Mobile Robots	103
B	Integrating Multimodal Perception into Ground Mobile Robots	105
C	Line Fitting-Based Corner-Like Detector for 2D Laser Scanners Data	107

List of Figures

1.1	The United Nations (UN) Sustainable Development Goals (SDGs)	8
2.1	Graphical model of the Simultaneous Localization and Mapping (SLAM) problem	12
2.2	Block diagram of the mapping and navigation systems of HectorSLAM	13
2.3	Composition of particles in FastSLAM in terms of path estimate and the set of estimators and associated covariances for individual feature locations	15
2.4	Edge connection in the graph formulation of the SLAM problem between vertexes x_i and x_j originated from a measurement z_{ij}	16
2.5	Preferred Reporting Items for Systematic reviews and Meta-Analyses (PRISMA) flow diagram for the selection process	21
2.6	Evolution of published records per year considering the 142 included records in the systematic literature review on long-term localization and mapping	23
2.7	Evolution of published records per year in Scopus and Web of Science databases on long-term localization and mapping	24
4.1	Error analysis of simple line regression models evaluating the fitting Root Mean Square Error (RMSE)	49
4.2	Example of the proposed corner-like feature extraction methodology on a laser point (x_j, y_j) at a distance scale $R_k = 0.15$ m	50
4.3	Example results for the proposed corner-like feature extraction methodology	52
4.4	Evaluation of the corner-like features observation ratio	53
4.5	Standard deviation analysis of the distance between corresponding corner-like features	53

List of Tables

2.1	Existent literature reviews, surveys, and tutorials on SLAM	18
2.2	Exclusion (E) criteria for the selection process of the systematic literature review	20
2.3	Quality Evaluation (QE) criteria and score range for the selection process of the systematic review	22
2.4	Data extraction items retrieved from literature works related to long-term localization and mapping in dynamic environments	26
5.1	Common M-estimator loss functions $\rho(s)$, their first-order derivative $\psi(s) = \rho'(s)$, and weight functions $w(s) = \psi(s)/(2s)$	66
5.2	Comparison of the three distance map variants (<i>Nano</i> , <i>Standard</i> , and <i>Sparse</i>) considered in the scope this PhD Thesis	76

List of Algorithms

5.1	Gauss-Newton minimization algorithm for Euclidean measurement and state spaces	63
5.2	Gauss-Newton minimization algorithm with robust estimators for Euclidean measurement and state spaces	66
5.3	Initialization procedure to resize the distance map data structures	73
5.4	Computation of the <i>Nano</i> distance map ($\mathcal{M}_{0,\text{dense}}$)	74
5.5	Computation of the <i>Standard</i> distance map ($\mathcal{M}_{1 2,\text{dense}}$)	75
5.6	Computation of the <i>Sparse</i> distance map ($\mathcal{M}_{1 2,\text{sparse}}$)	76

Abbreviations

3DGS 3D Gaussian Splatting. 36

ADRA AI, Data and Robotics Association. 3

AI Artificial Intelligence. 3, 18, 19

AMCL Adaptive Monte Carlo Localization. 51

AMR Autonomous Mobile Robot. 1–4, 11, 38

ARMA Auto-regressive Moving Average Model. 37, 40

ATE Absolute Trajectory Error. 30, 32, 35–37

BDVA Big Data Value Association. 3

BEV Bird-Eye-View. 32

BoW Bag-of-Words. 34

BPF Bayesian Persistence Filter. 37

CAD Computer-Aided Design. 31

CAGR Compound Annual Growth Rate. 1

CLAIRE Confederation of Laboratories for Artificial Intelligence Research in Europe. 3

CML Concurrent Mapping and Localization. 11

CNN Convolutional Neural Network. 34

CRF Conditional Random Field. 36

CRIS Centre for Robotics in Industry and Intelligent Systems. 7

DBN Dynamic Bayesian Network. 13, 14

DBSCAN Density Based Spatial Clustering of Applications with Noise. 46

DDE Dynamic Data Extraction. 25

DE Data Extraction. 22, 23

DEL Dynamic Edge Link. 33

DFR-SLAM Dynamic visual Feature Removal SLAM. 32

EKF Extended Kalman Filter. 14, 15, 17, 18, 36, 46

ELLIS European Laboratory for Learning and Intelligent Systems. 3

EU European Union. 3

EurAI European Association for Artificial Intelligence. 3

FEUP Faculty of Engineering, University of Porto. 7, 24

FFT Fast Fourier Transform. 40

FOV Field-of-View. 4, 51–55

FreMEn Frequency Map Enhancement. 37

FSH Feature Stability Histogram. 30

GHT Generalized Hough Transform. 59

GICP Generalized-ICP. 58, 59

GTSAM Georgia Tech Smoothing and Mapping. 15

- HMM** Hidden Markov Model. 36, 37, 40
- HOG** Histogram of Oriented Gradients. 35
- ICP** Iterative Closest Point. 6, 13, 31, 33, 38, 57–60, 68, 70, 78
- iiLab** Industry and Innovation Laboratory. 7
- IMU** Inertial Measurement Unit. 35
- INESC TEC** Institute for Systems and Computer Engineering, Technology and Science. 7, 24
- iSAM** incremental Smoothing And Mapping. 15
- ISDOR** Incremental Static-referenced Dynamic Object Removal. 35
- LC-CRF** Long-term Consistent Conditional Random Field. 32
- LiDAR** Light Detection And Ranging. 2, 13, 32–35, 37, 38
- LOAM** Lidar Odometry and Mapping. 34
- LTD** Long-Term Dynamics. 35
- LTM** Long-Term Memory. 2, 30, 34, 39
- LTS-NET** Long-Term Stability NETwork. 38
- MAP** Maximum A Posteriori. 60, 61
- MCL** Monte Carlo Localization. 31, 37, 38
- MLE** Maximum Likelihood Estimation. 16, 59, 61
- NDT** Normal Distribution Transform. 13, 33, 38, 58–60
- NeRF** Neural Radiance Field. 36
- NICP** Normal Iterative Closest Point. 13, 58, 59
- NN** Nearest Neighbor. 32, 68–70, 73–76
- ORB** Oriented FAST and Rotated BRIEF. 35
- PCA** Principle Component Analysis. 46, 58, 80
- PPS** Product, Processes, or Services. 4, 7
- PRISMA** Preferred Reporting Items for Systematic reviews and Meta-Analyses. 17, 19–21, 39
- RANSAC** RANdom SAmple Consensus. 46, 49
- RBPF** Rao-Blackwellized Particle Filter. 13, 14, 38
- RMSE** Root Mean Square Error. 34–36, 48, 49
- RNN** Recurrent Neural Network. 34, 35
- ROS** Robot Operating System. 9, 16, 46, 51, 54, 55
- RPE** Relative Pose Error. 36
- RProp** Residual Propagation. 59
- SDF** Signed Distance Field. 31, 68
- SDG** Sustainable Development Goal. 7, 8
- SGC-VSLAM** Semantic and Geometric Constraints Visual SLAM. 35
- SIVO** Semantically Informed Visual Odometry. 34
- SLAM** Simultaneous Localization and Mapping. 1, 2, 4–7, 9, 11–19, 21, 24, 34, 36, 38–41, 51, 55, 57, 59
- SM** Sensory Memory. 2, 30, 39
- SRIDA** Strategic Research, Innovation, and Deployment Agenda. 3
- STD** Short-Term Dynamics. 35
- STM** Short-Term Memory. 2, 30, 39
- SURF** Speeded Up Robust Features. 54
- SVD** Singular Value Decomposition. 46
- TORO** Tree-based netwORk Optimizer. 15

TSDf Truncated Signed Distance Field. 72

UN United Nations. 8

w.r.t. with respect to. 5, 7, 58, 59, 62, 67, 69, 71, 72, 77–79

WP Work Package. 3, 4, 7

Chapter 1

Introduction

This research, *RicoSLAM: localization and mapping in dynamic environments*, aims to advance the current state of the art in **Simultaneous Localization and Mapping (SLAM)**. This work focuses on improving localization and mapping performance when performing **SLAM** in dynamic environments, using wheeled odometry and 2D laser scanner data from sensors onboard industrial ground mobile robots. **SLAM** in dynamic environments is closely tied to lifelong **SLAM** and long-term operation of mobile robots, enabling robots to maintain accurate localization and mapping over extended periods while adapting to continuous environmental changes.

This chapter is divided into seven sections. First, Section 1.1 contextualizes this research work and describes the primary challenges related to localization and mapping in dynamic environments. Section 1.2 outlines the key motivations and reasons for undertaking this research. Next, Section 1.3 introduces the problem that this PhD Thesis seeks to address. Section 1.4 presents the main research questions and hypotheses that guided the research work. Then, Section 1.5 defines the scope and objectives of the PhD Thesis. Section 1.6 highlights the unique contributions this research makes to the literature, emphasizing its novelty and impact. Lastly, Section 1.7 outlines the document's organization. A final note is provided regarding the funding that supported the work carried out throughout the PhD Thesis.

1.1 Context

Nowadays, **Autonomous Mobile Robots (AMRs)** are being increasingly deployed in several industries, including logistics, retail, warehousing, automotive and aerospace manufacturing, agriculture, and healthcare. **AMRs** may automate tasks such as material transport, automated inspection, and security monitoring. By improving operational efficiency, flexibility and productivity, **AMRs** help companies achieve higher efficiencies and reduce production costs. Consequently, these improvements contribute to the growth of the global economy and increase the worldwide interest in the **AMR** market. According to a study by ABI Research, the global robotics market is estimated to be roughly US\$50 billion in 2025, growing at a **Compound Annual Growth Rate (CAGR)** of about 14% to reach approximately US\$111 billion by 2030. This market growth is driven by labour workforce shortages, supply chain reformation, and high turnover rates, with the concurrent increasing consumer demand. Mobile robots in particular can mitigate these challenges by expediting the loading/unloading of trucks, automating inventory management, and performing last-mile delivery services [1].

An **AMR** must interpret its surroundings by extracting meaningful data from its sensors to localize itself within the environment, enabling autonomous navigation [2]. **SLAM** addresses this perception and localization

challenge by concurrently incorporating the robot state estimation into the mapping process to build a representation of the environment as perceived by the robot. The state is typically described by the robot's pose (position and orientation), although other variables, such as velocity and calibration parameters, may also be included in the state. As for the environment representation, the mapping process involves creating a map incrementally, built as the robot explores unknown areas or refines it through passages in known locations, to describe the environment in which the robot operates. The map representation serves multiple purposes, such as path planning, providing a visual reference for human operators, while allowing the robot to reset its localization error when revisiting known areas [3].

In a static scene, where all elements of the environment remain fixed and unchanging over time, the robot would only need to map once because the map and sensor data would remain consistent with the environment throughout time. However, industrial and intralogistics locations, agricultural sites, warehouses, and service-oriented applications, such as shopping centers and healthcare institutions, have inherently moving elements in the scene (humans, objects), reconfiguration of the environment (reorganization of machines, patient beds), and appearance variations (lighting, weather, seasonal, or daytime changes). The deployment of **AMRs** in these varying conditions are a challenge for the **SLAM** system, where the system should decide how (consider the most current state, the evolution of the environment dynamics, or only the most static changes?) and when to update the map (when the variations occur, or only after a certain time?). This challenge is also known as the stability-plasticity dilemma, where a lifelong **SLAM** system should both adapt to new environment changes and preserve old states over time [4]. Furthermore, the map would grow indefinitely when gathering new information from the environment without any map management system, raising the question of how to bound map size (discard old information, compress representations, or limit the map growth?). This ever-growing problem poses another challenge for long-term **SLAM** due to the limited computational resources of the mapping vehicle. Indeed, the map size should be dependent on the environment area and not on the trajectory length of the mapping process nor operation time, only growing when the robot would explore unknown locations [5].

Currently, most methods rely on conventional **SLAM** map representations and then augment them with mechanisms to reason about dynamics. A recurring strategy for handling scene dynamics is to organize information into **Sensory Memory (SM)**, **Short-Term Memory (STM)**, and/or **Long-Term Memory (LTM)** [4, 6–9]. Another common trend is to partition the map into static or permanent, semi-static or movable, and dynamic layers, where static structures remain permanent over time, semi-static ones change only between visits at the same location, and dynamic layers represent sensor observations that move or are highly unstable in the environment [10–20]. Moreover, identification of scene dynamics is typically made through data consistency analysis, evaluating scan-to-map or scan-to-scan mismatches [6, 10, 11, 15, 17, 21–29], reprojection errors [19, 30–33], motion inconsistencies in tracked features [8, 19, 31, 33–36], ray-tracing discrepancies [4, 9, 11, 21, 22, 25, 37–39], and multi-session inconsistencies [24, 29, 38]. Clustering may improve dynamics identification through grouping nearby dynamic observations [7, 16, 28, 35, 38, 40, 41]. Object detection and semantic segmentation, typically from visual or 3D **Light Detection And Ranging (LiDAR)** data, may facilitate the identification of scene dynamics, especially when neural networks are used to classify objects in sensor observations, considering some of them dynamic, movable, or static [8, 12–14, 18–20, 33–36, 42–47]. Also, dynamic prediction may infer whether a region of the environment is stable, periodically occupied, or likely to change, to support more robust localization systems or even path-planning approaches that avoid highly dynamic areas [15, 16, 48–52]. Still, most methods rely primarily on static maps, excluding dynamic elements from the pose estimation process, and rarely maintain a continuously updated map representation.

1.2 Motivation

In May 2021, a joint partnership between major European organizations – **Big Data Value Association (BDVA)**, **Confederation of Laboratories for Artificial Intelligence Research in Europe (CLAIRE)**, **European Laboratory for Learning and Intelligent Systems (ELLIS)**, **European Association for Artificial Intelligence (EurAI)**, and **euRobotics** – founded the **AI, Data and Robotics Association (ADRA)**, formed from the European Partnerships in Cluster 4 (digital, industry, and space) in Horizon Europe. Since then, **ADRA**, in collaboration with the European Commission and other community members, has worked on the **Strategic Research, Innovation, and Deployment Agenda (SRIDA)** [53]. The most recent version, published in February 2024, reflects on European strategies and goals that **Artificial Intelligence (AI)**, data, and robotics can address on global challenges. This version outlines the **ADRA**'s Vision 2030, the global challenges and primary strategic objectives that the **ADRA** partnership seeks to address, major trends and gaps, as well as recommendations for the **European Union (EU)**'s 2025–2027 strategic plan in alignment with the **SRIDA**.

SRIDA [53] identifies the need for a new wave of **AI**, data, and robotics methods that can tackle complex societal and industrial challenges. Driven by the accelerating progress of **AMRs**, these challenges could lead to the development of self-supervised robot perception and autonomous control mechanisms, enabling predictive and safe robot operation without human intervention. Moreover, **SRIDA** acknowledges that autonomy in robotics is currently constrained by short-term, localized environments with controlled uncertainty. There is an urgent need for higher levels of autonomy, particularly in complex, unknown, and dynamic environments. Indeed, these environments often include unanticipated scenarios, which robots must be capable of handling and adapting to. Consequently, robots may present significant risks to both people and the environment, underscoring the need for reliable and trustworthy autonomous operations. As a result, **SRIDA** motivates research on improving localization and mapping processes for industrial mobile robots within dynamic environments, aiming to increase their levels of autonomy. This expansion of the range of scenarios where robots can operate without human intervention contributes to the long-term operation and reliable, lifelong autonomy of **AMRs**.

Furthermore, the logistics industry has high dynamic environments with constantly shifting elements, such as pallets, boxes, workers, moving **AMRs** (base platforms, rack and cart lifting robots, car pulling **AMRs**, moving pallet platforms, conveyor robots, mobile manipulators, pallet jack **AMRs**, tugger robots to tow carts like a convoy, and autonomous forklifts [54]), and manually operated mobile machines. Additionally, logistics warehouse environments may have featureless characteristics that make it difficult for **AMRs** to localize within the environment over time. One situation that is particularly critical to address in logistics environments involving robots' localization is cases involving random slotting warehouse strategies, which optimize space and enhance efficiency in modern operations. In random slotting, products are stored in any available space in the warehouse based on their characteristics (weight and size), thereby reducing replenishment time and increasing warehouse efficiency [55]. Subsequently, the interpretation and processing of dynamic changes in the environment by **AMRs** would enable the maintenance of an up-to-date map representation of the robots' surroundings in logistics environments. An up-to-date representation would allow the use of dynamic elements in the scene that are temporarily static, thereby improving the robot localization estimation in the environment.

Finally, the **Work Package (WP) 10** of the *GreenAuto – Green Innovation for the Automotive Industry* [56] project aims to develop technologies that reduce the production costs of electric vehicles, which are currently higher than those of traditional internal combustion engine alternatives. The use of autonomous logistic platforms on the factory floor, which either directly or indirectly assist in the production process, is already

well established. However, there is significant room for improvement, particularly in extending the applicability of those systems to less structured environments. WP10 proposes the development of a navigation system for AMRs that is robust in complex environments. One of the key deliverables from WP10 is the Product, Processes, or Services (PPS) 18, a navigation system for mobile robots designed to be both robust and efficient in automotive shop floor settings. Consequently, PPS18 motivates the research and development of localization and mapping algorithms that account for dynamic elements in the environment, aiming to enhance the robustness of the final autonomous system.

1.3 Problem Statement

As discussed in Section 1.1, a map representation in static environments does not require updates, as it remains consistent with the surrounding environment that has been explored. However, most use cases for AMRs involve dynamic environments with moving elements or reconfigurations that affect the consistency of the map over time. Consequently, the system must determine how and when to update the map, a challenge known as the stability-plasticity dilemma. In this context, a long-term SLAM system must adapt to new changes while preserving older features, particularly when the environment reverts to previous states [4].

This PhD Thesis addresses the problem of localizing a mobile robot within dynamic indoor environments. This problem is directly related to SLAM, given the need to update the map concurrently with localization upon dynamic changes in the robot's surroundings. Furthermore, this research is not limited to a specific type of dynamics, for example, only detecting semi-static dynamics – changes between running sessions, *i.e.*, changes occurring outside of the sensors' Field-of-View (FOV) at different time instants that, when observed again by the sensors, appear static but elements were effectively moved in the scene – or only dealing with highly dynamic elements (*e.g.*, moving persons or objects). Instead, this thesis studies the detection and removal of dynamic observations from the latest sensor data by comparing them against a local reference map, addressing both online removal and offline annotation of dynamic elements to prevent their integration into the map.

In terms of the sensor setup, this PhD Thesis assumes that the robot is equipped with at least one 2D laser scanner and, optionally, wheeled odometry data. This sensor setup is commonly found in AMRs within the logistics industry, where wheeled odometry is typically estimated by counting the ticks from the motor encoders while taking into account the robot's kinematic model [57]. In addition, AMRs often use 2D lasers not only for navigation but also as safety devices, in compliance with the ISO 10218-1:2025 [58] and ANSI/RIA R15.06-2025 [59] standards, which ensure a safe stopping distance and help prevent collisions with obstacles. Also, the usage of 2D lasers as safety devices allows AMRs to be compliant with ISO 3691-4:2023 [60], where, *e.g.*, an AMR must detect persons by passing on the two following tests: detecting a cylinder of 200 mm diameter and 600 mm length placed horizontally on the ground; and detect another cylinder of 70 mm diameter and 400 mm length set vertically. Consequently, this thesis focuses solely on range data.

1.4 Research Questions and Hypotheses

In response to the problem outlined in Section 1.3, this thesis proposes some research questions aimed at addressing and solving the identified problem. For each of these research questions, we formulate hypotheses that guided the investigation toward feasible solutions.

Research Question 1. *How to improve the mapping results when performing **SLAM** with a ground mobile robot in a dynamic indoor environment?*

In dynamic environments, **SLAM** algorithms can be affected by the presence of moving objects and changing surroundings. One key contributing factor to reduced **SLAM** accuracy is the inability of traditional static map representations to stay updated with real-time changes in the environment, leading to mismatches between current sensor observations and previously mapped data. As a result, the correspondence between these observations and the reference environment representation becomes weaker, leading to degraded map quality. Another issue arises from the growing amount of information to process over time. As the robot moves, the accumulation of data from sensors can lead to redundant information, especially in environments where sensor noise, matching errors, or outlier observations occur frequently. This redundancy manifests in the map as artificially thickened structures, commonly referred to as “fat” walls, where repeated scans acquired under slightly misaligned pose estimates cause point smearing along surfaces that should appear as sharp, thin boundaries in the map. Furthermore, as local pose errors accumulate over the robot’s trajectory, global map consistency degrades, reducing the likelihood of successfully closing loops when the robot revisits a previously mapped area. As a result, the global map may diverge over time from the true environment layout.

Given the previous challenges, we propose research hypotheses aimed at improving mapping results during **SLAM** in dynamic indoor environments. Dynamic observations, such as moving persons or displaced objects, act as outliers in the scan matching process, corrupting the pose estimation and increasing the drift accumulated over the robot’s trajectory. The primary focus is therefore on removing dynamic observations from the environment representation to reduce the noise introduced by both the sensor and the matching process, continuously updating the map with the latest sensor observations to increase the inliers matching ratio related to the correspondence between the current observation and the reference representation, while exploring concurrently offline global bundle adjustment techniques to remove distortions from the global map.

Hypothesis 1.1. *Exploring the application of ray tracing techniques to update local map representations online continuously, based on the latest sensor observations. These techniques can help determine whether a sensor observation corresponds to a new point or an update to an outdated one, depending on whether the latest ranging data is closer than the existing map point or if the sensor observation is passing through an existing feature. Also, the continuous online update of local maps based on the latest observations may improve the online estimation for the robot pose and its stability during **SLAM**, by increasing the inliers ratio of the sensor observation alignment *with respect to (w.r.t.)* local map.*

Hypothesis 1.2. *Researching techniques for robot relocalization on existing graphs representations, avoiding the addition of redundant data to the map. Linking graph growth to the spatial extent of the environment explored, this approach could help avoid unnecessary graph expansion, keeping the graph size manageable and improving computational efficiency without the graph size being dependent on the operation time.*

Hypothesis 1.3. *Investigating the offline removal of dynamic points using ray tracing techniques during the generation of the environment representation. This representation could be leveraged by other algorithms for more accurate mobile robot localization by excluding transient or moving objects from the map.*

Hypothesis 1.4. *Studying offline global map refinement techniques, removing dynamic points in the process, while trying to remove map distortions from the online-based map representation.*

Research Question 2. *How to integrate environment dynamics in the localization process when performing **SLAM** with a ground mobile robot in a dynamic indoor environment?*

The localization process in **SLAM** is also dependent on the quality of the sensor data, especially when dynamic elements are present in the environment. Moving objects, such as people, vehicles, or machinery, can introduce significant challenges during the data registration step. When these dynamic observations are considered during data registration without any outliers rejection technique, could introduce errors on pose estimation, potentially causing the robot's localization to drift or even fail entirely. In order to mitigate the effects of environment dynamics in the localization process, we explore several strategies formulated by the following research hypotheses. These hypotheses focus on extracting geometric features from the 2D laser scanner data, using robust kernels during the data registration process to reduce the influence of outliers while applying distance maps to avoid the traditional association step to speed up the data registration process. Additionally, implementing local map strategies to continuously update the local environment representation with the latest sensor observations can help maintain a more accurate knowledge of the robot's surroundings, improving localization stability in highly dynamic environments.

Hypothesis 2.1. *Exploring the use of geometric feature extraction, such as corners or lines, from 2D laser scanner data for localization purposes, rather than relying solely on raw ranging data. This approach may help filter out irrelevant data, providing more meaningful and stable features for localization, even in dynamic environments.*

Hypothesis 2.2. *Researching the use of distance map data structures for 2D point cloud registration. Distance maps can provide access to the nearest-neighbor point distance, enabling an **Iterative Closest Point (ICP)** [61]-like registration process without the need for the traditional association step. By applying robust kernels, this approach can improve the handling of dynamic environment point observations by considering them as outliers in the alignment optimization process or assigning them a lower weight.*

Hypothesis 2.3. *Exploring local map strategies to continuously update the local environment representation with the latest sensor observations. Unlike scan-based localization, which may be more prone to false positives or negatives, local maps facilitate loop closure detection through a scan-to-map alignment rather than scan-to-scan approaches that consider singular scans. Moreover, local maps can enhance the localization stability of mobile robots in highly dynamic environments, as they are continuously updated locally with the latest sensor observations, while allowing the incorporation of ray tracing-based techniques to detect dynamic elements and update the local representation accordingly.*

1.5 Scope and Objectives

Given the problem statement in Section 1.3, alongside the research questions outlined in Section 1.4, and considering the current advancements and state of the art in localization and mapping in dynamic environments as discussed in Chapter 2, this research thesis aims to explore and advance developments in **SLAM** using wheeled odometry and 2D laser scanner data, particularly focused on indoor industrial environments. We assume that the plant floor is planar, allowing just a 2D pose estimation to localize the robot in the environment.

In order to achieve the thesis's main goal, the following efforts were undertaken:

- modifications to existing mobile platforms to enable data acquisition for the purposes of this PhD Thesis;

- development of a line fitting-based corner-like feature extraction algorithm to capture corners geometric information from the environment on 2D laser scanners data;
- 2D point cloud data registration algorithms based on unsigned distance map data structures, proposing both point-to-point and point-to-plane error metrics;
- a full distance map-based graph **SLAM** framework, where the front-end estimates the robot pose *w.r.t.* a local reference representation (*e.g.*, singular scan or a local map), generating odometry and loop closure constraints, while the back-end of the framework implements a pose graph optimization using the `srrg2_solver` [62];
- ray tracing-based techniques to deal with dynamic observations on the 2D point cloud map representation of the environment;
- offline global bundle adjustment formulation based on distance maps to try to remove distortions from the global map.

Furthermore, while the degree-granting institution is the **Faculty of Engineering, University of Porto (FEUP)**, the host institution of this PhD Thesis is **INESC TEC – Institute for Systems and Computer Engineering, Technology and Science**. Specifically, all work conducted in this thesis was performed at **iiLab – Industry and Innovation Laboratory**, within **CRIIS – Centre for Robotics in Industry and Intelligent Systems** from **INESC TEC**. Throughout the PhD journey, the research was executed alongside **PPS18 – 3D navigation system for mobile robotic equipment**, part of **WP10 – automated logistics for the automotive industry**, within the *GreenAuto Agenda: Green Innovation for the Automotive Industry* [56], as a contribution for the long-term **SLAM** research topic. This **PPS** involves several partners, including **FEUP**, **INESC TEC**, and the leading partner, Flowbotic Mobile Systems [63], responsible for the development of an autonomous forklift for the automotive industry.

Within the scope of the 2030 Agenda for sustainable development (see Figure 1.1), agreed by 195 nations in September 2015 [64], this research work falls under the following **Sustainable Development Goals (SDGs)**:

- **SDG8**: Decent work and economic growth;
- **SDG9**: Industry, innovation and infrastructure;
- **SDG12**: Responsible Consumption and Production.

1.6 Contributions

This doctoral research advances localization and mapping for mobile robots in dynamic industrial environments by combining state of the art research, experimental evaluation, and system-level integration. The work resulted in peer-reviewed publications covering perception and feature extraction from 2D laser scanners, multimodal sensing integration, localization in dynamic environments, and **SLAM** benchmarking. Beyond the scientific contributions, this thesis also includes the public release of a dataset to support reproducibility and benchmarking, as well as the development of open-source software.

1.6.1 Main contributions

- [65] R.B. Sousa, H.M. Sobreira, and A.P. Moreira. “A systematic literature review on long-term localization and mapping for mobile robots”. In: *Journal of Field Robotics* 40.5 (2023), pp. 1245–1322. DOI: [10.1002/rob.22170](https://doi.org/10.1002/rob.22170).



Figure 1.1: The United Nations (UN) Sustainable Development Goals (SDGs) [64].

- [66] R.B. Sousa, H.M. Sobreira, M.F. Silva, and A.P. Moreira. “Line fitting-based corner-like detector for 2D laser scanners data”. In: *2024 10th International Conference on Automation, Robotics and Applications (ICARA)*. Athens, Greece, 2024, pp. 210–216. doi: [10.1109/ICARA60736.2024.10553006](https://doi.org/10.1109/ICARA60736.2024.10553006).
- [67] R.B. Sousa, H.M. Sobreira, J.G. Martins, P.G. Costa, M.F. Silva, and A.P. Moreira. “Integrating multimodal perception into ground mobile robots”. In: *2025 IEEE International Conference on Autonomous Robot Systems and Competitions (ICARSC)*. Funchal, Portugal, 2025, pp. 104–111. doi: [10.1109/ICARSC65809.2025.10970176](https://doi.org/10.1109/ICARSC65809.2025.10970176).

1.6.2 Complementary contributions

- [68] J.D. Ribeiro, R.B. Sousa, J.G. Martins, A.S. Aguiar, F.N. dos Santos, and H.M. Sobreira. “Indoor benchmark of 3-D LiDAR SLAM at iilab – industry and innovation laboratory”. In: *IEEE Access* 13 (2025), pp. 212421–212442. doi: [10.1109/ACCESS.2025.3643753](https://doi.org/10.1109/ACCESS.2025.3643753).
- [69] M.S. Lopes, A. Cordeiro, R.B. Sousa, J.A. Beça, P. Costa, J.P.C. Souza, and M.F. Silva. “CARGO: container-unloading autonomous robot with generalized operation”. In: *2026 12th International Conference on Automation, Robotics and Applications (ICARA)*. Istanbul, Turkey, 2026, pp. 1–6.

1.6.3 Public datasets

- [70] J.D. Ribeiro, R.B. Sousa, J.G. Martins, A.S. Aguiar, F.N. dos Santos, and H.M. Sobreira. *IILABS 3D: iilab indoor LiDAR-based SLAM dataset*. iilab – industry and innovation lab: INESC TEC – Institute for Systems, Computer Engineering, Technology, and Science, 2025. doi: [10.25747/VHNJ-WM80](https://doi.org/10.25747/VHNJ-WM80).

1.6.4 Open-source Software

- https://github.com/INESCTEC/inesctec_mrdt_line_fit_based_corner_detector_2d_laser: C++ implementation compatible with **Robot Operating System (ROS)** 1 of the line fitting-based corner-like detector for 2D lasers [66];
- https://github.com/INESCTEC/inesctec_mrdt_hangfa_discovery_q2: repository containing all the documentation (hardware, 3D modelling, electronics, network configuration, sensor setup) and code (firmware, C++ nodes compatible with both **ROS** 1 and **ROS** 2) associated with the modifications made on the Hangfa Discovery Q2 omnidirectional platform [71] to integrate multimodal perception [67];
 - https://oss.inesctec.pt/inesctec_mrdt_hangfa_discovery_q2/: GitHub.io public website to help researchers replicate the multimodal perception integration discussed in [67].
- <https://github.com/sousarbarb/slr-lt-lm-mr>: repository containing all the documentation, data, and Python scripts used in the scope of the systematic literature review on long-term localization and mapping [65];
- https://github.com/sousarbarb/hector_slam: added an offline node to the original Hector **SLAM** [72] open-source repository (https://github.com/tu-darmstadt-ros-pkg/hector_slam) to process multiple **ROS** bag files, generate the 2D occupancy grid map, and save the robot pose data into TUM format [73];
- https://github.com/sousarbarb/slam_gmapping: same addition made to the `hector_slam` **ROS** package for the GMapping [74] open-source implementation;
- https://github.com/sousarbarb/stage_ros: native **ROS** wrapper for the Stage Simulator [75–77], added compatibility with both **ROS** 1 and **ROS** 2 versions;
- https://gitlab.com/sousarbarb_srrg-software/srrg_system: added CMake compatibility with both **ROS** 1 and **ROS** 2 versions of the packages `srrg2_core`, `srrg2_qgl_viewport`, and `srrg2_solver` [62].

1.7 Document Structure

The remaining of this document is organized as follows (**[TODO: update later the following text]**):

- **Chapter 2** presents the **SLAM** problem, covering scan matching, particle filters, and pose graph formulations. The chapter also reviews literature on long-term **SLAM** in dynamic environments, including map representation, map matching, dynamic object detection, and dynamics prediction;
- **Chapter 3** describes the various mobile platforms and datasets used in this PhD Thesis, including robots such as the 5dpo omnidirectional robot, the Avantgarde tricycle and Flowbotics differential autonomous forklifts, and also details the IILABS 3D and other datasets from the literature;
- **Chapter 4** presents a corner-like feature detection algorithm for 2D laser-scanner data and evaluates its geometric invariance to various viewpoints;
- **Chapter 5** explores the concept of using distance maps for data registration, discussing optimization algorithms, covariance estimation, and point-to-point and point-to-plane alignment error formulations;
- **Chapter 6** presents the research conducted in this PhD Thesis on **SLAM** in dynamic environments, focusing on techniques such as online **SLAM**, dynamics removal, and offline global bundle adjustment;
- **Chapter 7** summarizes the conclusions and contributions of this research and offers final remarks on future work, including potential areas for further exploration and involvement in the scientific community.

Funding

This work has received funding from National Funds through the Portuguese funding agency, FCT – Fundação para a Ciência e a Tecnologia, within scholarship 2021.04591.BD (DOI: [10.54499/2021.04591.BD](#)) and project LA/P/0063/2020 (DOI: [10.54499/LA/P/0063/2020](#)). This work has also been co-financed by Component 5 – Capitalisation and Business Innovation, integrated in the Resilience Dimension of the Recovery and Resilience Plan within the scope of the Recovery and Resilience Mechanism (MRR) of the European Union (EU), framed in the Next Generation EU, for the period 2021–2026, within project GreenAuto, with reference 54.

Chapter 2

Literature Review

This chapter discusses the current literature on **SLAM** in dynamic environments. Our systematic literature review [65] of the broader topic of long-term localization and mapping (see Appendix A) analyzes 142 works covering appearance invariance, modeling the environment dynamics, map size management, multi-session, and computational topics such as parallel computing and timing efficiency. The review article [65] serves as the foundation for this chapter’s discussion. While this chapter adopts a similar systematic methodology, the discussion focuses on the scope of this PhD Thesis, namely, how the literature addresses scene dynamics in localization and mapping processes. As a result, this chapter extends the discussion of dynamics modelling presented in the review [65] and updates the results of the search inquiry to scientific databases, yielding an analysis of 53 studies, compared with the 32 originally covered in that review related to dynamics modelling.

Hence, this chapter is divided into four sections. First, Section 2.1 formulates the traditional **SLAM** problem and presents three of its most common solutions: scan matching, particle filtering, and pose graph optimization. Next, Section 2.2 contextualizes the elaboration of a systematic literature review on long-term localization and mapping within the scope of this PhD Thesis, covering the main purpose of the study, the review methodology, and the results overview. Section 2.3 reviews the current literature on long-term localization and mapping in dynamic environments, focusing on how the literature addresses scene dynamics through map representation, map matching, dynamic object detection, and prediction modelling. Finally, Section 2.4 presents a summary of this chapter and the main conclusions.

2.1 Simultaneous Localization and Mapping (SLAM)

2.1.1 Problem introduction

SLAM is one of the most fundamental problems in robotics. This problem, also previously known as **Concurrent Mapping and Localization (CML)**, may be formulated by Definition 2.1.1. When the robot does not have access to a representation of the environment nor know its pose, **SLAM** enables an **AMR** to build a map of its surroundings while simultaneously localizing itself. **SLAM** may also be used to build a map in applications assuming a static scene, even if the map is only continuously updated upon the mapping process [78].

Definition 2.1.1 ([2, p. 349]). “...the robot must not only create a map, but it must also do this while moving and localizing to explore the environment. In the robotics community, this is often called the *Simultaneous Localization and Mapping (SLAM) problem*.”




Figure 2.1: Graphical model of the **Simultaneous Localization and Mapping (SLAM)** problem: (a) online SLAM; (b) full SLAM [78]. Light gray circles represent observed inputs (measurements $z_{1:t}$ and control variables $u_{1:t}$), while nodes enclosed in a rounded rectangle indicate the jointly estimated variables: in (a), only the current pose x_t and the map m ; in (b), the entire trajectory $x_{1:t}$ alongside the map m .

Two main formulations define **SLAM** from a probabilistic perspective, also illustrated in Figure 2.1. The first is online **SLAM**, where the algorithm estimates the posterior over the momentary pose (x_t at a time instant t) along with the map (m), as defined in Equation (2.1), provided the sensor measurements ($z_{1:t}$) and the control variables ($u_{1:t}$). The online term is due to only involving the estimation of variables that persist at the current time instant t . Several algorithms incrementally implement online **SLAM**, *i.e.*, discarding past measurements and controls once they have been processed. The other probabilistic formulation is called full **SLAM**, given the goal to compute a posterior over the entire path ($x_{1:t}$) along with the map, as formulated in 2.2, instead of just the current pose (x_t). The difference between online and full **SLAM** is that the former is obtained by marginalizing out the past poses ($x_{1:t-1}$) from the full **SLAM** posterior [78].

$$p(x_t, m \mid z_{1:t}, u_{1:t}) \quad (2.1)$$

$$p(x_{1:t}, m \mid z_{1:t}, u_{1:t}) \quad (2.2)$$

In the case of the algorithm requiring correspondences between measurements and map landmarks, the inclusion of the correspondence variables (c_t or $c_{1:t}$) in the estimation problem may be useful for reasoning if the perceived feature is the same one as previously detected and already represented in the map or if it is a new landmark. Consequently, Equation (2.3) and Equation (2.4) formulate the posterior estimation for online and full **SLAM**, respectively [78].

$$p(x_t, m, c_t \mid z_{1:t}, u_{1:t}) \quad (2.3)$$

$$p(x_{1:t}, m, c_{1:t} \mid z_{1:t}, u_{1:t}) \quad (2.4)$$

Although computing a full posterior captures all the information to be known about the map and the pose or the path, this estimation is practically unfeasible due to the high dimensionality of the continuous parameter space and the large number of discrete correspondence variables [78]. So, next, the discussion on scan matching, particle filter, and pose graph implementations of **SLAM** focuses on the practical aspects of performing continuous localization and mapping.




Figure 2.2: Block diagram of the mapping and navigation systems of HectorSLAM [72].

2.1.2 Scan matching

A basic approach for solving the **SLAM** problem is continuously matching the sensor scans to the map learned until the current instant. An example is HectorSLAM [72] algorithm, where the block diagram describing that system is shown in Figure 2.2. This algorithm is considered to be online **SLAM** due to only estimating the current pose and the map. Also, the method integrates incrementally the sensor readings to allow discarding those in subsequent iterations. Furthermore, HectorSLAM [72] aligns the beam endpoints of a 2D **LiDAR** with the current map using a matcher inspired by an alignment algorithm based on computer vision [79]. The matcher estimates the homogeneous transformation between the sensor scan and the map (*i.e.*, scan-to-map alignment) to update the latter with the latest scan, while also updating the current estimation of the robot pose. Although the matcher is specific to the HectorSLAM [72] algorithm, it does not invalidate using other matchers, such as the PerfectMatch [80], **ICP** [61], **Normal Iterative Closest Point (NICP)** [81], **Normal Distribution Transform (NDT)** [82], or others, to obtain the transformation between the 2D **LiDAR** scan and the map, allowing to perform **SLAM** with a scan matching methodology.

The main advantage of scan matching **SLAM** algorithms is the lower use of computational resources compared to more complex approaches. These algorithms do not require an exhaustive pose search or data association search with previous scans (the matching with all preceding scans is implicit in matching the current one with the map) [72]. However, as in the case of HectorSLAM [72], a pure scan matching approach does not implement loop closure (see Definition 2.1.2). This limitation prevents correcting the map when recognizing a previously mapped location, which means there is no reduction in the uncertainty of the map estimate, nor any correction for map distortions caused by incremental alignment errors.

Definition 2.1.2 ([2, p. 188]). “The problem of location recognition is also called *loop-detection*, because a *loop* is a closed trajectory of a vehicle that returns to a previously-visited point. The problem of minimizing the error at the loop closure is instead called *loop-closing*.”

2.1.3 Particle filter

The graphical model of the **SLAM** problem shown in Figure 2.1 is equivalent to a discrete **Dynamic Bayesian Network (DBN)** (see Definition 2.1.3). A sampling-based inference algorithm for **DBNs** is the particle filter. However, these filters suffer from the problem of dimensionality, scaling exponentially [83]. The assumption of conditional independence between map measurements allows using **Rao-Blackwellized Particle Filters (RBPFs)** to solve the **SLAM** problem. These filters use particles to represent the posterior probability over some variables

(e.g., the robot path, correspondence variables) and Gaussians (or other parametric probabilistic distribution functions) to represent all other variables (e.g., the feature positions) [78]. This marginalization of some variables implemented by **RBPFs** improves the computational efficiency compared to standard particle filters, reducing the space size for the sampling step [83].

Definition 2.1.3 ([84]). “A convenient way to describe this structure is via the *Dynamic Bayesian Network (DBN)*... A Bayesian network is a graphical model that describes a stochastic process as a directed graph. The graph has one node for each random variable in the process, and a directed edge (or arrow) between two nodes models a conditional dependence between them”

An example of a **RBPF-based SLAM** is FastSLAM [85, 86]. This algorithm decomposes the **SLAM** problem into a robot localization one and the estimation of map landmarks conditioned on the robot pose estimate. The decomposition and factoring of the posterior probability ($p(x_t, m \mid z_{1:t}, u_{1:t}, c_{1:t})$) followed by the FastSLAM [85, 86] algorithm is shown in Equation (2.5). The algorithm assumes the conditional independence of the features (m_n) given the robot path ($x_{1:t}^{[M]}$), i.e., knowing the location of one feature does not add any additional information about the location of other features if the robot path is known. Moreover, Figure 2.3 represents the composition of particles implemented in FastSLAM. Each particle k contains an estimate of the robot path until the current time instant t and a set of **Extended Kalman Filters (EKFs)** with a mean ($\mu_n^{[k]}$) and covariance ($\Sigma_n^{[k]}$) of the estimation, one for each feature $n = 1..N$ represented in the map m . The basic steps of the FastSLAM [85, 86] algorithm are the following ones [78]:

- for each particle $k = 1..M$:
 1. **retrieval**: retrieve a pose $x_{t-1}^{[k]}$ from the particle set;
 2. **prediction**: sample a new pose $x_t^{[k]} \sim p(x_t \mid x_{t-1}^{[k]}, u_t)$ using the probabilistic motion model of the robot (the improved version, FastSLAM 2.0 [86], incorporates also the sensor measurements z_t^i in the prediction step);
 3. **measurement update**: for each observed feature obtained from a measurement z_t^i , identify the correspondence n in the particle map, and incorporate z_t^i into the corresponding **EKF**, updating the mean ($\mu_n^{[k]}$) and covariance ($\Sigma_n^{[k]}$) of the corresponding map feature accordingly;
 4. **importance weight**: compute the importance weight $w^{[k]}$ for the particle.
- **resampling**: sample, with replacement, the M particles, where the probability of sampling being proportional to $w^{[k]}$, to incorporate the measurements z_t^i into the robot pose $x_t^{[k]}$ (due to only considering the most recent control u_t in the previous steps, only in FastSLAM 1.0 [85]).

$$\begin{aligned}
 p(x_t, m \mid z_{1:t}, u_{1:t}, c_{1:t}) &= p(x_t \mid z_{1:t}, u_{1:t}, c_{1:t}) p(m \mid x_t, z_{1:t}, c_{1:t}) = \\
 &= p(x_t \mid z_{1:t}, u_{1:t}, c_{1:t}) \prod_{n=1}^N p(m_n \mid x_t, z_{1:t}, c_{1:t})
 \end{aligned} \tag{2.5}$$

The adaptation of FastSLAM [85, 86] to occupancy grid-based **SLAM** is implemented by the GMapping [74] algorithm. This adaptation requires the modification of three modules relative to FastSLAM [85, 86]. First, GMapping [74] considers the occupancy grid as the map representation for the **SLAM** problem, where each particle contains a grid map. Next, the computation of a measurement likelihood ($p(z_t \mid x_t^{[k]}, m_{t-1}^{[k]})$) uses the beam endpoint sensor model, characteristic of range-bearing sensors such as 2D lasers, instead of integrating the measures into an **EKF** as in FastSLAM [85, 86]. This model considers the individual scan beams independent from each other, using the distance between the beam endpoint and the closest obstacle in the map to that point

	robot path	feature 1	feature 2	...	feature N
Particle $k = 1$	$x_{1:t}^{[1]} = \{(x \ y \ \theta)^T\}_{1:t}^{[1]}$	$\mu_1^{[1]}, \Sigma_1^{[1]}$	$\mu_2^{[1]}, \Sigma_2^{[1]}$	$\dots$	$\mu_N^{[1]}, \Sigma_N^{[1]}$
Particle $k = 2$	$x_{1:t}^{[2]} = \{(x \ y \ \theta)^T\}_{1:t}^{[2]}$	$\mu_1^{[2]}, \Sigma_1^{[2]}$	$\mu_2^{[2]}, \Sigma_2^{[2]}$	$\dots$	$\mu_N^{[2]}, \Sigma_N^{[2]}$
		$\vdots$			
Particle $k = M$	$x_{1:t}^{[M]} = \{(x \ y \ \theta)^T\}_{1:t}^{[M]}$	$\mu_1^{[M]}, \Sigma_1^{[M]}$	$\mu_2^{[M]}, \Sigma_2^{[M]}$	$\dots$	$\mu_N^{[M]}, \Sigma_N^{[M]}$

Figure 2.3: Composition of particles in FastSLAM [85, 86] in terms of path estimate and the set of estimators and associated covariances for individual feature locations [78].

to compute the beam's individual likelihood. Lastly, the update of each particle's map ($p(m_t^{[k]} | z_t, x_t^{[k]}, m_{t-1}^{[k]})$) computes a new occupancy grid map. This update estimates the probability of each individual cell, assuming the independence of each other, using a binary Bayes filter and the inverse model of the sensor measurements [78].

A key advantage of particle filter-based **SLAM** is the decisions on data associations being made on a per-particle basis, instead of the most likely one as in scan matching or **EKF** approaches. This characteristic represents inherently multi-hypotheses associations in the filter, making it more robust to data association problems. Another advantage is particle filters being capable to model non-linear robot motions without requiring model approximations. However, a disadvantage of particle filters is relative to loop closure detection. Even though these filters are able to detect previously visited places, unlike scan matching approaches such as HectorSLAM [72], this place recognition ability is only maintained through the diversity in the particle sets, possibly maintaining larger sets and requiring more computational resources [78].

2.1.4 Pose graph

Another solution for the **SLAM** problem leveraging its graphical model (illustrated in Figure 2.1) is pose graph **SLAM**, also known as GraphSLAM. This optimization-based solution consists of two subsystems, as in filter-based approaches such as particle filters and **EKF**. In filter-based **SLAM**, the filters implement a prediction stage usually based on the vehicle's motion model, and then, the correction phase updates the predicted state by matching the sensor observation to the map. As for optimization-based solutions such as the pose graph formulation, a first module also known as the front-end is responsible to identify constraints (*e.g.*, odometry, loop closure, landmark observations) of the problem based on sensor data by finding correspondences between new observations and the current map. Then, a second subsystem known as the back-end computes or refines the poses of the graph and the map minimizing the coherence errors as a whole. This second module usually requires optimization algorithms such as Gauss-Newton or Levenberg-Marquardt for minimizing the objective function [87]. However, there are available in the literature optimization frameworks more specific for the pose graph formulation, namely, **Tree-based netwORk Optimizer (TORO)** [88], **incremental Smoothing And Mapping (iSAM)** [89, 90], **Georgia Tech Smoothing and Mapping (GTSAM)** [91, 92], g^2o [93], or `srrg2_solver` [62].

Furthermore, the edge connection in GraphSLAM between vertexes x_i and x_j is illustrated in Figure 2.4. Considering z_{ij} as a virtual measurement obtained from raw sensor data between those two vertexes, let $\bar{z}_{ij}$ and




Figure 2.4: Edge connection in the graph formulation of the **SLAM** problem between vertexes x_i and x_j originated from a measurement z_{ij} [84].

Ω_{ij} be the mean and the information matrix of the virtual data, respectively. This virtual measurement actually represents the transformation that maximizes the overlap between the observations acquired from the pose nodes i and j . The log-likelihood l_{ij} (see Definition 2.1.4) of the measurement z_{ij} can be defined as in Equation (2.6), where $\hat{z}_{ij}(x_i, x_j)$ is the prediction of the virtual measurement given the configuration of the nodes x_i and x_j . Then, the configuration of nodes $x^* = (x_1, \dots, x_T)^T$ can be found formulating a **Maximum Likelihood Estimation (MLE)** problem that minimizes the negative log-likelihood function $F(x)$ (see Equation (2.7)) of all the observations, seeking to solve the Equation (2.8) [84]. For exemplification purposes, two possible examples of virtual measurements are wheeled odometry and 2D laser scan matching data. Both cases would define the prediction of the measurement as $\hat{z}_{ij} = x_j - x_i$. The virtual measurement would be $z_{ij} = g(x_i, u_{ij})$ for wheeled odometry, where $g(x_i, u_{ij})$ represents the relative transformation between x_i and x_j estimated by the odometry estimator, while using the wheeled odometry data u_{ij} between those two nodes. As for the scan matching case, the measurement would be $z_{ij} = h(z_i, z_j)$ representing the matching between the laser scans z_i and z_j acquired from the nodes x_i and x_j , respectively.

Definition 2.1.4 ([94, p. 224]). “...since the natural logarithm is a one-to-one function on $\mathbb{R}^+$, no information about θ is lost when the log-likelihood function $l(\theta) = \ln[p(\theta)]$ is used in place of $p(\theta)$. Moreover, for many probability models, it is easier to work with the log-likelihood function than with the likelihood function...”

$$l_{ij} \propto [z_{ij} - \hat{z}_{ij}(x_i, x_j)]^T \Omega_{ij} [z_{ij} - \hat{z}_{ij}(x_i, x_j)] \quad (2.6)$$

$$F(x) = \sum_{i,j} e_{ij}^T \Omega_{ij} e_{ij} \quad (2.7)$$

$$x^* = \underset{x}{\operatorname{argmin}} F(x) \quad (2.8)$$

A pose graph **SLAM** implementation is the **SLAM** Toolbox framework [95], also available in the official navigation stack of the **ROS** [96] (in version 2). This toolbox is intended for 2D **SLAM** using a 2D laser scanner in conjunction with wheeled odometry obtained from the wheel encoders of the robot's motors. The optimization framework used for the pose graph back-end is the Ceres Solver [97]. Although **SLAM** Toolbox [95] states

having lifelong mapping capabilities, these are only relative to being able to update existing maps and serialize the data for use in other mapping sessions, not for dynamic environments or appearance invariance.

Overall, a key difference between pose graph **SLAM** compared to other solutions for the **SLAM** problem, such as filter and scan matching approaches, is that the former has access to the full data when building the map, being able to apply improved linearization and data association techniques. Indeed, in GraphSLAM, all data can be used to linearize and calculate correspondences, *i.e.*, can revise past data association and linearize more than once [78]. This key difference is pointed as a possible reason to graph **SLAM** tending to give better results than filters, where the latter is more subject to linearization issues [87]. However, the pose graph formulation shows a linear growth over time in the graph's size requiring techniques for constraining the graph size, and thus, limit the computational resources required for continuously performing GraphSLAM [78].

2.2 Systematic Literature Review

Next, our systematic literature review on long-term localization and mapping for mobile robots [65] (see Appendix A) analyzes 142 works, following the **Preferred Reporting Items for Systematic reviews and Meta-Analyses (PRISMA)** [98] guidelines. Instead of focusing on a single challenge, the review [65] covers several topics, including robustness to appearance changes, dynamics modelling, map sparsification to cope with the continuously increasing size of environment representations, multi-session mapping, and additional computational issues, such as parallel and distributed computing and memory management. The review [65] further provides a comparative analysis of the public datasets and experimental data adopted in the selected studies, focusing on scene conditions, sensing modalities, and their spatial and temporal properties. In addition, we made available a public GitHub repository¹ containing all the documentation and scripts used throughout the review process to enhance the repeatability and reproducibility of our review by the scientific community.

Although our review [65] covers a broader topic than the scope of this PhD Thesis, this section presents the systematic methodology followed in the review that underpins the discussion in Section 2.3. Lastly, this section examines the growing interest in long-term localization and mapping, based on the original search query used in the review, while adapting the query to investigate further the growing interest in handling scene dynamics.

2.2.1 Purpose of the study

The main studies reviewing the **SLAM** literature are presented in Table 2.1. In terms of introducing the problem formulation, the studies focus on explaining different frameworks for performing **SLAM**. Durrant-Whyte and Bailey [99, 100] discuss Bayesian-based probabilistic formulations, namely, **EKF** and the Rao-Blackwellized particle filter frameworks, also categorized as filtering approaches. In contrast, Grisetti *et al.* [84] explain in detail a smoothing formulation, the graph-based **SLAM**, which estimates the full trajectory of the robot from a set of sensor measurements. Thrun [101] introduces both probabilistic (**EKF** and particle filter) and smoothing (graph) formulations. Focusing on vision sensorization, Scaramuzza and Fraundorfer [102, 103] present an extensive tutorial on visual odometry for estimating relative motion from visual data, where the study discusses camera modeling and calibration, motion estimation, and feature matching. Yousif *et al.* [104] extend the previous visual odometry tutorial to include methodologies for the vision-based **SLAM** problem. Although

¹<https://github.com/sousarbarb/slr-ltlm-mr>

Table 2.1: Existent literature reviews, surveys, and tutorials on **SLAM** [65].

Year	Topic	Reference
2006	Probabilistic formulations (EKF , particle filter)	Durrant-Whyte and Bailey (2006) [99], Bailey and Durrant-Whyte (2006) [100]
2008	Probabilistic and pose graph formulations	Thrun (2008) [101]
2010	Graph SLAM	Grisetti <i>et al.</i> (2010) [84]
2011	Observability, convergence, consistency (feature-based SLAM)	Dissanayake <i>et al.</i> (2011) [105]
2012	Visual odometry	Scaramuzza and Fraundorfer (2011) [102], Fraundorfer and Scaramuzza (2012) [103]
2014	Underwater navigation and localization	Paull <i>et al.</i> (2014) [106]
2015	Visual SLAM	Yousif <i>et al.</i> (2015) [104]
2016	Observability, convergence, consistency (feature and graph-based SLAM)	Huang and Dissanayake (2016) [107]
	Multi-robot SLAM	Saeedi <i>et al.</i> (2016) [108]
	Visual place recognition	Lowry <i>et al.</i> (2016) [109]
	SLAM literature overview	Cadena <i>et al.</i> (2016) [3]
2017	Autonomous vehicles	Bresson <i>et al.</i> (2017) [87]
2018		Kuutti <i>et al.</i> (2018) [110]
2018	AI for long-term autonomy	Kunze <i>et al.</i> (2018) [111]
	Long-term sensorization	Zaffar <i>et al.</i> (2018) [112]
2020	Deep learning	Fayyad <i>et al.</i> (2020) [113]
	Multi-robot search and rescue	Queralta <i>et al.</i> (2020) [114]
2021	Self-driving vehicles	Badue <i>et al.</i> (2021) [115]
2022	Underground navigation	Ebadi <i>et al.</i> (2024) [116]

the introductions mentioned here provide comprehensive explanations of the problem formulation itself, none of them focus the discussion on the possible long-term challenges of **SLAM**.

Moreover, other studies focus on theoretical aspects of the **SLAM** formulation. Dissanayake *et al.* [105] discuss the fundamental properties of **SLAM**: observability (whether the problem is solvable), convergence (whether the state uncertainty converges to a finite value), and consistency (whether the estimated state is unbiased). Huang and Dissanayake [107] extend the previous work [105] by providing an in-depth explanation of the fundamental properties and defining criteria for the performance evaluation of **SLAM** algorithms in terms of consistency, accuracy, and computational efficiency. In contrast to [105], which focuses mainly on filtering-based **SLAM**, Huang and Dissanayake [107] also discuss the properties for smoothing formulations. Both studies focus only on theoretical aspects, without addressing long-term localization and mapping.

In terms of surveys on **SLAM** trends, Cadena *et al.* [3] provide a broad overview of metric and semantic **SLAM** works. The work briefly discusses localization and mapping robustness in terms of loop closure validation and handling dynamic environments, and **SLAM** scalability in terms of pose graph sparsification and parallel and distributed computing. By contrast, Lowry *et al.* [109] focus on topological **SLAM**, providing a comprehensive review of visual place recognition. Although the study discusses the challenges of navigation across varying conditions, it is limited to vision sensors. Bresson *et al.* [87] survey trends in single- and multi-vehicle **SLAM** and in large-scale experiments for autonomous vehicles. Although the study compares the methods on accuracy, scalability, availability, recovery, map updateability, and scene dynamics, their study only refers to approaches composed of at least odometry and a mapping module, not discussing localization-only

algorithms. Also, the discussion is more focused on loop closure and relocalization modules and on leveraging existing data, rather than on methodologies for addressing the long-term challenges of continuous **SLAM**.

Similar to [87], Kuutti *et al.* [110] and Badue *et al.* [115] analyze trends for self-driving vehicles. While Kuutti *et al.* [110] focus on sensorization and cooperative localization between vehicles, Badue *et al.* [115] concentrate on the architecture of autonomous driving systems. However, neither of those studies discuss the challenges of achieving long-term **SLAM**. Saeedi *et al.* [108] present a review of multi-robot **SLAM**. Even though the authors identify large-scale, dynamic environments, multi-session operations, and agent scalability as challenges for multi-robot **SLAM**, their study does not provide an overview of existing methodologies to address these problems. Paull *et al.* [106] and Ebadi *et al.* [116] also present overviews of the **SLAM** literature, but are more specific in their domains: the former focuses on autonomous underwater navigation, and the latter on **SLAM** in extreme underground environments. However, those two studies do not discuss long-term **SLAM**.

The surveys by Kunze *et al.* [111] and Zaffar *et al.* [112] focus on some challenges of long-term autonomy. Kunze *et al.* [111] provides a brief overview of how **AI** can enable the long-term operation of autonomous systems. Although the survey discusses challenges such as environments with varying appearance and the dynamics of moving elements, the discussion is limited to **AI**-related work. The work only briefly analyzes localization and mapping tasks, while also focusing on reasoning, human-robot interaction, and planning. As for Zaffar *et al.* [112], the study focuses solely on reviewing, discussing, and comparing sensors with respect to sensor lifetime, field operability, ease of replacement, and suitability for various environments.

Nevertheless, none of the studies presented in Table 2.1 describe their methodology for selecting the works for discussion. This limitation prevents other researchers from repeating the review process, *e.g.*, to update existing studies and maintain an up-to-date understanding of the current state of the art in **SLAM**.

As a result, *A systematic literature review on long-term localization and mapping for mobile robots* [65] (see Appendix A) was written in the scope of this PhD Thesis. The review follows the **PRISMA** [98] guidelines, defining a systematic methodology to ensure the repeatability of the selection and data extraction processes while allowing future updates over the discussion and the review's methodology. Furthermore, our systematic review does not focus on any specific strategy or publication time interval. These considerations improve the study's coverage of the long-term localization and mapping topic, compared to the existent **SLAM** surveys.

In summary, the following questions are the focus of our systematic review on long-term localization and mapping literature [65]:

- main challenges inherent to lifelong **SLAM**;
- main strategies for accomplishing long-term operations with mobile robots;
- public datasets commonly used for evaluating long-term localization and mapping algorithms;
- how the researchers evaluate the performance of autonomous systems in long-term operations.

2.2.2 Review methodology

A systematic literature review uses explicit, rigorous, and reproducible systematic methods to synthesize the findings of studies related to a particular research question, topic area, or phenomenon of interest. This type of review ensures the quality and trustworthiness of its findings by presenting a complete, organized, and summarized analysis of all works considered, while allowing others to replicate or update the review. The most common standard for performing a systematic review is the **PRISMA** statement [98]. Although **PRISMA** was initially designed to evaluate the effects of health interventions, its methodology checklist items are general

Table 2.2: Exclusion (E) criteria for the selection process of the systematic literature review [65].

E#	Criteria	Statement
E1	Index	Papers not indexed in a scientific publication venue
E2	Language	Full-text of the papers not published in English
E3	Subject Area	Papers not classified in the databases as Computer Science, Engineering, Mathematics, or Multidisciplinary
E4	Short Papers	Papers classified as short papers accordingly to the publication venue
E5	Gray, Secondary, and Tertiary Literature	Books, preprints, reports, reviews, thesis, ...
E6	Availability	Full-text of the papers not available in digital libraries
E7	Dataset	Papers that focus only on data collection
E8	Coverage	Papers using only odometry for localization
E9	Scope	Papers that focus on different and not related subjects

and applicable to other subject areas [117–120]. Thus, the methodology used in our systematic review on long-term localization and mapping, extensively discussed in Appendix A, follows the PRISMA guidelines. This methodology is composed of the following steps:

1. definition of eligibility criteria;
2. conceptualization of the search strategy;
3. selection process:
 - (a) identification;
 - (b) screening;
 - (c) quality assessment.
4. data extraction;
5. synthesis and analysis of the included works.

First, we excluded eligible studies according to the exclusion criteria presented in Table 2.2. These criteria focus on the type of paper, availability, and relevance to the research topic of long-term localization and mapping. Studies must be published in scientific venues to ensure peer review. Only primary literature and English-language papers available in digital libraries are considered to ensure accessibility and reproducibility. Additionally, studies not categorized in the scientific databases under *Computer Science*, *Engineering*, *Mathematics*, or *Multidisciplinary* subject areas are excluded, as these fields are closely related to the research topic. We also excluded works that focus solely on datasets, as these are not considered original research, as well as odometry-only approaches, which are not viable for long-term operations. Lastly, we excluded papers that are not directly related to long-term localization and mapping.

Next, our search strategy for the systematic literature review aimed to identify relevant data sources and define the search fields and base string to consider during the process of obtaining the included records. The review considered multiple databases to increase the overall scientific coverage, namely, ACM [121], Dimensions [122], IEEE Xplore [123], INSPEC [124], Scopus [125], and Web of Science [126]. Furthermore, the *Title*, *Abstract*, and *Keywords* were the fields used to obtain search results and focus the search on the most relevant summary information. The *Keywords* field includes the author keywords, database-indexed terms, and, if available, uncontrolled keywords. Also, we did not impose any date restrictions to avoid missing important studies and to provide a more comprehensive discussion on the topic. Finally, we last performed the search on May 17, 2022, for the review. The search query used the following base string, combining terms related to




Figure 2.5: Preferred Reporting Items for Systematic reviews and Meta-Analyses (PRISMA) flow diagram for the selection process [65].

robots, vehicles, localization, mapping, **SLAM**, and long-term operations, leveraging the asterisk (*) wildcard to consider multiple synonyms.

(robot* OR vehicle*) AND ((locali* AND map*) OR "slam") AND
("long term" OR "life long" OR lifelong)

Figure 2.5 summarizes the selection process for the review. During the identification phase, the search strategy was applied to all data sources, using Parsifal [127] to identify duplicate records across databases. The 232 excluded records in this phase include works that do not meet criteria E4 (auxiliary Python scripts to identify short papers from BibTex bibliography data) and E7 (searching for the dataset term in the titles). We have gone through a double-check procedure to ensure these criteria are applied correctly. As a result, we screened 589 records, requiring us to read their titles and abstracts to determine eligibility for the quality assessment phase. Initially rejected papers were subject to a further assessment to validate the exclusion. This assessment included analyzing the experimental results and conclusions presented in the original works to either confirm or reverse the exclusion decision. Overall, we rejected 178 studies in the screening phase, leaving 411 eligible works for quality assessment.

The quality evaluation of the 411 eligible works followed the 8 Quality Evaluation (QE) criteria presented in Table 2.3. All of them are subjective criteria derived from the analysis of the eligible works, focusing on the detail and relevance of the proposed methodologies, experimental results, and long-term consideration. Considering the possible scores for the QE criteria in Table 2.3, each work can only have a maximum score of 10.0. After we evaluated the 411 eligible works, our first observation was that works with a non-detailed or appropriate methodology, results' discussion, or conclusions should not be included in the review. Another observation is relative to rejecting works that do not consider either localization or mapping problems, or do

Table 2.3: Quality Evaluation (QE) criteria and score range for the selection process of the systematic review [65].

QE#	Criteria	Score
QE1	Does the paper have an updated state of the art on long-term localization and mapping?	{0.0, 0.5, 1.0}
QE2	Is the methodology appropriate and detailed?	{0.0, 0.5, 1.0}
QE3	Does the methodology consider both localization and mapping problems?	{0.0, 1.0, 2.0}
QE4	Is the hardware and/or software used in the experiments detailed?	{0.0, 0.5, 1.0}
QE5	Does the paper presents any kind of long-term experimental results?	{0.0, 2.0}
QE6	Does the paper presents comparative results with other methods and/or ground-truth data?	{0.0, 1.0}
QE7	Does the work's implementation and/or the data used in the experiments are publicly available?	{0.0, 1.0}
QE8	Is the discussion of the results and conclusions appropriate and detailed?	{0.0, 0.5, 1.0}

not present any long-term experimental results. This rejection is due to the focus of the review on long-term localization and mapping. Furthermore, the quality assessment phase should include a cut-off score to filter out works with low quality scores. Therefore, we considered the following two reasons to reject a record in the quality assessment phase of the selection process, resulting in 175 studies for the data extraction phase:

1. QE2, QE3, QE5, QE8: reject works with a 0.0 (no compliance) score;
2. cut-off score: reject works with a score lower or equal to 7.5/10.0.

Finally, the data extraction process involved analyzing the eligible records from the quality assessment phase and extracting relevant information. In the scope of the systematic review on long-term localization and mapping, we considered the following **Data Extraction (DE)** items required for each record:

- [DE1] Long-term considerations – long-term factors the works consider in their proposed approach and experiments. Considering the knowledge obtained in the previous phases of the review's methodology, we considered the following factors for categorizing the included works:
 - appearance: varying conditions, appearance changes;
 - dynamics: environment dynamics, dynamic elements;
 - sparsity: map pruning, redundant data removal;
 - multi-session: map management;
 - computational: memory management, efficiency.
- [DE2] Localization – how the robot localizes itself and the type of localizer;
- [DE3] Mapping – type of the map;
- [DE4] Multi-robot – if the proposed methodologies consider multi-robot systems;
- [DE5] Execution mode – offline, online, if requires both, or if no information on this item;
- [DE6] Environment and domain – type of environment (indoor, outdoor) and domains (air, ground, water) tested with the proposed methodologies;
- [DE7] Sensory setup – which sensors considered in the methodologies;
- [DE8] Non-public experiments – if the authors performed experiments or tests with non-public data;
- [DE9] Ground-truth – how ground-truth for non-public data is obtained or its type, if available;
- [DE10] Distance and time characteristics – relative to the non-public experiments if available, as follows:
 - total distance (km) of the non-public experiments;
 - path (km), in the case of repetitive paths;




Figure 2.6: Evolution of published records per year considering the 142 included records in the systematic literature review on long-term localization and mapping. The time interval is between the smallest publication year found in the included records (2002) and the year of last full inquiry’s date (2022). The latter is with a diagonal hatch pattern due to the fact that the last full inquiry did not consider the whole year [65].

- total time (h) in terms of continuous operation;
- time interval (day/week/month/year, or d/w/m/y) between the first and the last run.
- [DE11] Datasets – if and which public datasets are used in the experiments;
- [DE12] Evaluation metrics – which metrics are used for evaluation.

Although the data extraction phase typically does not exclude records, 33 of the 175 works analyzed had extended versions of the proposed methodologies, which we excluded to maintain the originality of the review. As a result, we included 142 original works for analysis and discussion in the systematic literature review study, representing 34.55% of the eligible records. The extracted data for these works is available in Appendix A.

2.2.3 Results overview

The publications per year evolution can be used to evaluate the relevance of the long-term localization and mapping topic. Figure 2.6 presents this evolution on the 142 included records in our systematic review. The final year, 2022, has its data with a diagonal hatch pattern, indicating that it does not represent a full year. Analyzing Figure 2.6, the long-term localization and mapping topic started to gained relevance in 2009, with 6 works, compared to only 2 in the previous years. From that year onwards, the graph shows an almost linear trend, reaching a maximum of 23 records in 2021, and already having 8 publications by May 17, 2022. This tendency indicates that long-term localization and mapping have gained interest over the years and, consequently, underscores the importance and relevance of the research topic to the scientific community.

Although the scope of our systematic literature review [65] includes methodologies for handling the dynamics of the scene, the original query did not focus on any particular challenge related to long-term localization




Figure 2.7: Evolution of published records per year in Scopus [125] and Web of Science [126] databases on long-term localization and mapping: (a) (robot* OR vehicle*) AND ((locali* AND map*) OR "slam") AND ("long term" OR "life long" OR lifelong); (b) (robot* OR vehicle*) AND ((locali* AND map*) OR "slam") AND ("long term" OR "life long" OR lifelong) AND (dynamic* OR mov*). Date of the last full inquiry: February 13, 2026. The time interval is between 2009 and the year of last full inquiry's date (2026). The latter is with a diagonal hatch pattern due to the fact that the last full inquiry did not consider the whole year.

and mapping. Figure 2.7 presents the evolution of publication year during the identification phase of a systematic review, using the *Title*, *Abstract*, and *Keywords* (authors' and indexed ones) as search fields, across two bibliographic databases: Scopus [125] and Web of Science [126]. The consideration of the previous two with INSPEC [124] could improve the coverage of the search results, as discussed in Appendix A, where the joint consideration of the three databases guarantees the identification of all 142 included records. However, the b-on² agreement provided by the Portuguese funding agency, FCT – Fundação para a Ciência e a Tecnologia, no longer has INSPEC [124] available to its participating entities, including FEUP and INESC TEC. Moreover, Figure 2.7a uses the exact inquiry string used in our systematic review. Figure 2.7b uses a different string to present the evolution of publications on the challenge of localization and mapping in dynamic environments, which aligns with the goals of this PhD Thesis. This scope reduction is accomplished by using the following search string:

(robot* OR vehicle*) AND ((locali* AND map*) OR "slam") AND
 ("long term" OR "life long" OR lifelong) AND (dynamic* OR mov*)

The part added to the original query (dynamic* OR mov*) aims to limit search results to works featuring dynamic or moving elements in the scene. Also, only works since 2009 were considered in the analysis, given that year was the one where the long-term localization and mapping started to gain traction in the literature.

When analyzing Figure 2.7a, the graph shows increasing interest among the scientific community in long-term localization and mapping. Since 2009, the graph shows a consistent upward trend, reaching a maximum of 113 and 81 published records in 2025 for Scopus [125] and Web of Science [126], respectively. The trend indicates that long-term SLAM is gaining prominence, underscoring its importance to the scientific community.

²<https://www.b-on.pt/>

This result conforms with the analysis of the year of publication for the 142 literature works included in our systematic review (see Figure 2.6).

Similarly, the graph shown in Figure 2.7b indicates growing interest in the topic of this PhD Thesis: localization and mapping in dynamic environments. From 2009 onwards, the topic has also gained relevance, achieving a maximum of 53 and 35 published works in 2025 for Scopus [125] and Web of Science [126], respectively. Although the overall number of publications related to dynamic environments is lower than for the broader long-term topic as expected, the growing interest in addressing the challenges of localization and mapping in dynamic environments underscores the relevance of this PhD Thesis.

2.3 Long-Term SLAM in Dynamic Environments

In order to focus the literature discussion in this chapter on localization and mapping in dynamic environments, this section follows the same systematic methodology outlined in Section 2.2.2. The discussion builds on the analysis already made in the review article [65] on dynamics modelling. This analysis examines how literature handles scene dynamics through map representation, map matching, dynamic object detection, and predictive modelling. However, similarly to Section 2.2.3, the original search query is extended with `dynamic*` OR `mov*` to update the included works from the original date of the last full search inquiry in the scientific databases, May 17, 2022, to February 13, 2026, while focusing on works that handle scene dynamics. Also, this last full inquiry includes only Scopus [125] and Web of Science [126] search results to facilitate the update to the discussion. As a result, after completing all stages of the review methodology described in Appendix A and briefly presented in Section 2.2.2, this section discusses 53 works from the literature, including the 32 studies related to handling scene dynamics already discussed in the review article.

Table 2.4 presents the data extraction results for the 53 works identified in the latest full inquiry on dynamics modelling. The **Dynamic Data Extraction (DDE)** items differ from those presented in 2.2.2, as they are designed to capture meaningful information within the scope of localization and mapping in dynamic environments. Accordingly, the **DDE** items required for each of the 53 records are as follows:

- [DDE1] Mapping – mapping process and type of the map;
- [DDE2] Localization – how the robot localizes itself and the type of localizer;
- [DDE3] Environment and domain – type of environment (indoor, outdoor) and domains (air, ground, water) tested with the proposed methodologies;
- [DDE4] Sensory setup – which sensors considered in the methodologies;
- [DDE5] Dynamics detection – how dynamics identified from sensor observations;
- [DDE6] Dynamics labelling – classification of dynamic observations;
- [DDE7] Dynamics statistics – statistics formulated in the methodologies related to dynamics modelling;
- [DDE8] Dynamics prediction – methodologies to predict scene dynamics in the environment;
- [DDE9] Dynamics removal – how dynamics are handled in localization and mapping processes.

Next, complementing the information available in Table 2.4, the following analysis discusses how the 53 works included in this section address scene dynamics through map representation, map matching, dynamic object detection, and prediction modelling.

Table 2.4: Data extraction items retrieved from literature works related to long-term localization and mapping in dynamic environments.

DDE:	1:	2:	3:	4:	5-9: Dynamics					
Ref.	Mapping	Localization	Indoor Outdoor	Air Ground Water	Sensors	Detection	Labelling	Statistics	Prediction	Removal
Biber and Duckett (2009) [4]	submap (2D point cloud, graph, 3DoF edges)	point cloud alignment (2D, 3DoF)	x	x	wheel odometry, laser (2D)	ray tracing, 1D robust statistics over dynamic sample sets (median, median of absolute differences), timescales	–	timescale (update policy)	–	dynamic sample set (timescale-based probability)
Dayoub <i>et al.</i> (2011) [6]	keyframe (graph, 6DoF edges)	feature matching (location)	x	x	camera (omni)	scan-to-map matching	sensory + short + long-term memory (SM, STM, LTM)	hit + miss observations	–	recall mechanism (forget unused features from LTM)
Walcott-Bryant <i>et al.</i> (2012) [10]	pose graph (graph, 2D point clouds, 3DoF edges; iSAM)	point cloud alignment (2D, 3DoF)	x	x	laser (2D)	scan-to-map matching	active + dynamic map, static + removed + added points	change score	–	disable node sectors, inactive nodes removal (all sectors turned off)
Bacca <i>et al.</i> (2013) [7]	keyframe (graph; TORO)	Bayesian (2D, 3DoF)	x	x	camera (omni), laser (2D)	k-means clustering	sensory + short + long-term memory (SM, STM, LTM)	feature stability histogram	–	STM pruning
Einhorn and Gross (2013) [128]	pose graph (graph, 2D/3D NDT; g^2o)	odometry	x	x	camera (mono, RGBD), laser (2D)	explicit free space mapping	–	–	–	explicit free space mapping
Tipaldi <i>et al.</i> (2013) [48]	grid (occupancy, 2D)	particle filter (2D, 3DoF)	x	x	laser (2D)	hidden Markov model (HMM)	–	mixing time of Markov chain	dynamic cells transition	–
Saarinne <i>et al.</i> (2013) [21]	3D NDT, grid (occupancy, 3D)	–	x	x	camera (RGBD), laser (3D)	ray tracing, scan-to-map matching (NDT-based)	–	–	–	ad-hoc sensor model
Pomerleau <i>et al.</i> (2014) [37]	point cloud (laser, 3D)	point cloud alignment (3D, 3DoF)	x	x	wheel odometry, laser (3D)	ray tracing approximation (sparse map cloud, kd-tree w/ spherical coordinates)	–	Bayesian (visibility, points proximity, motion patterns)	–	down-weight dynamic points
Einhorn and Gross (2015) [40]	pose graph (graph, 2D/3D NDT)	odometry	x	x	camera (mono, RGBD), laser (2D)	occupancy probability (short vs static / long-term)	dynamic cell clusters	static occupancy probability	–	–
Rapp <i>et al.</i> (2015) [49]	grid (occupancy, 2D)	particle filter (2D, 3DoF)	x	x	odometry, radar	semi-Markov chain extension	–	cell reliability (Lévy process)	expected retention time on cell state	–
Santos <i>et al.</i> (2016) [51]	spatio-temporal grid (occupancy, 3D)	–	x	x	camera (RGBD)	fast Fourier transform (FFT, FreMEEn) (static vs periodic changes)	–	mean value, dominant frequencies	fast Fourier transform (FFT, FreMEEn) (periodic changes)	–
An <i>et al.</i> (2016) [129]	pose graph (graph)	EKF (2D, 3DoF)	x	x	wheel odometry, camera (gray, mono), laser (2D)	feature matching	–	edge weight (dynamic edge link)	–	nodes removal (average edge weight threshold)
Biswas and Veloso (2017) [9]	feature (2D line segments), varying graphical network (VGN, Ceres)	maximum likelihood estimation (MLE)	x x	x	wheel odometry, laser (2D), camera (RGBD)	feature ray tracing	dynamic (only for obstacle avoidance) + short (movable) + long-term features (DF, STF, LTF)	–	–	dynamic features removal (localization)
Krajník <i>et al.</i> (2017) [52]	grid (occupancy, 2D/3D)	–	x x	x	camera (RGBD), laser (2D)	fast Fourier transform (FFT, FreMEEn) (static vs periodic changes)	–	mean value, dominant frequencies	fast Fourier transform (FFT, FreMEEn) (periodic changes)	–
Bescos <i>et al.</i> (2018) [18]	keyframe (graph)	bundle adjustment (3D, 6DoF)	x x	x	camera (color, mono, stereo, RGBD)	pixel segmentation (Mask R-CNN), multi-view geometry changes (depth)	dynamic or movable classes (person, bicycle, car, motorcycle, airplane, bus, train, truck, boat, bird, cat, dog, horse, sheep, cow, elephant, bear, zebra and giraffe)	–	–	background inpainting
Sun <i>et al.</i> (2018) [130]	grid (occupancy, 3D)	point cloud alignment (3D, 6DoF)	x	x	laser (3D)	recurrent neural network (LSTM RNN) for cell semantic state transition	–	–	–	–

Table 2.4 (Continued)

DDE:	1:	2:	3:		4:	5–9: Dynamics							
Ref.	Mapping	Localization	Indoor	Outdoor	Air	Ground	Water	Sensors	Detection	Labelling	Statistics	Prediction	Removal
Lázaro <i>et al.</i> (2018) [22]	submap (2D point cloud, graph, 3DoF edges; g^2o)	point cloud alignment (2D, 3DoF)	x			x		wheel odometry, laser (2D)	ray tracing, scan-to-scan matching	–	–	–	cloud merging, outdated nodes removal
Boniardi <i>et al.</i> (2019) [23]	pose graph (graph, 2D point clouds, 3DoF edges), CAD floor plans	point cloud alignment (2D, 3DoF)	x			x		laser (2D)	scan-to-map matching	–	–	–	outdated nodes removal
Wang <i>et al.</i> (2019) [19]	keyframe (graph)	bundle adjustment (3D, 6DoF)	x	x		x		camera (RGBD)	semantic segmentation (FCIS), feature motion consistency check	static background, moveable class (person, vehicles, cup, chair, etc.)	–	–	moveable points removal
Zhang <i>et al.</i> (2019) [11]	grid (signed distance field, 2D)	point cloud alignment (2D, 3DoF)	x			x		laser (2D)	ray tracing (Bresenham algorithm), scan-to-map matching	static + semi-static + dynamic objects	–	–	dynamic objects removal
Ganti and Waslander (2019) [12]	keyframe (graph)	bundle adjustment (3D, 6DoF)		x		x		camera (stereo)	semantic segmentation (Bayesian SegNet)	static (road, sidewalk, building, wall/fence, pole, traffic light, traffic sign, vegetation, terrain) + dynamic (sky, person/rider, car, truck/bus, motorcycle/bicycle, void) classes	–	–	–
Song <i>et al.</i> (2019) [45]	keyframe (graph)	odometry (3D, 6DoF)		x		x		camera, IMU	pixel-wise motion attribute segmentation (unstable vs moving vs static, MD-Net)	unstable (sky, vegetation, terrain) + moving (human, vehicle, traffic light) + static (ground, flat, building, wall, fence, guard rail, bridge, tunnel, pole, pole group, traffic sign)	–	–	–
Pan <i>et al.</i> (2019) [41]	feature (point clusters)	odometry, reprojection minimization (3D, 6DoF)		x		x		camera (mono), laser (3D)	point clusters registration	–	clusters occurrence	–	dynamic point clusters removal
Wang <i>et al.</i> (2020) [50]	grid (occupancy, 2D)	particle filter (2D, 3DoF)	x			x		wheel odometry, laser (2D)	time series modelling	–	autoregressive moving average (ARMA)	autoregressive moving average (ARMA; static + periodic + aperiodic changes)	–
Yang <i>et al.</i> (2020) [42]	keyframe (graph)	visual odometry (3D, 6DoF)	x			x		camera (RGBD)	object detection (YOLOv3), geometry constraints	dynamic class (people)	–	–	–
Ding <i>et al.</i> (2020) [38]	point cloud (3D)	bundle adjustment (3D, 6DoF)		x		x		camera (color, stereo), laser (3D)	multi-session consistency, point clustering	–	observations count	–	intra (RANSAC, ray tracing) + inter-session (map matching) dynamic points removal
Yue <i>et al.</i> (2020) [43]	point cloud (3D), grid (occupancy, 3D)	–		x		x		camera (color, mono, thermal), laser (3D)	object detection (YOLOv3), 3D point cloud projection	dynamic class (people)	–	–	–
Singh <i>et al.</i> (2021) [46]	topological (BoW)	feature matching (location)	x			x		camera (stereo, RGBD)	pixel segmentation	dynamic class (sky, person, car, etc.)	–	–	down-weight dynamic class
Zhu <i>et al.</i> (2021) [13]	grid (occupancy, 2D)	particle filter (2D, 3DoF)	x			x		wheel odometry, camera, laser (3D)	object detection (YOLOv3)	static + semi-dynamic (chairs, parked cars) + dynamic (moving people / cars) objects, static + semi-dynamic maps	–	–	down-weight semi-dynamic objects removal
Guo and Liu (2021) [14]	keyframe (graph)	feature matching (3D, 6DoF)	x			x		camera (RGBD)	pixel segmentation (Mask R-CNN, depth region growing refinement), multi-view geometry segmentation (depth)	dynamic vs static / ground class	–	–	background inpainting

Table 2.4 (Continued)

DDE: 1: 2: 3: 4: 5-9: Dynamics										
Ref.	Mapping	Localization	Indoor Outdoor	Air Ground Water	Sensors	Detection	Labelling	Statistics	Prediction	Removal
Thomas <i>et al.</i> (2022) [39]	sparse grid (hashmap, occupancy, 2D)	point cloud alignment (3D, 6DoF)	x	x	wheel odometry, laser (3D)	ray tracing (NN training), feature semantic labelling (KPConv)	feature semantic labelling (ground, permanent, movable, dynamic)	–	dynamic motion label (KPConv)	movable (localization) + dynamic (localization, mapping) points removal
Du <i>et al.</i> (2022) [32]	keyframe (graph)	reprojection minimization (3D, 6DoF)	x	x	camera (RGBD)	long-term consistent conditional random field (LC-CRF)	–	observations count, average reprojection error, symmetric epipolar distance	–	dynamic features removal
Xing <i>et al.</i> (2022) [8]	keyframe (graph)	feature matching (3D, 6DoF)	x x	x x	camera (gray, stereo), IMU	object detection (MobileNet V2 SSD), feature motion consistency check	short + long-term dynamics (STD, LTD)	–	–	–
Chen <i>et al.</i> (2022) [44]	point cloud (3D)	odometry (3D, 6DoF)	x	x	laser (3D)	object detection (PointPillars)	dynamic points (vehicles, pedestrians, cyclists, etc.)	–	–	dynamic objects removal
Kim and Kim (2022) [24]	keyframe (graph; iSAM2, GTSAM)	point cloud alignment (3D, 6DoF)	x	x	laser (3D), IMU	scan-to-map matching, multi-session consistency	high + low dynamic (HD, LD), positive (new points) + negative (disappeared) difference (PD, ND)	–	–	positive + negative difference maps (PD, ND)
Wu <i>et al.</i> (2022) [30]	keyframe (graph)	feature matching (3D, 6DoF)	x	x	camera (RGBD)	reprojection error (Student's t-distribution modelling), conditional random field (CRF) minimization	–	static score (Student's t-distribution)	–	dynamic features removal
Fu <i>et al.</i> (2022) [131]	feature (plane, SDF)	–	x	x	camera (RGBD)	planes 2D height map comparison, planes SDF 3D validation	2D ternary change mask (changed, unchanged, unexplored)	–	–	–
Liu <i>et al.</i> (2022) [31]	keyframe (graph)	bundle adjustment (3D, 6DoF)	x	x	camera (RGBD)	motion consistency, reprojection error	–	static probability (reprojection error)	–	dynamic features removal
Deng <i>et al.</i> (2023) [15]	keyframe (graph)	bundle adjustment (3D, 6DoF)	x	x	camera (stereo, RGBD)	scan-to-map matching	dynamic + semi-static (periodic, abrupt changes) + static points	observations count, survival time, point persistence	Bayesian persistence filter (persistence probability), time series modeling (semi-static points)	dynamic points removal
Nielsen and Hendebay (2023) [132]	feature (multi-hypotheses tree)	–	x	x	wheel odometry, laser (2D)	multi-hypotheses testing	–	negative information (missed detections) + new landmark observations + false measurements + landmark entering FOV	–	hypotheses pruning / merging
Soares <i>et al.</i> (2023) [20]	keyframe (graph)	bundle adjustment (3D, 6DoF)	x	x	camera (RGBD)	object detection (YOLOv4), feature repopulation (back vs foreground)	dynamic (person) + movable (chair) + static objects	–	–	covariance static < movable < dynamic measurements
Stefanini <i>et al.</i> (2023) [25]	grid (occupancy, 2D)	–	x	x	wheel odometry, laser (2D)	scan-to-map matching, ray tracing	–	cells changed flag count	–	–
Yu <i>et al.</i> (2023) [47]	keyframe (graph)	bundle adjustment (3D, 6DoF)	x	x	camera (RGBD)	semantic segmentation (lightweight Mask R-CNN + MobileNetV2), motion probability model	dynamic class (humans, animals, cars)	motion probability model (semantic mask + optical flow)	–	dynamic features removal
Hroob <i>et al.</i> (2024) [26]	grid (occupancy, 3D), point cloud (3D)	point cloud alignment (3D, 6DoF)	x	x	laser (3D), IMU	scan-to-map matching, long-term stability network (LTS-NET, PointNet++-based)	stable vs unstable points	point spatiotemporal stability score (LTS-NET)	point stability score (unsupervised learning, based on recorded sessions)	stability score threshold removal (localization)
Hroob <i>et al.</i> (2024) [27]	grid (occupancy, 3D), point cloud (3D)	point cloud alignment (3D, 6DoF)	x	x	laser (3D)	scan-to-map matching, stability score inference (modified MinkUNet14), scan-map discrepancy segmentation	stable vs unstable points	point spatiotemporal stability score (modified MinkUNet14)	point stability score (unsupervised learning, based on recorded sessions)	stability score threshold removal (localization)

Table 2.4 (Continued)

DDE:	1:	2:	3:			4:	5-9: Dynamics						
Ref.	Mapping	Localization	Indoor	Outdoor	Air	Ground	Water	Sensors	Detection	Labelling	Statistics	Prediction	Removal
Wang <i>et al.</i> (2024) [16]	grid (occupancy, 2D)	particle filter (2D, 3DoF)	x		x			wheel odometry, laser (2D)	object clustering (k-means), time series modelling	semi-static vs dynamic map	autoregressive moving average (ARMA), dynamic object density	autoregressive moving average (ARMA; static + periodic + aperiodic changes)	–
Yang <i>et al.</i> (2024) [29]	point cloud (3D)	–		x	x			laser (3D)	multi-scan consistency (OctoMap), multi-session consistency (k-NN spatial filter)	–	–	–	positive + negative difference maps (PD, ND)
Cheng <i>et al.</i> (2025) [34]	3D Gaussian splatting	gradient-based optimization (rendered RGB + depth + semantic maps)		x		x		camera (RGBD), laser (3D)	object segmentation (SegmentAnything), motion + temporal consistency	entering + vanishing + full observation	–	–	3D Gaussian-based encoding (dynamic objects)
Hu <i>et al.</i> (2025) [35]	object map (3D point cloud, color histogram, semantic category)	feature matching (3D, 6DoF)	x		x			camera (RGBD)	object detection (YOLOX), motion consistency, foreground vs background clustering (DB-SCAN)	–	object static probability	–	dynamic points weighting (pose optimization)
Jia <i>et al.</i> (2025) [36]	3D Gaussian splatting	feature matching (3D, 6DoF)	x		x			camera (RGBD)	semantic segmentation (FastInst), motion consistency	–	–	–	dynamic objects removal (3DGS model)
Li <i>et al.</i> (2025) [17]	grid (hashmap, dynamic i-Octree)	point cloud alignment (3D, 6DoF)		x		x		laser (3D), IMU	scan-to-map matching	prior (multi-session) + temporary (short-term) + static map points	–	–	–
Wang <i>et al.</i> (2025) [33]	feature (point, line, plane)	bundle adjustment (3D, 6DoF)	x		x			camera (RGBD), IMU	semantic segmentation (DeepLab v3), motion consistency (optical flow), structural regularity check (Manhattan world)	–	dynamic probability (dynamic filter)	–	dynamic features removal (dynamic probability threshold)
Zhu <i>et al.</i> (2025) [28]	sparse grid (hashmap, 2D+3D, i-Octree)	odometry (3D, 6DoF)		x		x		laser (3D), IMU	scan-to-map matching, object clustering (region growing)	–	–	–	dynamic objects removal (scan filtering), ghost trails retroactive removal (i-Octree search)

2.3.1 Map representation

Inspired by the human memory, Dayoub *et al.* [6] and Bacca *et al.* [7] adapt the multi-store model of Atkinson and Shiffrin [133] for mapping. This model divides memory into three stores: **Sensory Memory (SM)**, which stores the most recent information from the sensors; **Short-Term Memory (STM)**; and **Long-Term Memory (LTM)**. Three mechanisms move information between memories: selective attention from **SM** to **STM**, rehearsal from **STM** to **LTM**, and the retrieval mechanism to move unused information from **LTM** back to **STM**. Dayoub *et al.* [6] implement two types of state machines for rehearsal and retrieval mechanisms of the **STM** and **LTM**. In rehearsal, an **STM** feature moves closer to **LTM** or moves back to the initial state (or is forgotten if already in that state), depending on whether it is observed consecutively or not. Similar for retrieval, where a feature in **LTM** moves to the initial state or closer to being forgotten if it is not observed in the current view. Consequently, **LTM** and **STM** retain the most static and dynamic features, respectively, based on their observability in the current view. In a changing environment, the method reduced localization failure rates compared to a static view. Instead of using a state machine, Bacca *et al.* [7] implement a **Feature Stability Histogram (FSH)** depending on the feature observability to distinguish between **STM** and **LTM** features using k-means clustering. This modification allows an input feature in **SM** to bypass **STM** and become part of **LTM**, depending on the feature strength. The method was able to filter out pedestrians from 2D laser and camera data and achieved a more accurate representation of the environment compared to a static approach.

Although Biber and Duckett [4] do not adopt specific memory mechanisms, they implement **STM** and **LTM** maps. The method implements a dynamic map as a set of local maps, each containing submaps that represent different timescales. The timescale parameter for each submap determines when to add samples from 2D laser scans probabilistically. The dynamics of the environment are represented by using 5 different timescales, where the smallest one (~ 3.1 s) represents an **STM** map updated at every instant, and the other 4 timescales (~ 0.43 operation runs, 0.43 days, 3.1 days, and 13.5 days) are **LTM** submaps updated after each session or daily. Instead of focusing only on the **LTM** maps as in Dayoub *et al.* [6] and Bacca *et al.* [7], Biber and Duckett [4] select the best representation of both **STM** and **LTM** maps that best explains the sensor data. In a 5-week experiment, localization improved with a dynamic map, in contrast to the degradation observed when only considering a static representation of the environment. Indeed, the usage of timescales led to static parts emerging as walls in **LTM** and dynamic elements disappearing from **STM** maps.

Similar to the works [6] and [7], two maps can provide a more stable and dynamic representation of the environment. In Walcott-Bryant *et al.* [10], an active map represents the most current state of the environment, including parts that did not change from previous passes and objects added to the environment. A dynamic map only saves the points of a 2D laser scan that changed over time. Wang *et al.* [19] use a tracking map with short-term static points and a long-term map containing only long-term static points, identified by a semantic segmentation module, with ORB-SLAM2 [134]. Also, Zhu *et al.* [13] create offline a semi-dynamic map and a static one, where the former has semi-dynamic objects (parked cars in a parking lot environment), and the latter has both static and semi-dynamic objects. The main goal of representing two different dynamics is to favor the more stable one in the long term. Walcott-Bryant *et al.* [10] and Wang *et al.* [19] use only static parts of the current environment representation (active and long-term maps, respectively) for localization. In the experiments, the approach [10] showed that it could identify static parts, though it was affected by false positives and negatives and by the blur effect in the 2D grid map. The method [19] improved the **Absolute Trajectory Error (ATE)** in a dynamic environment over ORB-SLAM2 [134] and DynaSLAM [18]. As for Zhu *et al.* [13],

their **Monte Carlo Localization (MCL)** framework (equivalent to particle-filter-based localization) reduces the weight assigned to observations of moved semi-dynamic objects. The method improved the localization in a parking lot compared to a standard **MCL**.

As for a multi-hypothesis representation, Nielsen and Hendeby [132] propose a framework to track landmark positions using Markov chains to model transitions between positions. The method [132] groups position hypotheses by landmark signature, and selects, according to incoming observations, the position within each group with the highest probability as the active one. The hypothesis likelihood depends on the match between the landmark measurements and the landmark in the map, as well as on the number of missed detections, *i.e.*, landmarks with a probability of being seen in the current pose but not represented in the observations. Although the work [132] maintained an up-to-date representation of the actual landmark positions in simulated and real experiments, the proposed approach always assume uniquely identifiable landmarks.

Li *et al.* [17] leverage a map representation based on voxelized maps and octree structures to handle environment dynamics. The work [17] organizes prior, temporary, and static map points within hashed voxel blocks, each containing a Dynamic i-Octree [135] for efficient local map loading, radius-neighbor search, and incremental map updates. The octree representation is not only a storage/indexing structure but also a mechanism that allows the system to adapt when the robot moves between mapped and unmapped areas or when the environment changes. This mechanism promotes temporary observations into the static map only when sufficient evidence accumulates. Experiments demonstrated the effectiveness of the method, in which a prior map was dynamically updated to reflect changes caused by parked vehicles and growing trees, while minimizing the inclusion of dynamic objects, such as passing vehicles and pedestrians.

2.3.2 Map matching

Another trend in the literature is the identification of the scene dynamics by comparing the current observation to the map. Assuming a prior vector map as a permanent map, Biswas and Veloso [9] determine the probability of long-term observed features by the 2D laser scan-to-map matching distance. Short-term features are identified by the scan-to-scan matching distance, while the remaining features are considered dynamic and not used for localization. Compared to **MCL** with a static map and to Tipaldi *et al.* [48], Biswas and Veloso [9] obtained lower localization error in a parking-lot environment. However, the method would not handle semi-static changes. Instead of using a permanent map, Zhang *et al.* [11] maintain an **Signed Distance Field (SDF)** representation based on a prior occupancy map. The proposed approach rejects dynamic points identified by range flow and updates the **SDF**-based map with semi-static changes observed in the scan-to-map difference. Compared to **MCL** in a semi-static environment, Zhang *et al.* [11] had lower pose errors and an improved representation of the environment.

Similarly, Fu *et al.* [131] use **SDFs** in mapping, proposing the PlaneSDF representation, in which a scene is a set of planes, each with an associated **SDF** volume. This volume describes geometric details of the objects attached to their associated plane. Then, a change detection algorithm compares 2D height maps and performs 3D geometric validation between two PlaneSDF-based maps to compute a change mask. However, the method [131] failed to detect changes in ambiguous situations, such as object occlusions. Boniardi *et al.* [23] detect semi-static changes by leveraging **ICP** scan-to-map consistency and a **Computer-Aided Design (CAD)** model of the environment, and update the map accordingly. The proposed approach [23] was capable of maintaining a consistent map even with substantial environment reconfiguration. Du *et al.* [32] minimize the

Gibbs energy defined on the proposed **Long-term Consistent Conditional Random Field (LC-CRF)** to detect dynamic points, assuming that these points often exhibit large reprojection errors in frame-to-map matching and that points tend to share dynamic properties with their neighbors. In a dynamic scene, **LC-CRF** [32] achieved lower **ATE** than ORB-SLAM [136]. Unlike **LC-CRF** [32], the **Dynamic visual Feature Removal SLAM (DFR-SLAM)** proposed by Liu *et al.* [31] does not look for similar dynamic properties in the features' neighbors. Instead, the method [31] assumes that static features exhibit consistent motion over a long time and that dynamic features exhibit distinct motion patterns, labeling them using the proposed motion-consensus filtering algorithm. **DFR-SLAM** [31] formulates a binary classification problem using a graph-cut optimization framework, assigning penalties to connected visual features across instances that do not share the same label. The method [31] had a lower **ATE** than ORB-SLAM2 [134] in dynamic scenarios.

Furthermore, Pan *et al.* [41] and Ding *et al.* [38] leverage clustering properties of the observations, evaluating the observation count when matching the scan with the map. The proposed approach segments the points of a 3D **LiDAR** point cloud into clusters, assuming that dynamic points do not appear frequently at the same location. The map only considers clusters that appear in the same location more than 10 times. As for Ding *et al.* [38], the method builds on the assumption that the dynamic and static parts of the environment exhibit clustering properties relative to their neighbors (similar to Du *et al.* [32]). The number of observations across multiple mapping sessions, combined with their consistency relative to neighbors, determines whether a map point is static throughout the sessions. Both representations of the environment in Pan *et al.* [41] and Ding *et al.* [38] were stable to structural changes in the environment.

Building on multi-session consistency, Kim and Kim [24] classify dynamic points into high- and low-dynamic categories. An initial processing step removes high-dynamic range points using a range-image-based discrepancy. This step is followed by a multi-session consistency check in which aligned sessions are compared via kd-tree searches to identify low-dynamic points. By parsing negative (removed data between consecutive sessions) and positive differences (new data appearing in the environment between sessions) instead of saving the whole mapping sessions, their representation achieves a 60% reduction in memory usage compared to maintaining full session histories. Similarly, Yang *et al.* [29] utilize a clean initial base map and identify changes in incoming sessions using a **k-Nearest Neighbor (NN)**-based spatial filter. The proposed approach segments points into base-difference, coexist, and session-difference categories, while implementing a map version-control system based on positive and negative differences to maintain the environment representation throughout time. To accelerate this process, Yang *et al.* [29] employ 2D **Bird-Eye-View (BEV)** descriptors to pinpoint intensity differences in 3D data. A similar memory reduction compared to Kim and Kim [24] was observed in the experiments, increasing memory efficiency by 40.1%–94.2% across datasets, while retaining at least 5–7% more static points compared to [24].

The concept of ray tracing is also used by the included works to handle dynamic changes. Lázaro *et al.* [22] use ray tracing to exploit free space information. When comparing two 2D point clouds from a viewpoint, ray tracing identifies new objects added to the scene (observed points closer to the viewpoint than the old ones) and outdated information (observed points farther away), allowing dynamic change detection and providing an up-to-date representation for localization. Also, Lázaro *et al.* [22] merge points on the same ray that are close to each other into a new point based on the weighted average of their covariances. Experiments showed improvements in map definition and removal of non-static obstacles. Stefanini *et al.* [25] adopt a 1-to-N ray-casting strategy to improve map merging in the face of environment dynamics. Rather than a standard point-to-point comparison, the method [25] compares measurement points against a set of expected hit points along virtual neighboring

range measurements to detect scene changes more robustly.

Given that 3D ray tracing is memory-intensive and requires dense map representations, Pomerleau *et al.* [37] use the sparse point cloud from a 3D laser directly. The map points are associated with every single reading via small conical apertures in spherical coordinates, updating the observed points closer than the mapped ones, and leaving the farther ones untouched. The approximation of ray tracing results is used to update the probability that points in the map are dynamic, using a Bayesian approach. The probability of being dynamic can be used in ICP to avoid trusting dynamic points. In the experiments, the method [37] produced a more precise and cleaner map than standard ICP matching. Instead of weighting the map points, An *et al.* [129] propose the **Dynamic Edge Link (DEL)** to model the dynamics of a pose graph's edges rather than the data itself. The observation of moving obstacles between two poses gradually decreases the weight of the respective edge until it no longer detects the obstacle. Integrating DEL in an exploration scheme, nodes with an edge weight average lower than a certain threshold, meaning frequent moving obstacles or changed structure near that node, are not considered for exploration because the robot may be unable to move to that position.

Although the standard **NDT** [82] representation does not model free space, the works of Einhorn and Gross [40, 128] and Saarinen *et al.* [21] use **NDT** with occupancy maps to explicitly model free space and adopt an exponential weighted moving average and covariance for new measurements, giving them stronger influence than old ones. Einhorn and Gross [128] propose a generic 2D/3D mapping using **NDT** and occupancy maps. The hit cells from scan-to-map alignment are updated incrementally with exponential weighting, based on the current observation. The other cells along the sensor beam, potentially empty, are updated using the standard update rule for occupancy maps, based on the log-odds of the occupancy value and the inverse range sensor model. Instead of using the standard occupancy map update, the sensor model in Saarinen *et al.* [21] relies on the inconsistency between observations and the map. Also, the occupancy value describes the confidence of the **NDT** based on past observations. As for Einhorn and Gross [40], the proposed approach defines two occupancy map probabilities: occupancy and statically occupied. The former is updated based on the sensor model (2D/3D generic beam sensor) and scan-to-map alignment, and the latter slowly adapts to high-probability static objects in the environment. The static occupancy probability follows the proposed ad-hoc model, which is parameterized to control how quickly the static occupancy probabilities are adapted, depending also on the occupancy probability itself. The works [128] and [40] handled both semi-static and dynamic changes, maintaining a consistent, up-to-date representation of the environment. In contrast, the method [21] favored long-term static structures in dynamic environments.

2.3.3 Dynamic objects detection

In order to detect dynamic objects, Yue *et al.* [43] propose a collaborative dynamic-mapping approach for human detection using visual and thermal images and a 3D **LiDAR**. The YOLOv3 object detection algorithm extracts the bounding boxes from the images relative to humans. Projecting the 3D point cloud onto images creates a static point cloud for localization and mapping of each robot, filtering out points corresponding to humans. In the experiments, removing dynamic objects produced a more accurate relative transformation of the collaborative maps than not removing them. Zhu *et al.* [13] also use YOLOv3 to extract bounding boxes of dynamic classes (parked cars) from an RGB image, and the projection of 3D **LiDAR** points allows the creation of the static and semi-dynamic maps required for localization. Even though the method [13] does not discard dynamic objects, the localization module reduces the weight of the corresponding observations.

Furthermore, Sun *et al.* [130] adapt the PointNet network for object recognition (pedestrian, cyclist, car, or background) to classify 3D **LiDAR** scan points. The proposed Recurrent-OctoMap maintains occupancy and semantic information, where the latter specifies the cell's semantic state and the prediction's probability. An **LTM Recurrent Neural Network (RNN)** learns the transition between cell states. In a long-term experiment, the method improved its 3D semantic map compared to a standard Bayes update. In contrast, Chen *et al.* [44] use PointPillars for dynamic object removal within a **LiDAR SLAM** system employing **Lidar Odometry and Mapping (LOAM)** [137]. Given the bounding boxes predicted by PointPillars, the proposed approach crops all points from the 3D point cloud within those bounding boxes for later use in localization and mapping. Comparing localization and mapping errors before and after introducing dynamic object removal, localization accuracy remained similar in the experiments, though the maps created were much cleaner and sharper. However, localization performance sometimes decreased with dynamic object removal due to feature loss caused by false detections and a possible lack of features after removal.

Pixel-wise semantic segmentation is another way to identify dynamic objects. In addition to the semantic-descriptor proposed by Singh *et al.* [46], which describes the spatial distribution of semantic entities in the scene, the method assigns lower weights to features detected in the sky and in dynamic classes (*e.g.*, person, car, *etc.*) from EdgeNet semantic segmentation. The proposed approach [46] achieved higher accuracy than SeqSLAM [138], FAB-MAP [139], and **Bag-of-Words (BoW)** [140]-based place recognition methods in a highly dynamic outdoor experiment. Similarly, Yu *et al.* [47] implement a lightweight pixel-wise semantic segmentation network using MobileNetV2 as a backbone and building upon the ORB-SLAM2 [134] framework, also focusing on dynamic object classes such as humans, animals, and cars. Instead of down-weighting dynamic classes, all dynamic visual features are removed from the image using semantic annotations to obtain a 3D static semantic map. Compared to base ORB-SLAM2 [134], localization accuracy improved significantly in dynamic scenes, with an average **Root Mean Square Error (RMSE)** reduction of 95%.

Instead of identifying object classes, Song *et al.* [45] propose the MD-Net **Convolutional Neural Network (CNN)** to segment a grayscale image into unstable, static, and moving pixel points, only using static points for localization. The localization error was reduced compared to not estimating the pixel dynamic attribute. Ganti and Waslander [12] propose the **Semantically Informed Visual Odometry (SIVO)** to improve the performance of ORB-SLAM2 [134] by using the Bayesian neural network SegNet for segmentation and for computing network uncertainty. SegNet is trained to distinguish different object classes for identifying dynamic objects (sky, car, truck/bus, person/rider, motorcycle/bicycle, and void) from static ones (road, traffic sign, building, wall/fence, pole, vegetation, sidewalk, traffic light, and terrain). The proposed method considers only static keypoints as input to ORB-SLAM2 [134], thereby reducing most of the state's uncertainty (including network uncertainty). **SIVO** removed uninformative and dynamic keypoints from the current frame. However, rejecting potential dynamic objects without verifying their motion reduced the module's localization performance in certain scenarios. Soares *et al.* [20] propose Changing-SLAM built on ORB-SLAM3 [141]. Each feature keypoint is initially classified as dynamic, movable, or static by the object detection thread (YOLOv4) based on semantic information, and two keypoints belonging to the same object cannot have different classifications. Then, the proposed method employs feature repopulation to correctly classify features within a bounding box that are not dynamic. Experiments validate the feature repopulation technique, which keeps keypoints belonging to the background but within the dynamic object class. Comparing the proposed method to DynaSLAM [18] and ORB-SLAM3 [141], even though DynaSLAM [18] had overall the best results, Changing-SLAM offered a good trade-off between accuracy and speed in dynamic environments.

Spatial consistency can improve the identification of moving objects. Bescos *et al.* [18] propose DynaSLAM, which uses multi-view geometry verification to improve the semantic labeling of potential dynamic classes generated by a Mask RNN. Although DynaSLAM [18] outperformed ORB-SLAM2 [134] in highly dynamic scenarios, its performance reduces in slower dynamics. Similarly, Wang *et al.* [19] also use geometric verification to improve the labeling of dynamic objects, which are obtained using a ResNet-based network. Verification is based on reprojection error. The method [19] improved the ATE over DynaSLAM [18] and ORB-SLAM2 [134] in a scenario with movable objects. Guo and Liu [14] also use a Mask RNN network for pixel-wise segmentation and object instance labels, similar to DynaSLAM [18]. Rather than using multi-view geometry verification solely to improve semantic labeling, it is also used to detect untrained moving objects. The approach [14] outperformed DynaSLAM [18] in terms of ATE in dynamic scenes.

Instead of using semantic segmentation, the Semantic and Geometric Constraints Visual SLAM (SGC-VSLAM) proposed by Yang *et al.* [42] uses YOLOv3 to extract bounding boxes for dynamic objects, thereby also improving ORB-SLAM2 [134]. A constraint based on epipolar geometry improves the labeling. SGC-VSLAM decreases the RMSE of the ATE by 96% compared to ORB-SLAM2 [134] in highly dynamic environments. However, like DynaSLAM [18], its performance decreased at lower dynamics. Also, Xing *et al.* [8] propose the DE-SLAM to deal with Short-Term Dynamics (STD) and Long-Term Dynamics (LTD) at the same time. A MobileNetV2 identifies bounding boxes for movable objects (cars, persons, etc.) and classifies them as STD. A motion check of STD elements recognizes all moving objects in the current keyframe. As for LTD, DE-SLAM uses Histogram of Oriented Gradients (HOG) features extracted from Oriented FAST and Rotated BRIEF (ORB) keypoints to improve its invariance to illumination changes. In the experiments, DE-SLAM improved the localization over ORB-SLAM2 [134] in a changing environment. All of these methods, which use geometric constraints to improve semantic labeling, rely solely on static features for localization and mapping.

Moreover, Wang *et al.* [33] formulate a dynamic filter on top of ORB-SLAM3 [141] that fuses information from semantic segmentation and structural constraints (a Manhattan-world assumption with structure constraints between planar features) to identify both semantically identifiable and unidentifiable dynamic objects. The authors use the DeepLab semantic segmentation algorithm, capable of identifying 20 objects, including people and chairs, and leverage optical flow to detect abnormal dynamic features. Then, the dynamic filter computes long-term dynamic object probabilities considering structure regularity (related to the distance error of feature points to the epipolar line constraint), Inertial Measurement Unit (IMU) measurements (reprojection error of feature points from IMU pose estimation), and forward/backward projection (reprojection error of co-visible features onto the current frame). In experiments, compared to ORB-SLAM3 [141], it achieved up to 92% average RMSE improvements. Omitting structural regularity from the proposed approach [33] led to a marginal reduction in localization accuracy, whereas excluding the dynamic filter module altogether resulted in notable trajectory deviation. Zhu *et al.* [28] introduce the Incremental Static-referenced Dynamic Object Removal (ISDOR) module to detect and suppress moving objects in real time from 3D LiDAR data. The module leverages static structures (*e.g.*, ground) to identify points lying above static regions, vertically suspended points with significant height, or points adjacent to previously flagged dynamic areas. Once a portion of a dynamic object is identified with high confidence, the complete object is recovered via region growing and clustering. In experiments, ISDOR segmented dynamic objects visually align with the ground-truth, accurately isolating dynamic points while preserving static structure, including removing ghost trails in the datasets.

Another common approach is to formulate static scores based on spatiotemporal consistency to identify and handle dynamic elements in the scene. Wu *et al.* [30] formulate a static score estimation of the landmark to

obtain initial static-dynamic labelling, and a **Conditional Random Field (CRF)** is constructed to refine the labels. Based on ORB-SLAM2 [134], the static score is estimated from the statistical properties of the reprojection error and then propagated to subsequent frames. Similar to Du *et al.* [32], the authors consider that adjacent landmarks in 3D space or 2D images tend to have the same dynamic-static label. A **CRF** that considers neighbor motion and historical observations is built on a fully connected graph of all landmarks, improving spatiotemporal consistency of static-dynamic labels. Comparing the accuracy and efficiency to the base ORB-SLAM2 [134] and to Du *et al.* [32], improved **ATE** and **Relative Pose Error (RPE)** by 80–90% over ORB-SLAM2 [134]. Although Du *et al.* [32] reported the lowest **RMSE** for **ATE** and **RPE**, the proposed work [30] achieved similar accuracy with only a few centimeters of disadvantage, while taking 5 times less computing time for dynamic landmark removal than Du *et al.* [32]. Hu *et al.* [35] employ YOLOX for object detection on input images, while compensating for missed detections using an **EKF** (predict the state of the bounding box of a certain object already mapped) and the Hungarian algorithm (associate **EKF**-predicted bounding boxes with those detected in the current frame). Authors formulate a static probability distribution over keypoints based on motion consistency (via projection and epipolar constraints). This static probability is used as a weight in camera pose estimation and optimization. Experiments compare the proposed approach [35] with several state-of-the-art algorithms, including DynaSLAM [18] and ORB-SLAM2 [134]. The proposed approach [35] achieved competitive localization accuracy while offering a good balance between real-time performance and localization accuracy.

More recently, the literature has explored dense environment representations and methods for handling moving elements. Cheng *et al.* [34] render the environment through a **3D Gaussian Splatting (3DGS)** model. The authors use the SegmentAnything model to generate a pixel-wise semantic image, allowing correspondence between the features and their respective object labels. The objects are also classified into entering (objects entering the frame), vanishing (objects leaving the frame), and full (objects fully visible within the frame). A weighting factor accounts for the temporal decay of feature relevance (older features having less influence). Semantic data from the SegmentAnything model is integrated and rendered using a **3DGS** formula. Then, the method [34] uses Gaussian-based encoding to represent dynamic objects, considering semantic attributes, pose, and motion velocity information, to avoid generating noisy Gaussian points from dynamic objects that can affect model convergence. The proposed approach [34] accurately rendered dynamic objects in experiments. Similarly, Jia *et al.* [36] propose the DGS-SLAM to reconstruct a static, consistent, and dense representation of the surroundings by excluding dynamic objects from the **3DGS** model. The proposed approach [36] identifies dynamic objects using motion consistency (a dynamic projection threshold to evaluate the motion state of feature points) and object segmentation, with FastInst instance segmentation to obtain pixel masks of objects. When generating new Gaussian points from keyframes, each point is assigned a label indicating whether it belongs to an object or a static background, enabling 3D-2D association. DGS-SLAM achieved the best pose estimation results across most experiments, including tracking, compared to other **3DGS**-based and **Neural Radiance Field (NeRF)**-based **SLAM** systems.

2.3.4 Prediction modeling

In the included works, Markov processes are used to predict the dynamics of the environment. Tipaldi *et al.* [48] use a dynamic occupancy grid and exploit the stationary distribution and the state holding time associated with **Hidden Markov Models (HMMs)** on a 2D grid. The method [48] uses past observations from each run to

learn the state transition probabilities iteratively and estimate the **HMM** parameters. Then, localization can infer how often a dynamic object is expected to be seen in the environment and for how long. Comparing the proposed **HMM**-based localization to **MCL** using a standard grid, the former had a lower localization failure rate than **MCL**. Indeed, work [48] could handle both high-dynamics (moving cars) and low-dynamics (parked cars). Rapp *et al.* [49] implement a semi-Markov process extended by a Levy process to model time dependence in the state-holding time of Markov processes, also predicting, as in Tipaldi *et al.* [48], the expected retention time for each cell in a specific state. In the experiments, the proposed model by Rapp *et al.* [49], integrated into **MCL**, improved the classic **MCL** in a dynamic environment.

Environment dynamics may exhibit periodic patterns. Assuming periodic change patterns, Krajník *et al.* [52] propose **Frequency Map Enhancement (FreMEn)** to model the probability of occupancy or feature visibility in a grid as a combination of harmonic functions associated with periodic processes. **FreMEn** [52] uses spectral analysis (Fourier transform) to compute the harmonic functions and predict the future state with a given confidence. In a changing environment, **FreMEn** [52] outperformed a static map in terms of localization error by selecting the most likely visible features at each location. Santos *et al.* [51] adopt **FreMEn** [52] within an exploration scheme, where the planner predicts which areas are more likely to change at a certain time and generates the subsequent locations to explore. The experimental results showed that accounting for environment dynamics increases the amount of information gathered compared to static models.

Unlike **FreMEn** [52], Wang *et al.* [50] model both aperiodic and periodic changes using an **Auto-regressive Moving Average Model (ARMA)** model [142]. This model describes time series as stationary stochastic processes in terms of polynomials. While **FreMEn** [52] can update its model recursively online, **ARMA** [50] is updated only once a day based on past observations. However, the model by Wang *et al.* [50] achieved higher prediction accuracy than **FreMEn** [52] and a lower localization failure rate than both **FreMEn** [52] and Tipaldi *et al.* [48]. Wang *et al.* [16] extend the original work of Wang *et al.* [50] by adding long-term path planning to generate feasible paths with low localization uncertainty and fewer dynamic objects, while introducing new spatial representation models (semi-static vs dynamic density maps) and improving parameter calculation. In the experiments, the authors simulate a large-scale parking lot with periodic and aperiodic changing car patterns. As already shown in Wang *et al.* [50], since **FreMEn** [52] is based on a spectral model, it cannot effectively handle non-periodic growth patterns compared to **ARMA** models [16, 50]. Also, when evaluated in a real parking lot, the system proposed by Wang *et al.* [16] was able to predict more accurately for most parking spaces than **FreMEn** [52], being more stable across various real-world change patterns.

Den *et al.* [15] build a time series for each map point to predict their future states. The proposed approach uses a **Bayesian Persistence Filter (BPF)** and survival analysis to estimate the likelihood that each map point will still exist at a future time. The system classifies the map points as static or semistatic based on its persistence probability, in which points with a persistence probability lower than a certain threshold are considered dynamic ones and removed from the map. In the experiments, the authors compared the proposed prediction algorithm [15] to **FreMEn** [52] and **ARMA** [142] models in terms of predictive capabilities. The proposed method outperformed other prediction methods. As for localization accuracy, **BPF**-based localization had lower **ATE** in low-dynamic environments compared to ORB-SLAM2 [134], ORB-SLAM3 [141], and DynaSLAM [18]. In high-dynamic scenes, although the proposed method improved accuracy over ORB-SLAM2 [134] and ORB-SLAM3 [141], DynaSLAM [18] had lower **ATE** metrics.

Instead of modeling the dynamics in the map, Thomas *et al.* [39] use a KPConv network to predict online dynamic motion labels for points given a single 3D **LiDAR** as input. The method [39] is a self-supervised learning

approach with two main modules: PointMap and PointRay. PointMap is an ICP-based SLAM algorithm to provide a point cloud map for the annotation process. PointRay uses a similar approach to Pomerleau *et al.* [37] to approximate ray tracing using spherical coordinates for obtaining the training annotation of dynamic labels: permanent (static points over all sessions), ground (to avoid ray tracing ground samples), and long-term (still objects in single sessions but relocated between sessions) and short-term (dynamic objects) movables, with the localization not considering the latter two. In simulation experiments, PointMap with the proposed prediction module yielded lower localization pose errors than the standard MCL algorithm.

Moreover, Hroob *et al.* [26] propose the Long-Term Stability NETwork (LTS-NET). This regression network estimates a stability value for each point in a 3D LiDAR point cloud, thereby providing a spatiotemporal stability measure. LTS-NET is built on PointNet++ and uses labels generated unsupervised from previous observations and the map, with the NDT [143] employed for scan matching. Once trained, the network is applied online to identify stable points in the current LiDAR session, filtering out unstable observations. Experiments showed that filtering scans based on long-term stable features significantly enhances localization performance and robustness in dynamic environments, enabling the network to identify and filter long-term stable objects. These experimental results led to improved localization performance compared to not using network inference on stability scores. As for Hroob *et al.* [27], the authors formulate a more advanced method that also relies on NDT [143] for scan matching, but replaces PointNet++ with a modified MinkUNet14 to process sparse 4D tensors containing spatial coordinates and time data. Rather than inferring stability primarily from training data, this method explicitly models the spatio-temporal discrepancy between the current LiDAR scan and the point cloud map. Scan and map tensors are merged into a unified sparse 4D representation, after which Hroob *et al.* [27] prune voxels containing only map timestamps, since the goal is to infer stability scores only for scan points. By directly exploiting scan-map discrepancies, the method [27] achieved stronger generalization in the experiments. Indeed, it outperformed LTS-NET [26] in segmenting unstable elements such as humans, vegetation, parked cars, and pedestrians, cases in which LTS-NET [26] was less effective, likely due to its stronger dependence on object shapes observed during training.

2.4 Summary and Conclusions

SLAM is one of the most fundamental problems in robotics. When a robot lacks prior knowledge of the environment and does not know its own pose, SLAM enables an AMR to build a map of its surroundings while simultaneously localizing itself within it. A basic strategy for solving the SLAM problem is to continuously align the latest sensor observation with the map up to the current instant, then integrate the new measurement into the map using the estimated alignment pose, as in scan-matching systems such as HectorSLAM [72]. However, a pure scan-matching approach does not explicitly address loop closure, which limits the correction of accumulated drift when the robot revisits a previously mapped place. Another popular family of approaches is based on particle filters, specifically Rao-Blackwellized Particle Filter (RBPF), as in GMapping [74] (adaptation of RBPF to 2D occupancy grid mapping). In these methods, the posterior distribution over part of the state, such as the robot's trajectory, is represented using particles. At the same time, the remaining variables are modeled using Gaussians or other parametric distributions, resulting in higher computational efficiency than standard particle filters [78]. More recently, pose-graph formulations have become a dominant paradigm for solving the SLAM problem. In graph-based SLAM, the front-end identifies constraints, such as odometric relations between consecutive poses, loop closures with previously visited poses, and landmark observations from one or

multiple viewpoints, by establishing correspondences between new observations and the current map. The back-end then refines the graph by jointly minimizing the inconsistency of all constraints [87]. An implementation of this formulation for 2D laser scanners and wheel odometry is **SLAM Toolbox** [95]. Because the pose-graph formulation retains access to the full set of measurements and constraints during map optimization, it allows the use of improved linearization, loop-closing, and data-association strategies, which is one of the main reasons why graph-based **SLAM** often outperforms purely filter-based approaches. Nevertheless, pose-graph methods usually exhibit at least linear growth of the graph with the explored area, which makes graph sparsification, submapping, and map-maintenance strategies important for long-term operation [65, 78].

In a static environment, the robot would only need to map it once, since the map and incoming sensor data would remain consistent over time. Nonetheless, real environments contain moving agents, temporary occlusions, furniture rearrangement, seasonal or illumination changes, and other sources of temporal variability. Long-term **SLAM** addresses precisely these issues. In our previous work [65], we conduct a systematic literature review on long-term localization and mapping following the **Preferred Reporting Items for Systematic reviews and Meta-Analyses (PRISMA)** [98] guidelines. To the best of our knowledge, before that publication, no other work had reviewed the broader topic of long-term **SLAM** in a unified way, covering appearance variation for localization and place recognition, the modeling and handling of scene dynamics, redundant data removal from maps, multi-session mapping, computational aspects of long-term operation, datasets, and evaluation metrics. Using a systematic methodology to ensure repeatability and reproducibility, we review 142 works across 6 scientific databases. Overall, the scientific community have shown increasing interest in long-term localization and mapping. The subtopic of **SLAM** in dynamic environments, which is the focus of this PhD Thesis, has also experienced clear growth over time.

Subsequently, this PhD Thesis builds on the results of our systematic review, narrowing the discussion to long-term **SLAM** in dynamic environments and updating the search date of the last full inquiry to February 13, 2026. This update results in 53 works analyzed in Section 2.3. A trend in the literature is that most methods still rely on conventional **SLAM** map representations, such as graphs, occupancy grids, point clouds, or feature maps, and then augment them with mechanisms to reason about dynamics. In other words, the map remains primarily a static structural representation, while dynamic behavior is encoded as an additional property of cells, points, features, or objects. A recurring strategy, especially in human memory-inspired works, is to organize information into **Sensory Memory (SM)**, **Short-Term Memory (STM)**, and/or **Long-Term Memory (LTM)** [4, 6–9]. These works use selective attention, rehearsal, retrieval, and pruning policies to determine which observations should be included in the persistent map and which should remain temporary or ignored. Another common trend is to partition the map into static, semi-static, and/or dynamic components [10–17], or equivalently into permanent, movable, and dynamic layers [18–20]. In these formulations, static or permanent structures are expected to remain stable over long periods; semi-static or movable elements are objects that may persist for some time but can change between visits; and dynamic elements are objects currently moving or highly unstable in the sensor field of view.

Furthermore, methods that exploit data consistency are common in the literature for identifying scene dynamics. One typical strategy is to align the current sensor observation with the current map or with previous observations and then detect mismatches. Depending on the sensing modality, these spatial inconsistencies manifest as scan-to-map or scan-to-scan mismatches with range sensors [6, 10, 11, 15, 17, 21–29], reprojection errors with visual modalities [19, 30–33], motion inconsistencies in tracked features [8, 19, 31, 33–36], ray-tracing discrepancies that reveal free-space discrepancies or newly added objects [4, 9, 11, 21, 22, 25, 37–39], and

multi-session inconsistencies across repeated traversals of the same place [24, 29, 38]. Clustering is often used as a complementary step to group nearby inconsistent observations and turn isolated outliers into more coherent dynamic hypotheses [7, 16, 28, 35, 38, 40, 41]. Also, semantic perception has become a central component of long-term **SLAM** in dynamic environments, particularly in vision-based approaches. Object detection and semantic segmentation are frequently used to identify categories that are likely to be dynamic or movable, such as people, vehicles, bicycles, or animals [8, 13, 20, 35, 42–44]. Pixel-wise segmentation is also employed in the literature to discriminate between dynamic and static image regions and avoid excessive feature depletion compared to bounding-box semantic segmentation [12, 14, 18, 19, 33, 34, 36, 45–47]. Indeed, bounding-box detectors often require additional geometric constraints, background/foreground refinement, or feature repopulation mechanisms to avoid incorrectly rejecting useful static features within coarse detections [20].

Statistics are widely used in the literature to describe and, in some cases, predict scene dynamics. Early methods often relied on relatively simple quantities, such as hit-and-miss counts [6], change scores [10], observation counts [15, 32, 38], or feature stability histograms [7], to estimate how often a cell, point, or feature was observed consistently over time. Later works introduced richer temporal models, including **HMM** [48], semi-Markov processes [49], **ARMA** models [16, 50], persistence filters [15], and frequency-analytical models based on the **Fast Fourier Transform (FFT)** [51, 52]. These models are typically not used to forecast detailed future trajectories of dynamic objects, but rather to estimate a persistence, reliability, or staticness score for map elements, or even to identify periodic temporal patterns. Possible uses of predicting dynamics include inferring whether a region of the environment is stable, periodically occupied, or likely to change. Such information can then support more robust localization systems, more selective map updates, and even planning strategies that avoid highly dynamic areas or unfavorable times for revisiting them.

Most methods in the literature aim to protect localization from transient or unstable observations by relying primarily on static maps, or at most on static plus semi-static components, while excluding clearly dynamic elements from the pose-estimation process. In some approaches, dynamic and semi-static observations are removed before being inserted into the long-term map [9–11, 15, 22, 23, 26–28, 30–33, 36, 38, 39, 41, 44, 47, 129]. In contrast, other approaches do not discard those observations. Instead, dynamic observations and possibly semi-static ones are down-weighted during scan or feature matching, or are assigned larger covariance to reduce their influence on the pose estimate [13, 20, 37, 46]. Vision-based methods often use masking and background inpainting to remove dynamic image regions while preserving a usable visual observation for tracking and mapping [14, 18]. Overall, the dominant trend is that long-term **SLAM** systems in dynamic environments remain fundamentally static-map-oriented. Rather than building a fully dynamic world model, they estimate the stability of observations and decide whether to keep them, temporarily down-weighted, or remove them altogether. This trend makes the notion of persistence, rather than motion alone, one of the central concepts in the current long-term **SLAM** literature for dynamic environments.

In conclusion, the literature on long-term **SLAM** in dynamic environments shows progress in modeling persistence, detecting dynamic elements, and improving localization robustness. Still, the literature also reveals several gaps that motivate this PhD Thesis. Most existing approaches remain fundamentally static-map-oriented: they typically identify, mask, down-weight, or remove dynamic observations. However, current methods rarely maintain a continuously updated map representation that explicitly replaces outdated structure with the most recent consistent observations. In addition, the literature tends to treat mapping quality, localization robustness, and long-term map maintenance as partially separate problems, with limited integration between online dynamic-aware operation and offline global refinement. Another issue is scalability. Pose-graph and map

representations often grow over time, even when the robot revisits known areas, introducing redundant data, increasing computational cost, and degrading map quality through accumulated noise and distortion. To address these limitations, this PhD Thesis proposes a unified **SLAM** framework for ground mobile robots in dynamic indoor environments that combines: online ray tracing-based map updating to remove outdated or redundant observations; relocalization on existing graph structures to prevent unnecessary graph expansion; offline generation of cleaner static maps by removing dynamic points; and offline global refinement to correct accumulated distortions. In parallel, this PhD Thesis researches geometric features, distance-map-based registration with robust kernels, and continuously updated local maps. The final goal is to actively incorporate environment dynamics instead of just filtering them out, to provide a more stable, scalable, and accurate **SLAM** solution in highly dynamic environments.

Chapter 3

Datasets

TBC

3.1 Mobile Platforms

TBC

3.1.1 5dpo three-wheeled omnidirectional robot

3.1.2 Modified Hangfa Discovery Q2 four-wheeled omnidirectional platform

3.1.3 IDEPA's tricycle autonomous robot

3.1.4 Avantgarde tricycle autonomous forklift

3.2 Faculdade de Engenharia da Universidade do Porto

TBC

3.3 IILABS 3D – iiLab Indoor LiDAR-based SLAM

TBC

IDEPA

Cappero

Avantgarde at COLEP

Kuka

3.4 Summary and Conclusions

TBC

Chapter 4

Line Fitting-Based Corner-Like Detector for 2D Laser Scanners Data

This chapter presents a novel line-fitting-based corner-like detector for 2D laser-scanning data [66], developed within the scope of this PhD Thesis (see Appendix C). Even though the most common way to represent 2D environments is via occupancy grid maps, an alternative is to represent the scene as a set of features. Accordingly, this chapter explores the extraction of features that may resemble geometric corners in the environment. Our methodology explores local line fitting at each scan point and the intersection of the fitted lines before and after the point, with respect to the angular scanning direction. Experimental results reveal viewpoint and geometric invariance of the corner-like features extracted using the proposed methodology, demonstrating the potential of this work for localization and mapping problems on mobile robots.

Thus, this chapter is divided into four sections. First, Section 4.1 introduces feature extraction algorithms for 2D laser scanner data from the literature. Section 4.2 outlines our methodology for extracting corner-like features from 2D laser data using local line fitting and multiple neighborhood-scale analysis. Then, Section 4.3 presents experimental results to investigate the observability ratio of the proposed corner-like feature detector, the viewpoint- and geometry-invariance properties of the extracted features, and the extraction of line segments in parallel with the proposed corner-like feature detector. Lastly, Section 4.4 presents a summary of this chapter, the main conclusions, and future work.

4.1 Background

Laser scanners allow building a dense representation of the environment as a 2D occupancy grid map, where each cell represents occupied, free, or unknown space [74]. Although grid maps can represent arbitrary objects, they are computationally intensive and require substantial memory [144]. An alternative method is to represent the environment as a set of features. In the case of 2D lasers, the features may include beacons [145], geometric cues (line [146–154] and curve [150–152, 155] segments, corners [146–152, 155–159], and free endpoints of segments [146, 151, 158, 159]), and semantic features such as parked vehicles [156, 160] and doors [147].

The correspondence between geometric features and the environment (*e.g.*, lines and corners) may be of interest for localization and mapping problems on mobile robots, as well as for user interaction with autonomous mobile systems. The current literature presents several algorithms for extracting line segments from 2D laser

scan data [146–154], including an open-source implementation based on the ROS [96] framework¹. As for corner detection, the most common methodology found in the literature is to extract corners from the intersection of line segments detected in the laser scan [146–149, 155–158, 160–162]. Another typical approach in the literature is to estimate curvature from the scan data [150–152, 159]. Next, both trends for extracting corner-like features are analyzed based on methods found in the literature.

4.1.1 Line fitting intersection

A trend in the literature is to extract corners at the intersections of lines detected in the 2D laser scan [146–149, 155–158, 160–162]. First, the laser data may be filtered to remove invalid values (zero range values [156]) and possible noise (*e.g.*, considering the minimum distance to the two neighbors of a certain point to remove isolated data [156], mean average filtering [158], or bilateral filters to remove non-linear noise [149]). Next, a segmentation step extracts sets of continuous points in the environment for line fitting estimation. Castellanos *et al.* [146] segment the scan into polygonal lines and subsequently divide them into segments using an iterative method. Falomir *et al.* [155] search for discontinuities in the data to form point subsets and use a recursive line-fitting algorithm [163] to obtain line segments. Fernández-Moral *et al.* [161] search for parameters of a 2D line that maximize the number of points supporting the line model using the **RANdom SAmple Consensus (RANSAC)** algorithm [164]. Mohamed *et al.* [148] use the line tracking algorithm [165] to iteratively segment lines in the scan and perform **Principle Component Analysis (PCA)** to compute the lines' parameters. Then, the intersection points between consecutive line segments define a corner feature in the environment.

Furthermore, Lin *et al.* [157] and Liu *et al.* [149, 158] use adaptive threshold segmentation to cluster range data into continuous spatial subsets. Although Lin *et al.* [157] state that an adaptive segmentation was used due to the non-uniform distribution of the 2D laser data, the authors do not specify the threshold used in their work. In Liu *et al.* [158] work, the threshold is based on the laser's angular resolution and the distance to the current point under analysis. Liu *et al.* [149] also establish a threshold depending on the angular resolution. Instead of considering only the distance of the point under analysis, the proposed approach [149] selects the larger distance between adjacent points. Even though Roumeliotis and Bekey [147] do not use a specific adaptive threshold, the proposed SEGMENTS algorithm with two-level **EKF** exhibits adaptive behavior. A stricter filter, in terms of the system's noise covariance, attempts to follow a straight line, while a more flexible **EKF** tries to follow a new segment during each scan. The four methods [147, 149, 157, 158] also define a corner feature as the intersection between consecutive line segments.

A specific application of corner detection with 2D laser data in the literature is vehicle detection [156, 160]. Jung *et al.* [156] cluster laser data using a fixed threshold between consecutive range values. Then, the method performs template matching for every point under analysis, assuming the point is the vertex of the L-shape of a vehicle in the scan. Jung *et al.* [156] use **Singular Value Decomposition (SVD)** to estimate the L-shape's line parameters and decompose the problem into line-curve estimation when the vehicle has a curved corner in the data. In contrast, Qu *et al.* [160] segment the data using **Density Based Spatial Clustering of Applications with Noise (DBSCAN)** [166]. The method performs L-shape fitting in two steps: searching for the L-shape vertices and corner-point localization, the latter based on **RANSAC** [164].

However, corner detection algorithms that rely on global search during line segment scanning may be limited to well-structured environments. Thus, we propose a corner-like detection method based on local line fitting,

¹https://github.com/kam3k/laser_line_extraction

without requiring global line detection or recursive data search. Also, the experiments presented in Section 4.3 demonstrate the ability of the proposed detector to extract more points of interest from the scan compared to the algorithms presented previously.

4.1.2 Curvature-based detection

Other approaches to detecting corners in 2D laser data estimate the curvature. Núñez *et al.* [150] propose an adaptive curvature estimator that compares the k -nearest points on both sides of the point being evaluated. The k size is adjusted based on the distance between possible corners. A segmentation step precedes the curvature estimation to detect breakpoints using the adaptive breakpoint detector proposed by Borges and Aldon [167]. Then, Núñez *et al.* [150] extract corner, line, and curve segment features from the scan data by analyzing whether there are peaks in curvature, flat segments with zero average curvature, and flat segments with non-zero average curvature, respectively. Vázquez-Martín *et al.* [151] use the adaptive curvature estimator of Núñez *et al.* [150] to segment the scan into clusters of homogeneous curvature. Even though the existence of a corner is determined from the curvature function, its keypoint definition is based on the intersection of the two line features that generate the corner, to improve robustness to noise. Also, Certad *et al.* [159] modify the work of Núñez *et al.* [150] to consider two curvature estimations based on different noise-related thresholds, thereby improving the rate of true detection for corner and edge features in the experiments.

Moreover, Feng *et al.* [152] and Zhao *et al.* [162] report a correlation between the angle difference between the two fitted lines on both sides of a scan point and the curvature value. Feng *et al.* [152] introduce the micro-tangent line function as the variance between the tangents at two consecutive points in the scan. Similarly to Núñez *et al.* [150], line and curve segments are extracted from consecutive range readings with variance close to zero and with constant value larger than a threshold, respectively, and corner features are extracted from peaks in the variance function. Instead of only considering two consecutive points, Zhao *et al.* [162] define two distance-based regions on both sides of a point to estimate direction vectors and, consequently, the angle difference between those vectors. Candidate corner points are selected based on a fixed threshold for the angle difference. Then, the corner point is estimated via non-minimal suppression.

Still, most curvature-based methods for corner feature extraction associate the keypoint with a laser reading. Given the discrete, radial nature of 2D laser, the estimated corner may vary farther away. Also, there is the question of whether the curvature estimators are invariant to scale, angular resolution, and different viewpoints. We propose a detector based on the stability of intersection points from fitted lines at different distance-based scales. The experiments in Section 4.3 focus on evaluating the invariance of the proposed corner-like features.

4.2 Corner-Like Feature Detector

We propose a feature extraction algorithm on 2D laser data performing local line fitting to detect corner-like points in the scan. In summary, first, we transform the scan data from polar to Cartesian coordinates. This transformation maintains the data's angular coordinate order. For each point in the scan, our methodology considers multiple distance-based scales to extract arrival and departure neighborhoods, depending on whether the latter are in the negative or positive directions relative to the point's angular coordinate, respectively. Next, a linear regression algorithm fits lines to each neighborhood. The intersection point between the arrival and departure neighborhoods is a candidate feature keypoint at a given scale. Then, the following indicators are

estimated to determine if the candidate keypoint is a corner-line feature or not: minimum number of points to fit a line, angle between arrival and departure lines, maximum **RMSE** allowed for the line fitting procedure, maximum distance of the candidate point to the intersection centroid on the multiple scales, and maximum distance to the laser point on analysis. The candidate keypoint is validated if it passes all previous indicators over a minimum number of consecutive scales, thereby improving robustness to noise in the laser data.

4.2.1 Linear least-squares regression

A generic definition for the 2D line model is presented in Equation (4.1), where (x, y) are the 2D point coordinates and γ_0 , γ_1 , and γ_2 are the model parameters. However, the solution for the generic definition estimated with a linear least-squares formulation from a set of 2D points would be a null vector.

$$\gamma_0 + \gamma_1 x + \gamma_2 y = 0 \quad (4.1)$$

An alternative formulation is to characterize the 2D line relative to one of the point coordinates, reducing from three to two parameters in the model. The 2D line formulation illustrated in Equation (4.2) is relative to the y coordinate with respect to x , where β_1 is the line slope and β_0 is the line intersection on the y -axis. The least-squares solution for the parameters β_1 and β_0 from a set of points $(x_i, y_i), i = 1, \dots, N$ is formulated in Equation (4.3) and Equation (4.4), where $(\bar{x}, \bar{y})$ is the centroid point of the data [168]. Still, this model would not be valid when $\sum_{i=1}^n (x_i - \bar{x})^2 = 0$, given the indeterminate form caused in the parameter $\hat{\beta}_1$ (e.g., all the data points with the same x coordinate and different y values).

$$y = \beta_0 + \beta_1 x \quad (4.2)$$

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} \quad (4.3)$$

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x} \quad (4.4)$$

The 2D line can be defined relative to x with respect to the y coordinate, as formulated in Equation (4.5). The parameters of this model are the line slope (α_1) and the line intersection on the x -axis (α_0). Nevertheless, this model has also a undetermined form when $\sum_{i=1}^n (y_i - \bar{y})^2 = 0$ (e.g., all the data points with the same y coordinate and different x values).

$$x = \alpha_0 + \alpha_1 y \quad (4.5)$$

$$\hat{\alpha}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (y_i - \bar{y})^2} \quad (4.6)$$

$$\hat{\alpha}_0 = \bar{x} - \hat{\alpha}_1 \bar{y} \quad (4.7)$$

Instead of only considering a single line model for data fitting, we select the best line model that fits the laser data. This selection is based on which model has the lower fitting **RMSE**, ε_{RMSE} . Equation (4.8) and Equation (4.9) represent the fitting **RMSE** of the 2D line models formulated in Equation (4.2) and Equation (4.5), respectively. The validation of the proposed selection approach for the line models is shown in Figure 4.1,




Figure 4.1: Error analysis of simple line regression models evaluating the fitting **Root Mean Square Error (RMSE)**, ϵ_{RMSE} , for lines with an inclination angle between -90° and 90° and a simulation step of 0.25° , a set of 100 data points equally distributed along each line, a distance Gaussian noise added to the data point coordinates of zero mean and 0.03 m standard deviation (same specifications as the Hokuyo UST-10LX sensor), and 1000 simulation samples per line [66].

where selecting the model with the lowest **RMSE** achieves the best fitting line estimation.

$$\epsilon_{RMSE, \{\hat{\beta}_0, \hat{\beta}_1\}} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i)^2} \quad (4.8)$$

$$\epsilon_{RMSE, \{\hat{\alpha}_0, \hat{\alpha}_1\}} = \sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \hat{\alpha}_0 - \hat{\alpha}_1 y_i)^2} \quad (4.9)$$

Other models and fitting algorithms could have been considered to fit a line to the 2D laser data, such as formulating a line as the vector to the normal of the infinite line and the position of the weighted mean of the contributing points [153], or using **RANSAC** to estimate the line parameters of the generic line model [164]. However, these models and fitting algorithms may require iterative-based optimization procedures and matrices inversions. In contrast, our methodology has the advantage of having closed-form equations for both the line parameters and the fitting error.

4.2.2 Feature extraction

In terms of extracting features from 2D laser data, first, we need to transform the scan data from polar to cartesian coordinates. This transformation preserves the data ordering by its angular coordinate to speed up the search for the closest consecutive laser points in the angular direction at a given distance-based neighborhood.

Next, we evaluate each of the laser points (x_i, y_i) , $i = 1, \dots, N$ as a possible feature keypoint. Figure 4.2 presents an example of the clustering and line fitting results following the proposed methodology on a laser point (x_j, y_j) at a distance scale $R_k = 0.15$ m, where this scale defines the distance-based neighborhood around the scan point. These results illustrate the clustering of the arrival $((x_i, y_i) \in \{(x_j, y_j)_{arr, k}\})$ and departure




Figure 4.2: Example of the proposed corner-like feature extraction methodology on a laser point (x_j, y_j) at a distance scale $R_k = 0.15$ m. Scan points (x, y) simulated considering a 2D laser scanner with 0.25° angular resolution, 0.06–10.0 m range, and Gaussian noise with zero mean and 0.03 m standard deviation (same specifications as the Hokuyo UST-10LX sensor) [66].

$((x_i, y_i) \in \{(x_j, y_j)_{\text{dep},k}\})$ neighborhoods, representing the closest consecutive laser points at the negative and positive sides of the point (x_j, y_j) in the laser angular direction, respectively. The intersection point between the fitted lines $line_{\text{arr},k}$ and $line_{\text{dep},k}$ is considered a candidature feature keypoint $((x_j, y_j)_{\text{cand},k})$. This intersection point is the result of the lines intersection depending on which line model is selected for each neighborhood:

$$\begin{aligned} & \bullet \quad y = \beta_0^{\text{arr},k} + \beta_1^{\text{arr},k} x \quad ; \quad y = \beta_0^{\text{dep},k} + \beta_1^{\text{dep},k} x \\ & \quad \begin{cases} x_{j,\text{cand},k} = (\hat{\beta}_0^{\text{dep},k} - \hat{\beta}_0^{\text{arr},k}) / (\hat{\beta}_1^{\text{arr},k} - \hat{\beta}_1^{\text{dep},k}) \\ y_{j,\text{cand},k} = \hat{\beta}_0^{\text{arr},k} + \hat{\beta}_1^{\text{arr},k} \cdot x_{j,\text{cand},k} \end{cases} \end{aligned} \quad (4.10)$$

$$\begin{aligned} & \bullet \quad x = \alpha_0^{\text{arr},k} + \alpha_1^{\text{arr},k} y \quad ; \quad y = \beta_0^{\text{dep},k} + \beta_1^{\text{dep},k} x \\ & \quad \begin{cases} y_{j,\text{cand},k} = (\hat{\alpha}_0^{\text{arr},k} \hat{\beta}_1^{\text{dep},k} + \hat{\beta}_0^{\text{dep},k}) / (1 - \hat{\alpha}_1^{\text{arr},k} \hat{\beta}_1^{\text{dep},k}) \\ x_{j,\text{cand},k} = \hat{\alpha}_0^{\text{arr},k} + \hat{\alpha}_1^{\text{arr},k} \cdot y_{j,\text{cand},k} \end{cases} \end{aligned} \quad (4.11)$$

$$\begin{aligned} & \bullet \quad y = \beta_0^{\text{arr},k} + \beta_1^{\text{arr},k} x \quad ; \quad x = \alpha_0^{\text{dep},k} + \alpha_1^{\text{dep},k} y \\ & \quad \begin{cases} y_{j,\text{cand},k} = (\hat{\alpha}_0^{\text{dep},k} \hat{\beta}_1^{\text{arr},k} + \hat{\beta}_0^{\text{arr},k}) / (1 - \hat{\alpha}_1^{\text{dep},k} \hat{\beta}_1^{\text{arr},k}) \\ x_{j,\text{cand},k} = \hat{\alpha}_0^{\text{dep},k} + \hat{\alpha}_1^{\text{dep},k} \cdot y_{j,\text{cand},k} \end{cases} \end{aligned} \quad (4.12)$$

$$\begin{aligned} & \bullet \quad x = \alpha_0^{\text{arr},k} + \alpha_1^{\text{arr},k} y \quad ; \quad x = \alpha_0^{\text{dep},k} + \alpha_1^{\text{dep},k} y \\ & \quad \begin{cases} y_{j,\text{cand},k} = (\hat{\alpha}_0^{\text{dep},k} - \hat{\alpha}_0^{\text{arr},k}) / (\hat{\alpha}_1^{\text{arr},k} - \hat{\alpha}_1^{\text{dep},k}) \\ x_{j,\text{cand},k} = \hat{\alpha}_0^{\text{arr},k} + \hat{\alpha}_1^{\text{arr},k} \cdot y_{j,\text{cand},k} \end{cases} \end{aligned} \quad (4.13)$$

The following indicators determine if a candidate keypoint is a corner-like feature or not: minimum of 5 points in each neighborhood, an absolute angle between arrival and departure fitted lines in the range of $[45^\circ, 135^\circ]$, maximum ε_{RMSE} in line fitting of 0.025 m, maximum distance of the lines intersection points on the multiple scales to the intersection centroid of 0.025 m, and maximum distance of the intersection point to the scan point being evaluated of 0.05 m. The latter indicator avoids having features not representing physical

points. The indicator thresholds were determined empirically through extensive testing across multiple environments, selecting values that maximize the number of detected features while maintaining geometric consistency. Then, the candidate keypoint is validated if the point passes all previous indicators in a minimum number of 3 consecutive scales, to improve robustness to noise on the laser data. The scales considered in the experiments presented in this chapter are the following ones: [0.05, 0.075, 0.1, 0.125, 0.15, 0.175, 0.2, 0.225, 0.25] m. However, the indicators' thresholds and the scale may be parametrized in run-time. The feature keypoint definition is the lines intersection centroid on the multiple scales.

4.3 Experimental Results and Discussion

The experiments were performed using a three-wheeled omnidirectional robot developed by the 5dpo robotics team [169]. This platform was adapted to have a 2D laser scanner, namely, the Hokuyo UST-10LX with 270° FOV, 0.25° angular resolution, and 40 Hz sampling rate. The sensor is aligned with the geometric center of the robot, and the wheel odometry parameters were estimated using the OptiOdom [57] calibration method. The implementation of the feature extraction algorithm is in C++ and ROS. The algorithm ran on a computer with a processor Intel Core i7-6700HQ with 8GB of RAM, and a Nvidia GPU GeForce GTX 960M.

In terms of the experimental methodology, first, a map of the scene was retrieved using the SLAM Toolbox [95] algorithm with a resolution of 0.03 m. Although this algorithm already estimates the robot pose relative to the map being created, this estimation may be improved if the whole map is used to localize the robot. So, the Adaptive Monte Carlo Localization (AMCL) [78] algorithm available in ROS was used to estimate the robot's localization relative to the map obtained from the SLAM process. With this localization estimation, the point clouds from the 2D laser scanner can be transformed to the map coordinate frame that is a static one, instead of being referenced to the moving laser frame. All parameters of the algorithms, launch configuration files, and scripts used in the experiments are also available in a public GitHub repository².

4.3.1 Feature extraction results

Figure 4.3 is an example of the feature extraction results. This example represents the feature keypoints $(x, y)^{t=t_i}$ at a time instant t_i , where the blue circle is the robot pose estimated by the AMCL [78] algorithm and the orange dotted lines represent the FOV of the sensor. The points $(x, y)^{\text{persist}}$ represent the locations of the 39 most persistent features in space extracted by the proposed methodology. These locations were found through manual selection of the features location considering their observability through the experiment (see Section 4.3.2). Further results relative to the features extraction with the proposed methodology are available in videos³.

When analyzing visually the example results, the persistent features correspond not only to well-defined 90° corners but also to other corner-like points of interest. This observation is shown both in Figure 4.3 and in the videos³, even detecting stable features in legs of passing people through the environment. As for the difference between the features detection at a certain time instant and the most persistent ones, this difference may be due to noise in the distance measures of the 2D laser.

²https://github.com/INESCTEC/inesctec_mrdt_line_fit_based_corner_detector_2d_laser

³<https://www.youtube.com/watch?v=desB9LX4hck&list=PLvp8fJUEPxYQjsuJAeiBVLEKuBU4Otu7e>




Figure 4.3: Example results for the proposed corner-like feature extraction methodology. Occupancy grid map obtained by using the SLAM Toolbox [95] algorithm, considering a map resolution of 0.03 m/px [66].

4.3.2 Features observability

Next, the observability of the features extracted when moving the robot through the scene shown in Figure 4.3 was evaluated by processing the features location in the global map frame, using the experimental methodology explained previously. Figure 4.4 represents a color map of the features observability ratio, considering the FOV of the sensor and the robot pose throughout time. This color map considers the same resolution of 0.03 m/px as the one used to localize the robot. Also, the observability map allowed the manual selection of the 39 persistent features also shown in Figure 4.3.

The first observation is the lower number of false positives. Other than the locations near the 39 most persistent features, the observability ratio of the other areas of the map is lower than 1.0%. Furthermore, there is a blur in the observability ratio near the location of the 39 persistent features, instead of being concentrated in a single pixel. This blur is a consequence of the uncertainty of the location estimation of the robot, explaining why the maximum observability ratio when considering the resolution of 0.03 m/px of the map was only 42.23%. Indeed, if the observability analysis considers the N8-neighborhood of each pixel, this maximum observability ratio increases to 100% for 6 features. For the same consideration of the N8-neighborhood, 20 and 16 features have an observability ratio greater or equal to 50% and 75%, respectively. These results indicate the possible usefulness of the proposed corner-like feature detector for tasks such as localization and mapping using 2D lasers, given the stability of the feature keypoints.

4.3.3 Features association distance

Even though the analysis on the features observability indicate if the features are consistently seen in space, this evaluation is highly dependent on the location estimator required to transform the point cloud to a static frame. Another possible analysis in the context of feature-based localization systems is to evaluate the distance between different features throughout time, assessing their relative position instead of the absolute one in space. So,




Figure 4.4: Evaluation of the corner-like features observation ratio, considering the robot pose along time, the **Field-of-View (FOV)** of the 2D laser scanner, and a map resolution of 0.03 m/px [66].




Figure 4.5: Standard deviation analysis of the distance between corresponding corner-like features: (a) association matrix; (b) histogram distribution [66].

Figure 4.5 presents the analysis on the feature association distance in terms of standard deviation to evaluate the uncertainty on the relative position between the 39 most persistent features selected previously (see Figure 4.4). The estimation of the standard deviation only considers valid contributing measures, *i.e.*, only consider the distance between two features if both are visible in the same time instant. This estimation was possible by associating the features observable in each time instant to one of the 39 persistent ones if the error distance to the location manually selected is lower than 0.06 m. The latter threshold was needed to account both the unknown uncertainty of the location estimator and the noise on the laser measures (zero mean and 0.03 m standard deviation according to the official specifications for the Hokuyo UST-10LX sensor).

The maximum standard deviation of features association distance obtained in the experiments is 0.051 m. Although this value is greater than the one characteristic of the sensor, most of the associations have a non-zero standard deviation lower or equal to 0.03 m, representing 91.5% of the possible associations between features. Note that only non-zero values must be considered because, due to the sensor not having a 360° FOV and to possible physical obstructions on the laser view, not all associations are physically possible (in this case, 20 associations are not possible). Even when considering associations with a deviation lower or equal to 0.03 m, the proportion remains very significant at 59.3%. Thus, these results on the feature association distance underline the usefulness of the proposed corner-like features detector for localization processes, given that the geometric relations between features remain consistent and invariant.

4.3.4 Integration with a line-segment feature extraction algorithm

The final experiment in the scope of this chapter is shown in a video⁴, where our corner-like feature extractor and an open-source line-segment feature extraction algorithm¹ run simultaneously. In most cases, there is a complementarity between our corner-like features and line-segment features. Indeed, there are situations in which no corners are detected while line features are present. At the same time, other situations in which the expected lines are too small and are therefore possibly rejected by the line extractor, leaving only corners to be detected. Although the situations shown in the video are also dependent on the environment surrounding the robot, this complementarity between different types of features is advantageous for localization and mapping processes, thereby improving the robustness and stability of the robot pose estimate by increasing the number of features observable in the environment. Still, the main goal of this experiment is to show that the proposed feature extraction algorithm may be run in parallel with other algorithms, leveraging the ROS [96] modular architecture of having independent nodes subscribing to the same data, in this case, the range data from the 2D laser scanner sensor.

4.4 Summary and Conclusions

This chapter addresses a feature extraction algorithm for corner-like features from 2D lasers data. Although the most common way to represent the environment from 2D lasers is via 2D occupancy grid maps, an alternative is to represent the scene as a set of features. These features may include beacons[145], geometric cues (line [146–154] and curve [150–152, 155] segments, corners [146–152, 155–159], and free endpoints of segments [146, 151, 158, 159]), and semantic features [147, 156, 160]. In the specific case of corner-like features, the most common approach in the literature is to intersect line-segment features to extract corners from 2D lasers data [146–149, 155–158, 160–162]. However, this methodology makes corner detection highly dependent on the line-detection algorithm and may be limited to well-structured environments. Another common approach is to estimate curvature from the scan data [150–152, 159]. Still, the influence of sensor noise on curvature estimation may lead to variant features and false positives.

Given the previous limitations, we propose a feature-extraction algorithm to detect points of interest in 2D laser scanners data that resemble physical corners. The proposed algorithm performs local line fitting at multiple distance scales around a given point, similar to Speeded Up Robust Features (SURF) [170]. These scales define arrival and departure distance-based neighborhoods around the points in the scan. The intersection

⁴<https://www.youtube.com/watch?v=7xMKtlvd0wU>

of the fitted lines defines candidate feature keypoints. Then, a set of indicators based on line-fitting error, the angle between the arrival and departure lines, and consecutive observations of the same candidate keypoint across different scales defines a corner-like feature. Our methodology is independent of global search during line segment scanning, allowing other feature extraction methods to run concurrently.

Experimental results show that the proposed corner-like features correspond not only to well-defined 90° corners but may also detect other corner-like points of interest. More points of interest yield more features that may help stabilize the robot’s localization when estimating its pose using the proposed corner-like features. Furthermore, we analyze the observability ratio of the corner-like features throughout the experiment, based on the robot pose and sensor **FOV**. When considering the robot pose uncertainty and sensor, 6 features had a 100% observability ratio, while 20 and 16 features had a ratio greater or equal to 50% and 75%, respectively. These results reveal viewpoint invariance of the proposed feature extractor algorithm. As for geometric invariance, when analyzing the association distance between the 39 most persistent features from the experiments, the maximum standard deviation of these distances for each pair of associated features was 0.051 m. In contrast, 91.5% of possible feature associations had a standard deviation of 0.03 m or less, which is equivalent to the sensor noise. These results demonstrate that geometric relations between corner-like features extracted through our methodology remain consistent and geometrically invariant.

Overall, our corner-like features demonstrate their usefulness for localization and mapping processes through their viewpoint- and geometry-invariant properties in the experiments. These features can be extracted simultaneously with line segments from 2D laser scanner data, and combining different feature modalities can provide complementary observations, thereby improving the robustness and stability of localization. A natural continuation of this work is the development of a feature-based 2D localization estimator for mobile robots using 2D laser scan data, which could later be integrated as a front-end within a **SLAM** framework. Although feature-based localization and mapping constitute a promising research direction, it also requires additional components, such as data association algorithms, uncertainty models for feature observations, possibly feature descriptors, and specific map data structures. These requirements are not commonly supported in **ROS** [96], in which the most widely adopted representation of a 2D environment is the occupancy grid map rather than a feature-based map. Therefore, the next chapters of this PhD Thesis focus on the localization and mapping of mobile robots using wheeled odometry and raw 2D laser scanner data, rather than relying on feature-based environment representations.

Chapter 5

Point Cloud Registration with Distance Maps

This chapter presents a novel approach for 2D point cloud registration using distance maps, developed within the scope of this PhD Thesis in collaboration with Prof. Dr. Giorgio Grisetti from Sapienza Università di Roma. Distance maps encode the Euclidean distance to the nearest reference point at every cell of a precomputed discrete grid, decoupling data association from the iterative optimization loop. By replacing the explicit nearest-neighbour search performed at every iteration of classical **Iterative Closest Point (ICP)** [61, 171] with a constant-time grid lookup on distance map-based data structures, the proposed formulation provides both point-to-point and point-to-plane error formulations compatible with Gauss-Newton optimization on the $SE(2)$ manifold. The proposed approach does not require explicit data association at each iteration.

This chapter is organized as follows. Section 5.1 covers point cloud registration methods from the literature and the non-linear least-squares framework used throughout this work. Section 5.2 formulates the proposed approach, introducing the distance map representation, the registration problem on the $SE(2)$ manifold, and both point-to-point and point-to-plane error formulations with their Jacobians. Section 5.3 presents experimental results on synthetic and real-world data, evaluating registration accuracy and computational performance. Lastly, Section 5.4 summarizes the main conclusions and directions for future work.

5.1 Background

Point cloud registration is a fundamental problem in mobile robotics, consisting of estimating the rigid-body transformation that best aligns two point clouds acquired by a range sensor. This problem arises in a wide range of applications, including **SLAM**, odometry estimation, and map-based localization. Thus, this section introduces the background underpinning the proposed distance map-based registration approach. Section 5.1.1 reviews 2D and 3D point cloud registration methods from the literature, ranging from classical **ICP** variants to probabilistic and precomputed spatial gradient formulations, motivating the use of distance maps as a unified and correspondence-free registration framework. Section 5.1.2 presents the non-linear least-squares framework used throughout this chapter, covering the Gauss-Newton algorithm, uncertainty estimation, robust estimators, and the generalization to non-Euclidean state spaces.

5.1.1 Point cloud registration

One of the most popular approaches for point cloud registration is the **Iterative Closest Point (ICP)** algorithm. Besl and McKay [61] address registration of 3D shapes represented as point clouds, introducing the standard **ICP**

algorithm. The following two steps can summarize this algorithm: (i) compute correspondences between the source ($\mathbf{p}_i$) and target ($\mathbf{z}_i$) cloud points (closest nearest neighbor based on Euclidean distance); and (ii) estimate the rigid transformation of the source *w.r.t.* target cloud ($\mathbf{x}$) that minimizes a point-to-point distance measure between corresponding points (see Equation (5.1)). Chen and Medioni [171] exploit the tendency of most range data to be locally planar and introduce a point-to-plane variant of **ICP**. The proposed method exploits surface normal information on the target cloud ($\mathbf{n}_{z_i}$). Instead of minimizing the difference between the projected source and the target points, the point-to-plane error metric minimizes the squared dot product of the surface normal with the difference between the expected projected source and the reference target points (see Equation (5.2)). Indeed, Rusinkiewicz and Levoy [172] compared **ICP** variants with respect to point selection, data association, weighting of associated pairs, outlier rejection, and error metrics. The point-to-plane [171] metric demonstrated significantly better performance than the point-to-point [61] formulation in the experiments.

$$\mathbf{x} = \underset{\mathbf{x}}{\operatorname{argmin}} \sum_{i=1}^M \|\mathbf{x}\mathbf{p}_i - \mathbf{z}_i\|^2 \quad (5.1)$$

$$\mathbf{x} = \underset{\mathbf{x}}{\operatorname{argmin}} \sum_{i=1}^M \|\mathbf{n}_{z_i}^T \cdot (\mathbf{x}\mathbf{p}_i - \mathbf{z}_i)\|^2 \quad (5.2)$$

Unlike traditional **ICP** methods, which rely on explicit point correspondences, the **Normal Distribution Transform (NDT)** proposed by Biber and Strasser [82] performs 2D point cloud registration by representing the environment with local normal distributions. The 2D space is discretized into cells. Each cell is associated with a Gaussian distribution defined by the mean and covariance of the cell points, which provide information about the surface those points represent. Then, Biber and Strasser [82] formulate 2D point cloud registration as the minimization of an error function derived from these distributions, optimized using the Gauss-Newton algorithm. Experimental results showed that **NDT** [82] was able to map large areas without requiring odometry. Magnusson *et al.* [143] later extended **NDT** to 3D point cloud registration by subdividing space into voxels, each modelled as a normal distribution estimated from the points contained within it. The alignment is obtained by evaluating the likelihood of points from one scan under the distributions of the other scan. Compared with point-to-point **ICP** [61] using kd-tree nearest-neighbor search [173], 3D-**NDT** [143] achieves lower computation cost by avoiding explicit correspondence search, while also offering a more compact representation of 3D point cloud compared to raw data.

While leveraging **NDT** [82, 143]-like representations of the scene, **Generalized-ICP (GICP)** [174] and **NICP** [81] still retain the **ICP** structure for cloud registration. **GICP** [174] retains the standard **ICP** structure for correspondence estimation, typically using Euclidean closest-neighbor search accelerated with a kd-tree [173]. The main modification lies in the minimization objective function, introducing a plane-to-plane error metric. **GICP** models local surface information around each point, estimated from neighboring points via **PCA**. By accounting for the covariance of the matched points, **GICP** weights residuals based on local geometric uncertainty. Experiments showed that **GICP** [174] was more robust to incorrect correspondences and more accurate than standard point-to-point and point-to-plane **ICP** [61, 171]. **NICP** [81] extends **GICP** [174] not only in the optimization stage but also during correspondence search. Rather than relying purely on Euclidean nearest-neighbor matching, **NICP** [81] uses a line-of-sight-based association step and rejects candidate pairs when normals are not well defined, when point distance exceeds a threshold, when curvature mismatch is too large, or when the normals are insufficiently aligned. In the optimization step, **NICP** [81] augments each point

with its normal and minimizes a Mahalanobis distance in this 6D space. In the experiments, **NICP** [81] was more accurate than **GICP** [174] and **NDT** [143] registration algorithms on the evaluated datasets.

Furthermore, Montesano *et al.* [175] present a probabilistic scan matching method for 2D point cloud registration. The method [175] builds a set of possible correspondences using an individual compatibility test based on the Mahalanobis distance between pairs of points, while accounting for correspondence uncertainty. Then, Montesano *et al.* [175] maximize the probability of scan alignment over the points correspondence distribution. The algorithm [175] showed faster convergence rates than standard **ICP** [61]. Instead of selecting a single set of correspondences as Montesano *et al.* [175], Censi [176] considers all possible correspondences and generates alignment hypotheses through the **Generalized Hough Transform (GHT)**, while avoiding explicit point-to-point correspondences. Censi [176] computes alignment within a probabilistic framework that incorporates a motion model and estimates for registration uncertainty, accounting for surface normals.

Similarly, the **GMapping** [74] and **Cartographer** [177] **SLAM** algorithms use **MLE** probabilistic formulations to estimate the scan alignment *w.r.t.* the map. The scan matcher in **GMapping** [74] seeks the most likely robot pose by matching the current observation against the map built so far. The **MLE** formulation estimates the robot pose given the map, the current scan, and an initial pose guess. This formulation performs a gradient-descent search over the observation likelihood function defined on the grid map, based on the beam-endpoint sensor model [78]. **Cartographer** [177] optimizes the scan pose relative to the current local submap using a **Ceres** [97]-based scan matcher. This matcher estimates the pose that maximizes the occupancy probabilities of the scan points projected onto the submap. The problem is formulated as a non-linear least-squares optimization, using a smooth version of the occupancy probability values associated with the projected scan points.

Precomputed spatial gradients may be leveraged for the 2D point cloud registration without requiring explicit point-to-point correspondences. **HectorSLAM** [72] uses spatial gradients over the 2D occupancy grid map ($\mathcal{M}$). The scan matching finds the pose that maximizes the occupancy probability of where the beam endpoints $\mathbf{z}_i$ land on occupied cells (see Equation (5.3)). The method uses bilinear interpolation over the occupancy map and approximates its spatial derivatives, applying then a Gauss-Newton update. **PerfectMatch** [80, 178] focuses on a soccer field markings environment for robot soccer. This method formulates the 2D registration problem based on distance maps, using the **Residual Propagation (RProp)** [179] gradient-descent algorithm. Each point is evaluated by the distance map ($\mathcal{D}_{\mathcal{M}}$) that gives the distance to the closest field marking through a robust M-estimator (ρ) (see Equation (5.4)). The algorithm considers that the distance map is over the entire soccer field environment, only considering scan-to-map alignment. Although distance maps are continuous and thus differentiable, the authors approximate the final estimation covariance based on the second derivative of the error assuming an environment with orthogonal walls, equivalent to the line markings of soccer fields. Sobreira *et al.* [178] benchmarked in a simulation environment **PerfectMatch** [80] against **ICP** [61] and **NDT** [82]. **PerfectMatch** [80] was faster than both **ICP** [61] and **NDT** [82], given that **PerfectMatch** [80] uses lookup tables for the error function. Additionally, the **PerfectMatch** [80] algorithm demonstrated a higher tolerance to orientation initialization errors than **ICP** [61] and **NDT** [82].

$$\mathbf{x} = \underset{\mathbf{x}}{\operatorname{argmin}} \sum_{i=1}^M [1 - \mathcal{M}(\mathbf{x}\mathbf{z}_i)]^2 \quad (5.3)$$

$$\mathbf{x} = \underset{\mathbf{x}}{\operatorname{argmin}} \sum_{i=1}^M \rho(\mathcal{D}_{\mathcal{M}}(\mathbf{x}\mathbf{z}_i)) \quad (5.4)$$

Overall, most **ICP** methods rely on explicit point-to-point correspondences, requiring two steps per optimization iteration: data association and pairwise error computation between source and target points. As for probabilistic-based approaches, these methods typically require solving a **Maximum A Posteriori (MAP)** problem and having a map estimate of the environment over time, which is not compatible with scan-to-scan registration [174]. In contrast, distance maps should be compatible with both scan-to-scan and scan-to-map alignment, providing a nearest neighbor framework between source and target clouds. While the Perfect-Match [80] algorithm is specific to scan-to-map registration, an adaptation of the original approach should be compatible with scan-to-scan alignment. Moreover, distance maps provide a precomputed lookup table that directly allows evaluating the registration error of each source point without requiring an explicit data association step. This property allows faster computation during point cloud registration than **ICP** [61] and **NDT** [82], as showed in Sobreira *et al.* [178].

5.1.2 Least-squares minimization

State estimation in mobile robotics requires inferring an unknown state from noisy sensor measurements [62, 84]. First, Section 5.1.2.1 formulates the problem as a **MAP** inference problem and demonstrates how, under Gaussian noise assumptions, it reduces to a non-linear least-squares problem. Next, Section 5.1.2.2 shows how the Gauss-Newton algorithm solves the non-linear least-squares problem through a first-order approximation from the Taylor series. Section 5.1.2.3 presents an analysis of the uncertainty of the recovered state estimate based on the approximate Hessian matrix retrieved from the Gauss-Newton solver. Then, Section 5.1.2.4 extends the non-linear least-squares problem to use robust estimators to handle outliers. Finally, Section 5.1.2.5 generalizes the non-linear least-squares problems for non-Euclidean state spaces.

5.1.2.1 Problem formulation

Least-squares minimization arises in robotics when estimating an unknown state from noisy sensor measurements [62]. Let $\mathbf{x}$ denote non-observable state variables and $\mathbf{z}$ the measurements. In the scope of this PhD Thesis, the measurements in this chapter are the observations of the environment acquired by the robot's on-board 2D laser scanner. The state is assumed to follow a prior distribution $p(\mathbf{x})$, while the conditional distribution of the measurements given the state, $p(\mathbf{z}|\mathbf{x})$, is assumed known, commonly referred to as the *observation model*. Then, the goal is to estimate the most probable state given the measurements, *i.e.*, the posterior $p(\mathbf{x}|\mathbf{z})$. This problem represents a **MAP** estimation problem, as shown in Equation (5.5) [180].

$$\mathbf{x}^* = \underset{\mathbf{x}}{\operatorname{argmax}} p(\mathbf{x}|\mathbf{z}) \quad (5.5)$$

By applying Bayes' rule, shown in Equation (5.6), the posterior distribution $p(\mathbf{x}|\mathbf{z})$ becomes proportional to the likelihood of the measurements given the state $p(\mathbf{z}|\mathbf{x})$ multiplied by the prior $p(\mathbf{x})$. If a uniform prior over the admissible values of $\mathbf{x}$ is assumed, then $p(\mathbf{x})$ is constant and the posterior $p(\mathbf{x}|\mathbf{z})$ depends only on the measurement distribution conditioned by the state $p(\mathbf{z}|\mathbf{x})$.

$$p(\mathbf{x}|\mathbf{z}) = \frac{p(\mathbf{z}|\mathbf{x})p(\mathbf{x})}{p(\mathbf{z})} \propto p(\mathbf{z}|\mathbf{x})p(\mathbf{x}) \propto p(\mathbf{z}|\mathbf{x}) \quad (5.6)$$

Next, considering M measurements $\mathbf{z}^{[i]}$ that are conditionally independent given $\mathbf{x}$, the likelihood $p(\mathbf{x}|\mathbf{z})$ results in Equation (5.7). Under the assumption that each measurement follows a Gaussian distribution with mean $\mu_{z|x}^{[i]} = \mathbf{h}^{[i]}(\mathbf{x})$ (predicted measurement given the state $\mathbf{x}$) and covariance $\Sigma^{[i]} = \mathbf{\Omega}^{[i]-1}$ (where $\mathbf{\Omega}^{[i]}$ represents the information matrix of each measurement), the individual likelihoods $p(\mathbf{z}^{[i]}|\mathbf{x})$ result in Equation (5.8), based on the probability density function of normal distributions.

$$p(\mathbf{x}|\mathbf{z}) \propto \prod_{i=1}^M p(\mathbf{z}^{[i]}|\mathbf{x}) \quad (5.7)$$

$$\begin{aligned} p(\mathbf{z}^{[i]}|\mathbf{x}) &= \mathcal{N}(\mathbf{z}^{[i]}; \mathbf{h}^{[i]}(\mathbf{x}), \mathbf{\Omega}^{[i]-1}) \\ &\propto \exp\left(-\frac{1}{2} \left(\mathbf{h}^{[i]}(\mathbf{x}) - \mathbf{z}^{[i]}\right)^T \mathbf{\Omega}^{[i]} \left(\mathbf{h}^{[i]}(\mathbf{x}) - \mathbf{z}^{[i]}\right)\right) \end{aligned} \quad (5.8)$$

Furthermore, Equation (5.9) defines the residual error $\mathbf{e}^{[i]}(\mathbf{x})$ for each measurement. The squared Mahalanobis distance $\chi^{2[i]}(\mathbf{x})$ of the residual error is formulated in Equation (5.10). Thus, Equation (5.11) illustrates the proportionality of the individual likelihoods $p(\mathbf{z}^{[i]}|\mathbf{x})$ over the squared Mahalanobis distance $\chi^{2[i]}(\mathbf{x})$.

$$\mathbf{e}^{[i]}(\mathbf{x}) = \mathbf{h}^{[i]}(\mathbf{x}) - \mathbf{z}^{[i]} \quad (5.9)$$

$$\chi^{2[i]}(\mathbf{x}) = \mathbf{e}^{[i]}(\mathbf{x})^T \mathbf{\Omega}^{[i]} \mathbf{e}^{[i]}(\mathbf{x}) \quad (5.10)$$

$$\begin{aligned} p(\mathbf{z}^{[i]}|\mathbf{x}) &\propto \exp\left(-\frac{1}{2} \mathbf{e}^{[i]}(\mathbf{x})^T \mathbf{\Omega}^{[i]} \mathbf{e}^{[i]}(\mathbf{x})\right) \\ &\propto \exp\left(-\frac{1}{2} \chi^{2[i]}(\mathbf{x})\right) \end{aligned} \quad (5.11)$$

Under the assumption of having a uniform prior, the original **MAP** problem formulated in Equation (5.5) can be rewritten as an **MLE** problem that maximizes the probability of the measurements given the state $p(\mathbf{z}^{[i]}|\mathbf{x})$ (see Equation (5.12)). Through the negative log-likelihood (see Definition 2.1.4), dropping constant factors, and assuming the measurements under Gaussian noise, the original **MAP** problem for robotics state estimation reduces to a **Maximum Likelihood Estimation (MLE)** problem, which in turn reduces to minimizing the sum of squared Mahalanobis residuals. This problem is equivalent to a weighted non-linear least-squares problem, as formalized in Equation (5.12).

$$\begin{aligned} \mathbf{x}^* &= \operatorname{argmax}_{\mathbf{x}} \prod_{i=1}^M p(\mathbf{z}^{[i]}|\mathbf{x}) \\ &= \operatorname{argmax}_{\mathbf{x}} \prod_{i=1}^M \exp\left(-\frac{1}{2} \mathbf{e}^{[i]}(\mathbf{x})^T \mathbf{\Omega}^{[i]} \mathbf{e}^{[i]}(\mathbf{x})\right) \\ &= \operatorname{argmin}_{\mathbf{x}} \sum_{i=1}^M \mathbf{e}^{[i]}(\mathbf{x})^T \mathbf{\Omega}^{[i]} \mathbf{e}^{[i]}(\mathbf{x}) \\ &= \operatorname{argmin}_{\mathbf{x}} \sum_{i=1}^M \chi^{2[i]}(\mathbf{x}) \end{aligned} \quad (5.12)$$

5.1.2.2 Gauss-Newton algorithm

Following the non-linear least-squares formulation commonly used in robotics state estimation, one possible solution is an iterative minimization of Equation (5.13) [62]. At each iteration k , the state is updated by a perturbation $\Delta \mathbf{x}_k$, as shown in Equation (5.14).

$$F(\mathbf{x}_k) = \sum_{i=1}^M \mathbf{e}^{[i]}(\mathbf{x}_k)^T \boldsymbol{\Omega}^{[i]} \mathbf{e}^{[i]}(\mathbf{x}_k) \quad (5.13)$$

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \Delta \mathbf{x}_k \quad (5.14)$$

In order to compute $\Delta \mathbf{x}_k$, each residual is linearized around the current estimate $\mathbf{x}_k$ via a first-order Taylor approximation, as formulated in Equation (5.15) [180, 181], where $\mathbf{J}^{[i]} \in \mathbb{R}^{m_i \times n}$ is the Jacobian of the i -th residual evaluated at $\mathbf{x}_k$, with m_i the dimension of the residual i and n the dimension of the state.

$$\mathbf{e}^{[i]}(\mathbf{x}_k + \Delta \mathbf{x}_k) \simeq \underbrace{\mathbf{e}^{[i]}(\mathbf{x}_k)}_{\mathbf{e}^{[i]}} + \underbrace{\frac{\partial \mathbf{e}^{[i]}(\mathbf{x})}{\partial \mathbf{x}} \bigg|_{\mathbf{x}=\mathbf{x}_k}}_{\mathbf{J}^{[i]}} \Delta \mathbf{x}_k \quad (5.15)$$

Next, Equation (5.16) substitutes the first-order Taylor approximation into the objective function and expands the products in Equation (5.13), where the scalar symmetry $\Delta \mathbf{x}^T \mathbf{J}^{[i]T} \boldsymbol{\Omega}^{[i]} \mathbf{e}^{[i]} = \mathbf{e}^{[i]T} \boldsymbol{\Omega}^{[i]} \mathbf{J}^{[i]} \Delta \mathbf{x}$ has been used to combine the two cross-terms. Also, note that $\mathbf{b}^{[i]} = \mathbf{J}^{[i]T} \boldsymbol{\Omega}^{[i]} \mathbf{e}^{[i]} \in \mathbb{R}^n$ is defined as a column vector. The notation $\mathbf{b}^{[i]T} \Delta \mathbf{x}$ in the expansion below reflects the transpose of that column vector.

$$\begin{aligned} \mathbf{e}^{[i]}(\mathbf{x}_k + \Delta \mathbf{x}_k)^T \boldsymbol{\Omega}^{[i]} \mathbf{e}^{[i]}(\mathbf{x}_k + \Delta \mathbf{x}_k) &\simeq \left(\mathbf{e}^{[i]} + \mathbf{J}^{[i]} \Delta \mathbf{x} \right)^T \boldsymbol{\Omega}^{[i]} \left(\mathbf{e}^{[i]} + \mathbf{J}^{[i]} \Delta \mathbf{x} \right) \\ &= \Delta \mathbf{x}^T \underbrace{\mathbf{J}^{[i]T} \boldsymbol{\Omega}^{[i]} \mathbf{J}^{[i]}}_{\mathbf{H}^{[i]}} \Delta \mathbf{x} + 2 \underbrace{\mathbf{e}^{[i]T} \boldsymbol{\Omega}^{[i]} \mathbf{J}^{[i]}}_{\mathbf{b}^{[i]T}} \Delta \mathbf{x} + \underbrace{\mathbf{e}^{[i]T} \boldsymbol{\Omega}^{[i]} \mathbf{e}^{[i]}}_{c^{[i]}} \end{aligned} \quad (5.16)$$

Summing over all measurements and defining $\mathbf{H} = \sum_{i=1}^M \mathbf{H}^{[i]} \in \mathbb{R}^{n \times n}$, $\mathbf{b} = \sum_{i=1}^M \mathbf{b}^{[i]} \in \mathbb{R}^n$, and $c = \sum_{i=1}^M c^{[i]} \in \mathbb{R}$, Equation (5.17) formulates the linearized cost derived in Equation (5.16) to a quadratic form in $\Delta \mathbf{x}$. The information matrices of the measurements $\boldsymbol{\Omega}^{[i]}$ are symmetric positive semi-definite. Consequently, the matrix $\mathbf{H}$ inherits the same property by construction, where $\mathbf{H}^T = \mathbf{H}$ [62].

$$F(\mathbf{x}_k + \Delta \mathbf{x}_k) \simeq \Delta \mathbf{x}_k^T \mathbf{H} \Delta \mathbf{x}_k + 2 \mathbf{b}^T \Delta \mathbf{x}_k + c \quad (5.17)$$

The optimal perturbation minimizes the quadratic approximation presented in Equation (5.17). Setting the gradient w.r.t. $\Delta \mathbf{x}$ to zero, as shown in Equation (5.18), yields the normal equations presented in Equation (5.19), whose solution $\Delta \mathbf{x}_k$ is then used to update the state via Equation (5.14), constituting one iteration of the *Gauss-Newton* method. The method linearizes the residual error functions around the current estimate and minimizes the resulting quadratic approximation of the weighted non-linear least-squares cost [62]. Algorithm 5.1 details the steps for a single iteration of the minimization procedure.

$$\mathbf{0} = \frac{\partial}{\partial \Delta \mathbf{x}} [\Delta \mathbf{x}_k^T \mathbf{H} \Delta \mathbf{x}_k + 2 \mathbf{b}^T \Delta \mathbf{x}_k + c] = 2 \mathbf{H} \Delta \mathbf{x}_k + 2 \mathbf{b} \quad (5.18)$$

Algorithm 5.1: Gauss-Newton minimization algorithm for Euclidean measurement and state spaces
(a single iteration of the iterative procedure)

Input: Current state estimation $\mathbf{x}_k$, measurements $C = \{\langle \mathbf{z}^{[i]}, \mathbf{\Omega}^{[i]} \rangle\}$
Output: State estimation for the next iteration step $\mathbf{x}_{k+1}$

```

1  $\mathbf{b} \leftarrow \mathbf{0}_{n \times 1}$ 
2  $\mathbf{H} \leftarrow \mathbf{0}_{n \times n}$ 
3 forall  $i = 1, \dots, M$  do
4    $\mathbf{e}^{[i]} \leftarrow \mathbf{h}^{[i]}(\mathbf{x}_k) - \mathbf{z}^{[i]}$ 
5    $\mathbf{J}^{[i]} \leftarrow \left. \frac{\partial \mathbf{h}^{[i]}(\mathbf{x})}{\partial \mathbf{x}} \right|_{\mathbf{x}=\mathbf{x}_k}$ 
6    $\mathbf{H} \leftarrow \mathbf{H} + \mathbf{J}^{[i]T} \mathbf{\Omega}^{[i]} \mathbf{J}^{[i]}$ 
7    $\mathbf{b} \leftarrow \mathbf{b} + \mathbf{J}^{[i]T} \mathbf{\Omega}^{[i]} \mathbf{e}^{[i]}$ 
8  $\Delta \mathbf{x}_k \leftarrow \text{solve}(\mathbf{H} \Delta \mathbf{x}_k = -\mathbf{b})$ 
9  $\mathbf{x}_{k+1} \leftarrow \mathbf{x}_k + \Delta \mathbf{x}_k$ 
10 return  $\mathbf{x}_{k+1}$ 

```

$$\mathbf{H} \Delta \mathbf{x}_k = -\mathbf{b} \quad (5.19)$$

5.1.2.3 Least-squares and uncertainty

At convergence, the approximate Hessian matrix $\mathbf{H}$ of the Gauss-Newton algorithm provides an estimate of the uncertainty for the recovered state. Consider a small perturbation $\Delta \mathbf{x}$ away from the optimum $\mathbf{x}^*$. Under the first-order linearization of each measurement function $\mathbf{h}^{[i]}$, the estimated measurement $\mathbf{h}^{[i]}(\mathbf{x}^* + \Delta \mathbf{x})$ around the optimum can be approximated by Equation (5.20), where $\mathbf{J}^{[i]} \in \mathbb{R}^{m_i \times n}$ is the Jacobian of the i -th measurement function evaluated at the optimum, m_i is the dimension of the i -th measurement, and n is the dimension of the state. This linearization is valid as long as the measurement function is sufficiently smooth and $\|\Delta \mathbf{x}\|$ is small.

$$\mathbf{h}^{[i]}(\mathbf{x}^* + \Delta \mathbf{x}) \simeq \mathbf{h}^{[i]}(\mathbf{x}^*) + \underbrace{\left. \frac{\partial \mathbf{h}^{[i]}(\mathbf{x})}{\partial \mathbf{x}} \right|_{\mathbf{x}=\mathbf{x}^*}}_{\mathbf{J}^{[i]}} \Delta \mathbf{x} \quad (5.20)$$

Next, Equation (5.21) represents the conditional probability of the i -th measurement given the perturbed state, where $\mathbf{z}^{[i]*} = \mathbf{h}^{[i]}(\mathbf{x}^*)$ is the predicted measurement at the optimum.

$$p(\mathbf{z}^{[i]} | \mathbf{x}^* + \Delta \mathbf{x}) = \mathcal{N}(\mathbf{z}^{[i]}; \mu_{z|x}^{[i]} = \mathbf{z}^{[i]*} + \mathbf{J}^{[i]} \Delta \mathbf{x}, \mathbf{\Omega}^{[i]-1}) \quad (5.21)$$

Since measurements are conditionally independent given the state, the joint conditional likelihood $p(\mathbf{z} | \mathbf{x}^* + \Delta \mathbf{x})$ over all M measurements factorizes into a product of Gaussian distributions of each individual measurement. As a result, the joint conditional likelihood $p(\mathbf{z} | \mathbf{x}^* + \Delta \mathbf{x})$ formulated in Equation (5.22) is also a Gaussian over the stacked mean and the block-diagonal information matrix, where $\mathbf{J} = [\mathbf{J}^{[1]T}, \dots, \mathbf{J}^{[M]T}]^T \in \mathbb{R}^{m \times n}$ is the stacked Jacobian, $m = \sum_{i=1}^M m_i$, and $\mathbf{\Omega}_z = \text{diag}(\mathbf{\Omega}^{[1]}, \dots, \mathbf{\Omega}^{[M]})$ is the block-diagonal information matrix reflecting the conditional independence of the measurements.

$$p(\mathbf{z} | \mathbf{x}^* + \Delta \mathbf{x}) = \mathcal{N}(\mathbf{z}; \mu_{z|x} = \mathbf{z}^* + \mathbf{J} \Delta \mathbf{x}, \mathbf{\Omega}_z^{-1}) \quad (5.22)$$

Assuming an uninformative (uniform) prior on $\Delta \mathbf{x}$ around the optimum $\mathbf{x}^*$, this corresponds to a zero information matrix with a zero mean for the $\Delta \mathbf{x}$, as shown in Equation (5.23). The distribution over the state perturbation $\Delta \mathbf{x}$ is not a separate variable, given that $\Delta \mathbf{x} = \mathbf{x} - \mathbf{x}^*$. Applying the Gaussian chain rule [180], Equation (5.24) defines the information matrix $\mathbf{\Omega}_{x,z}$ for the joint distribution over $(\Delta \mathbf{x}, \mathbf{z})$. The prior contributes no information to the top-left block, so $\mathbf{\Omega}_{xx} = \mathbf{J}^T \mathbf{\Omega}_z \mathbf{J}$.

$$p(\Delta \mathbf{x}) = \mathcal{N}(\Delta \mathbf{x}; \mathbf{0}, \mathbf{\Omega}_x^{-1}), \quad \mathbf{\Omega}_x = \mathbf{0} \quad (5.23)$$

$$\mathbf{\Omega}_{x,z} = \begin{pmatrix} \mathbf{\Omega}_x + \mathbf{J}^T \mathbf{\Omega}_z \mathbf{J} & \mathbf{J}^T \mathbf{\Omega}_z \\ \mathbf{\Omega}_z \mathbf{J} & \mathbf{\Omega}_z \end{pmatrix} = \begin{pmatrix} \mathbf{J}^T \mathbf{\Omega}_z \mathbf{J} & \mathbf{J}^T \mathbf{\Omega}_z \\ \mathbf{\Omega}_z \mathbf{J} & \mathbf{\Omega}_z \end{pmatrix} = \begin{pmatrix} \mathbf{\Omega}_{xx} & \mathbf{\Omega}_{xz} \\ \mathbf{\Omega}_{zx} & \mathbf{\Omega}_{zz} \end{pmatrix} \quad (5.24)$$

When conditioning the state posterior $\mathbf{x}$ on the measurements $\mathbf{z}$ from the joint probability and under the Gaussian assumption, Schön and Lindsten [182] demonstrate that the mean and the covariance matrix for the state posterior given the measurements can be formulated by Equation (5.25) and Equation (5.26), respectively. Since at the optimum state the residual $\mathbf{z} - \mathbf{z}^* \simeq \mathbf{0}$, the correction term vanishes in the mean of the state posterior given the measurements.

$$\begin{aligned} \mu_{x|z} &= \mu_x - \mathbf{\Omega}_{xx}^{-1} \mathbf{\Omega}_{xz} (\mathbf{z} - \mu_z) \\ &= \mathbf{x}^* - \mathbf{\Omega}_{xx}^{-1} \mathbf{\Omega}_{xz} (\mathbf{z} - \mathbf{z}^*) \\ &\simeq \mathbf{x}^* \end{aligned} \quad (5.25)$$

$$\begin{aligned} \Sigma_{x|z} &= \Sigma_{xx} - \Sigma_{xz} \Sigma_{zz}^{-1} \Sigma_{zx} \\ &= \mathbf{\Omega}_{xx}^{-1} \end{aligned} \quad (5.26)$$

Therefore, the covariance of the optimized state $\mathbf{x}^*$ is estimated by Equation (5.27), where $\mathbf{\Omega}_{xx} = \mathbf{J}^T \mathbf{\Omega}_z \mathbf{J} = \mathbf{H}$ is the approximate Hessian recovered from the Gauss-Newton solver. The approximation sign reflects the fact that $\mathbf{H}$ is itself a first-order approximation to the true Hessian of the cost function, obtained by discarding the second-order residual terms in the linearization. The diagonal entries of $\Sigma_{x|z}$ represent the marginal variances of each state component, and the off-diagonal entries encode the correlations induced by the measurement model. This result requires $\mathbf{H}$ to have full column rank, *i.e.*, the stacked Jacobian $\mathbf{J}$ must be full column rank. When this condition does not hold (*e.g.*, when some state components are unobservable given the available measurements), $\mathbf{H}$ is singular and the covariance is undefined for those directions [180].

$$\Sigma_{x|z} = \mathbf{\Omega}_{xx}^{-1} = (\mathbf{J}^T \mathbf{\Omega}_z \mathbf{J})^{-1} \approx \mathbf{H}^{-1} \quad (5.27)$$

5.1.2.4 Robust estimators

The weighted non-linear least-squares algorithm formulated in Equation (5.12) assumes that all measurements follow a Gaussian noise model. Under this assumption, the cost $F(\mathbf{x}) = \sum_{i=1}^M \chi^{2[i]}(\mathbf{x})$ grows quadratically with the residual error $\mathbf{e}^{[i]}(\mathbf{x})$, meaning that a single outlying measurement with a large $\chi^{2[i]}$ can dominate the sum and severely corrupt the optimum. M-estimators address this limitation by replacing the quadratic penalty with a robust loss function $\rho(\cdot)$ that down-weights large residuals, resulting in the robust cost $F_\rho(\mathbf{x})$ presented in

Equation (5.28), where $s^{[i]}$ is the scalar Mahalanobis distance of the i -th measurement [183].

$$F_\rho(\mathbf{x}) = \sum_{i=1}^M \rho\left(s^{[i]}(\mathbf{x})\right), \quad s^{[i]}(\mathbf{x}) = \sqrt{\chi^{2[i]}(\mathbf{x})} \quad (5.28)$$

The standard least-squares cost $F(\mathbf{x})$ is recovered as a special case of $F_\rho(\mathbf{x})$ when $\rho(s) = s^2$. Consider the gradient of the standard least-squares with respect to $\mathbf{x}$ presented in Equation (5.29), using the identity $\frac{\partial}{\partial \mathbf{x}}(\mathbf{e}^{[i]}(\mathbf{x})^T \mathbf{\Omega}^{[i]} \mathbf{e}^{[i]}(\mathbf{x})) = 2 \cdot \mathbf{J}^{[i]T} \mathbf{\Omega}^{[i]} \mathbf{e}^{[i]}$ because $\mathbf{\Omega}^{[i]}$ is symmetric.

$$\begin{aligned} \nabla_{\mathbf{x}} F &= \sum_{i=1}^M \frac{\partial}{\partial \mathbf{x}} \left(\mathbf{e}^{[i]}(\mathbf{x})^T \mathbf{\Omega}^{[i]} \mathbf{e}^{[i]}(\mathbf{x}) \right) = \sum_{i=1}^M 2 \cdot \mathbf{J}^{[i]T} \mathbf{\Omega}^{[i]} \mathbf{e}^{[i]} \\ &= \sum_{i=1}^M \nabla_{\mathbf{x}} F^{[i]} \end{aligned} \quad (5.29)$$

When expanding the gradient of the robust cost $F_\rho(\mathbf{x})$ presented in Equation (5.30), the derivation chain rule is applied using the gradient identity from Equation (5.29), and the influence of the robust estimator ρ appears as a scalar weight w of each measurement factor in $F_\rho(\mathbf{x})$ when compared to the gradient of the standard cost $F(\mathbf{x})$. This influence results in a re-weighted least-squares problem, introducing the effective information matrix for each measurement factor $\tilde{\mathbf{\Omega}}^{[i]}(\mathbf{x}_k)$ formulated in Equation (5.31).

$$\begin{aligned} \nabla_{\mathbf{x}} F_\rho &= \sum_{i=1}^M \frac{\partial}{\partial \mathbf{x}} \left(\rho\left(s^{[i]}(\mathbf{x})\right) \right) = \sum_{i=1}^M \rho'\left(s^{[i]}(\mathbf{x})\right) \cdot \frac{\partial \left(\sqrt{\chi^{2[i]}(\mathbf{x})} \right)}{\partial \mathbf{x}} = \sum_{i=1}^M \frac{\rho'\left(s^{[i]}(\mathbf{x})\right)}{2 \cdot s^{[i]}(\mathbf{x})} \cdot \nabla_{\mathbf{x}} F^{[i]} \\ &= \sum_{i=1}^M w(s^{[i]}(\mathbf{x})) \cdot \nabla_{\mathbf{x}} F^{[i]}, \quad w(s) = \frac{\rho'(s)}{2s} \end{aligned} \quad (5.30)$$

$$\tilde{\mathbf{\Omega}}^{[i]}(\mathbf{x}_k) = w^{[i]}(\mathbf{x}_k) \mathbf{\Omega}^{[i]}, \quad (5.31)$$

Subsequently, Equation (5.32) defines the robust approximate Hessian matrix $\tilde{\mathbf{H}}$ and gradient vector $\tilde{\mathbf{b}}$ for the Gauss-Newton method to solve least-squares problems with robust estimators. The only modification with respect to the standard formulation is the substitution $\mathbf{\Omega}^{[i]} \leftarrow \tilde{\mathbf{\Omega}}^{[i]}$, so any sparse linear solver used for the standard Gauss-Newton step can be reused without modification. Algorithm 5.2 details the steps for a single iteration of the Gauss-Newton minimization procedure with robust estimators.

$$\tilde{\mathbf{H}} = \sum_{i=1}^M \mathbf{J}^{[i]T} \tilde{\mathbf{\Omega}}^{[i]} \mathbf{J}^{[i]}, \quad \tilde{\mathbf{b}} = \sum_{i=1}^M \mathbf{J}^{[i]T} \tilde{\mathbf{\Omega}}^{[i]} \mathbf{e}^{[i]}(\mathbf{x}_k). \quad (5.32)$$

At convergence, the robust approximate Hessian $\tilde{\mathbf{H}}$ provides a covariance estimate analogous to Equation (5.27), as presented in Equation (5.33). Given that outlier factors receive $w^{[i]} \approx 0$, they contribute negligibly to $\tilde{\mathbf{H}}$, not inflating the reported uncertainty of the inlier-driven solution [183].

$$\tilde{\Sigma}_{\mathbf{x}|\mathbf{z}} \approx \tilde{\mathbf{H}}^{-1}. \quad (5.33)$$

Lastly, Table 5.1 summarizes popular choices for the robust loss function ρ , together with their first-order

Algorithm 5.2: Gauss-Newton minimization algorithm with robust estimators for Euclidean measurement and state spaces (a single iteration of the iterative procedure)

Input: Current state estimation $\mathbf{x}_k$, measurements $C = \{\langle \mathbf{z}^{[i]}, \mathbf{\Omega}^{[i]} \rangle\}$, robust loss function ρ

Output: State estimation for the next iteration step $\mathbf{x}_{k+1}$

```

1  $\tilde{\mathbf{b}} \leftarrow \mathbf{0}_{n \times 1}$ 
2  $\tilde{\mathbf{H}} \leftarrow \mathbf{0}_{n \times n}$ 
3 forall  $i = 1, \dots, M$  do
4    $\mathbf{e}^{[i]} \leftarrow \mathbf{h}^{[i]}(\mathbf{x}_k) - \mathbf{z}^{[i]}$ 
5    $\mathbf{J}^{[i]} \leftarrow \left. \frac{\partial \mathbf{e}^{[i]}(\mathbf{x})}{\partial \mathbf{x}} \right|_{\mathbf{x}=\mathbf{x}_k}$ 
6    $s^{[i]} \leftarrow \sqrt{\mathbf{e}^{[i]T} \mathbf{\Omega}^{[i]} \mathbf{e}^{[i]}}$ 
7    $w^{[i]} \leftarrow \frac{\rho'(s^{[i]})}{2 \cdot s^{[i]}}$ 
8    $\tilde{\mathbf{\Omega}}^{[i]} \leftarrow w^{[i]} \mathbf{\Omega}^{[i]}$ 
9    $\tilde{\mathbf{H}} \leftarrow \tilde{\mathbf{H}} + \mathbf{J}^{[i]T} \tilde{\mathbf{\Omega}}^{[i]} \mathbf{J}^{[i]}$ 
10   $\tilde{\mathbf{b}} \leftarrow \tilde{\mathbf{b}} + \mathbf{J}^{[i]T} \tilde{\mathbf{\Omega}}^{[i]} \mathbf{e}^{[i]}$ 
11  $\Delta \mathbf{x}_k \leftarrow \text{solve}(\tilde{\mathbf{H}} \Delta \mathbf{x}_k = -\tilde{\mathbf{b}})$ 
12  $\mathbf{x}_{k+1} \leftarrow \mathbf{x}_k + \Delta \mathbf{x}_k$ 
13 return  $\mathbf{x}_{k+1}$ 

```

Table 5.1: Common M-estimator loss functions $\rho(s)$, their first-order derivative $\psi(s) = \rho'(s)$, and weight functions $w(s) = \psi(s)/(2s)$. The tuning constant c controls the inlier/outlier threshold [184].

Estimator	$\rho(s)$	$\psi(s)$	$w(s)$
L^2	s^2	$2s$	1
L^1	$ s $	$s/ s $	$1/(2 s)$
L^2_{trunc}	$\begin{cases} s^2 & \text{if } s < c \\ c^2 & \text{otherwise} \end{cases}$	$\begin{cases} 2s & \text{if } s < c \\ 0 & \text{otherwise} \end{cases}$	$\begin{cases} 1 & \text{if } s < c \\ 0 & \text{otherwise} \end{cases}$
L^1_{trunc}	$\begin{cases} s & \text{if } s < c \\ c & \text{otherwise} \end{cases}$	$\begin{cases} s/ s & \text{if } s < c \\ 0 & \text{otherwise} \end{cases}$	$\begin{cases} 1/(2 s) & \text{if } s < c \\ 0 & \text{otherwise} \end{cases}$
Huber	$\begin{cases} s^2 & \text{if } s < c \\ 2c s - c^2 & \text{otherwise} \end{cases}$	$\begin{cases} 2s & \text{if } s < c \\ 2c \cdot s/ s & \text{otherwise} \end{cases}$	$\begin{cases} 1 & \text{if } s < c \\ c/ s & \text{otherwise} \end{cases}$
Cauchy	$c^2 \ln\left(1 + \frac{s^2}{c^2}\right)$	$\frac{2s}{1 + (s/c)^2}$	$\frac{1}{1 + (s/c)^2}$
Tukey	$\begin{cases} c^2 \left[1 - \left(1 - \frac{s^2}{c^2}\right)^3\right] & \text{if } s < c \\ c^2 & \text{otherwise} \end{cases}$	$\begin{cases} 6s \left(1 - \frac{s^2}{c^2}\right)^2 & \text{if } s < c \\ 0 & \text{otherwise} \end{cases}$	$\begin{cases} 3 \left(1 - \frac{s^2}{c^2}\right)^2 & \text{if } s < c \\ 0 & \text{otherwise} \end{cases}$
Geman-McClure	$\frac{s^2}{c^2 + s^2}$	$2s \cdot \frac{c^2}{(c^2 + s^2)^2}$	$\frac{c^2}{(c^2 + s^2)^2}$

derivative $\psi(s) = \rho'(s)$ and weight function $w(s) = \psi(s)/(2s)$ [184]. The tuning constant c controls the threshold between the inlier region, where the loss behaves near-quadratically, and the outlier region, where large residuals are down-weighted.

5.1.2.5 Non-Euclidean spaces

State estimation in robotics frequently involves state spaces that are not flat Euclidean spaces. For example, robot poses include orientations that are elements of the special orthogonal group $SO(n)$ ($SO(2)$ and $SO(3)$ for 2D and 3D rotation matrices), or full rigid-body transformations that are elements of the special Euclidean group $SE(n)$, combining rotations and translations in $\mathbb{R}^n$ ($SE(2)$ and $SE(3)$ for 2D and 3D poses). These groups are smooth manifolds that do not support standard Euclidean addition and subtraction of state variables. Thus, the geometry of the state space must be accounted for explicitly during optimization [62, 180].

A smooth manifold $\mathcal{M}$ is a space that locally resembles a Euclidean space $\mathbb{R}^n$. At each point $\mathbf{x} \in \mathcal{M}$, the *tangent space* $\mathcal{T}_{\mathbf{x}}(\mathcal{M}) \cong \mathbb{R}^n$ provides the appropriate linear structure for computing increments, required for iterative least-squares minimization. A small perturbation $\Delta\mathbf{x} \in \mathbb{R}^n$ expressed in the tangent space is mapped back onto the manifold through the *exponential mapping*, as formulated in Equation (5.34). The operator $\oplus$ replaces standard Euclidean addition throughout the optimization, ensuring that every iteration remains on $\mathcal{M}$, while preserving geometric constraints intrinsic to the state representation (*e.g.*, the orthogonality constraint on rotation matrices, or the unit-norm constraint on quaternions). The inverse operation maps an element $\mathbf{y} \in \mathcal{M}$ back to the tangent space at $\mathbf{x}$ through the *logarithmic map*, as formulated in Equation (5.35) [180]. In the Euclidean case ($\mathcal{M} = \mathbb{R}^n$), the operators $\oplus$ and $\ominus$ reduce to standard vector addition and subtraction operations.

$$\mathbf{x} \oplus \Delta\mathbf{x} \triangleq \exp_{\mathbf{x}}(\Delta\mathbf{x}) \in \mathcal{M} \quad (5.34)$$

$$\mathbf{y} \ominus \mathbf{x} \triangleq \log_{\mathbf{x}}(\mathbf{y}) \in \mathcal{T}_{\mathbf{x}}(\mathcal{M}) \quad (5.35)$$

Given the current state estimate $\mathbf{x}_k \in \mathcal{M}$ and a tangent-space perturbation $\Delta\mathbf{x} \in \mathbb{R}^n$, the first-order Taylor expansion of the residual is formulated by Equation (5.36). Equation (5.37) represents the *logarithmic map* of the predicted measurement $\mathbf{h}^{[i]}(\mathbf{x}_k)$ into the tangent space at the reference measurement $\mathbf{z}^{[i]}$. The *manifold Jacobian* (see Equation (5.38)) is the derivative of the residual *w.r.t.* $\Delta\mathbf{x} \in \mathbb{R}^n$, evaluated at the origin of the tangent space (equivalent to $\Delta\mathbf{x} = 0$) [62]. The choice of left versus right-sided perturbation update in Lie groups affects the form of the Jacobian. Hence, *exponential* and *logarithmic map* operations should be consistent throughout the implementations.

$$\mathbf{e}^{[i]}(\mathbf{x}_k \oplus \Delta\mathbf{x}) \approx \mathbf{e}^{[i]}(\mathbf{x}_k) + \mathbf{J}^{[i]}(\mathbf{x}_k)\Delta\mathbf{x} \quad (5.36)$$

$$\mathbf{e}^{[i]}(\mathbf{x}_k) = \mathbf{h}^{[i]}(\mathbf{x}_k) \ominus \mathbf{z}^{[i]} \in \mathcal{T}_{\mathbf{z}^{[i]}}(\mathcal{M}) \quad (5.37)$$

$$\mathbf{J}^{[i]}(\mathbf{x}) = \left. \frac{\partial \mathbf{e}^{[i]}(\mathbf{x} \oplus \Delta\mathbf{x})}{\partial \Delta\mathbf{x}} \right|_{\Delta\mathbf{x}=0} \in \mathbb{R}^{m_i \times n} \quad (5.38)$$

As a result, the definition of the approximate Hessian and gradient vectors based on the Jacobian computed *w.r.t.* the tangent-space perturbation rather than a Euclidean increment (see Equation (5.38)) and the error from the *logarithmic map* of the measurement and prediction estimation make the iterative least-squares on manifolds similar to Algorithm 5.1 and Algorithm 5.2, differing only in how $\mathbf{H}$ and $\mathbf{b}$ are computed. Once $\Delta\mathbf{x}^*$ has been computed in the tangent space through standard linear solvers as in the original Gauss-Newton method, the state is retracted back onto the manifold via the *exponential map*, as formulated in Equation (5.39).

$$\mathbf{x}_{k+1} = \mathbf{x}_k \oplus \Delta\mathbf{x}^* = \exp_{\mathbf{x}_k} \left(\text{solve} \left(-\mathbf{H}_k^{-1} \mathbf{b}_k \right) \right) \quad (5.39)$$

5.2 Registration Methodology

This section presents the proposed approach for 2D point cloud registration using distance maps, formulated as a non-linear least-squares problem and solved via the Gauss-Newton algorithm on the $SE(2)$ manifold. First, Section 5.2.1 introduces the distance map representation, its spatial derivatives, the associated representation levels, and the data structures used for computation and querying. Next, Section 5.2.2 states the registration problem, defining the pose parametrization and the manifold update. Then, Section 5.2.3 derives the point-to-point error formulation and its Jacobian matrix via right-sided perturbation on $SE(2)$. Lastly, Section 5.2.4 derives the point-to-plane error formulation and its Jacobian, together with the surface normal estimation procedure used to extract normals from the moving point cloud.

5.2.1 Distance maps

Distance maps, also referred in the literature as distance fields or transforms, are scalar fields that associate to every cell of a discrete grid the Euclidean distance to the nearest element of a given reference set. In the context of point cloud registration, the reference set is the target point cloud and the distance map is precomputed only once before the iterative optimization begins, in contrast to classical **ICP** methods [61, 171]. This precomputation encodes data association implicitly, removing the explicit **Nearest Neighbor (NN)** search that traditional **ICP** methods [61, 171] perform at every iteration and replacing it with a constant-time grid lookup.

Therefore, Section 5.2.1.1 defines the distance map over the reference point cloud and discusses how it removes the explicit data association step required by classical **ICP** methods [61, 171]. Section 5.2.1.2 derives the first- and second-order spatial derivatives of the distance map, providing both analytical expressions and finite-difference approximations used in the implementation of the proposed approach. Section 5.2.1.3 describes the three representation levels of increasing information content considered in this work. Lastly, Section 5.2.1.4 details the computation pipeline and the data structures underlying each variant of the distance map representations considered in the scope of this PhD Thesis.

5.2.1.1 Overview

Let $\mathcal{P} = \{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_n\} \subset \mathbb{R}^2$ represent a reference 2D point cloud, comprising n 2D points acquired by a planar laser scanner. These points are expressed in the sensor coordinate frame S , where each point is $\mathbf{p}_i = (x_i, y_i)^T$. The distance map $\mathcal{D} : \mathbb{R}^2 \rightarrow \mathbb{R}_0^+$ is defined as the function that maps any query point $\mathbf{y} \in \mathbb{R}^2$ to its Euclidean distance to the nearest point in $\mathcal{P}$, as formulated in Equation (5.40), where $\mathbf{p}^*(\mathbf{y}) = \operatorname{argmin}_{\mathbf{p}_i \in \mathcal{P}} \|\mathbf{y} - \mathbf{p}_i\|_2$ is the **NN** in the reference cloud $\mathcal{P}$. Note that $\mathcal{D}(\mathbf{y}) \geq 0$ for all $\mathbf{y}$, making this an *unsigned distance function*, as opposed to **SDFs** used in volumetric reconstruction. An important property of *unsigned distance fields* is that they are continuous everywhere, but not everywhere differentiable [185].

$$\mathcal{D}(\mathbf{y}) = \{\mathbb{R}^2 \rightarrow \mathbb{R}_0^+ : \|\mathbf{y} - \mathbf{p}^*(\mathbf{y})\|_2\} \quad (5.40)$$

As discussed in Section 5.1.1, classical **ICP**-based registration approaches [61, 171] implement a two-step iterative process: data association and optimization steps. These methods require reevaluating the correspondences at every iteration, since the pose estimate evolves and changes the transformed positions of the moving points, potentially altering the data association results. In contrast, the distance map formulation removes this coupling entirely. The distance map $\mathcal{D}$ is precomputed over a discrete grid with a given δ (m/px) resolution

covering the environment from the reference cloud $\mathcal{P}$ prior to the registration process. Subsequently, this structure is queried during optimization without requiring any further reference to $\mathcal{P}$. The structure $\mathcal{D}$ encodes information on the **NN** point from $\mathcal{P}$ in each cell of the distance map. Data association is thus resolved implicitly within the error function evaluation itself, transforming a moving point by the current pose estimate. At the resulting point location, looking up into $\mathcal{D}$ yields information (e.g., distance, its derivatives) on the nearest reference point in $O(1)$, with no explicit correspondence search.

5.2.1.2 Spatial derivatives

While the distance map $\mathcal{D}$ is continuous everywhere but not everywhere differentiable [185], its spatial derivatives exist almost everywhere and are central to the registration error formulations derived in Section 5.2.3 and Section 5.2.4. This subsection derives the first- and second-order spatial derivatives of $\mathcal{D}$, providing both analytical expressions and their finite-difference approximations used in our implementation.

First-order derivative (gradient vector) The gradient $\nabla \mathcal{D}(\mathbf{y}) \in \mathbb{R}^2$ of the distance map at a query point $\mathbf{y}$ can be obtained analytically wherever $\mathcal{D}$ is smooth, i.e., at points admitting a unique **NN** in $\mathcal{P}$. At such a point, $\mathbf{p}^*(\mathbf{y})$ is unique and the gradient is a unit vector pointing from the nearest reference point toward $\mathbf{y}$, as given in Equation (5.41), which results from differentiating Equation (5.40) w.r.t. $\mathbf{y}$. Indeed, the unit-norm property $\|\nabla \mathcal{D}\|_2 = 1$ holds almost everywhere for *unsigned distance fields*, with the exception of points equidistant from two or more reference points in $\mathcal{P}$, where the gradient is not defined (e.g., at the center of a circumference). This property represents the Eikonal equation, a fundamental property of Euclidean distance fields [185].

$$\nabla \mathcal{D}(\mathbf{y}) = \frac{\mathbf{y} - \mathbf{p}^*(\mathbf{y})}{\|\mathbf{y} - \mathbf{p}^*(\mathbf{y})\|_2}, \quad \|\nabla \mathcal{D}(\mathbf{y})\|_2 = 1 \quad (5.41)$$

In the discrete grid setting, the gradient is estimated numerically using central finite differences along the y_1 and y_2 directions of $\mathbf{y} = (y_1, y_2)^T$ on the distance map $\mathcal{D}$, as formulated in Equation (5.42), with cell resolution $\delta > 0$ and one-sided (forward or backward) differences used at grid borders.

$$\left. \frac{\partial \mathcal{D}}{\partial y_1} \right|_{\mathbf{y}} \approx \frac{\mathcal{D}(y_1 + \delta, y_2) - \mathcal{D}(y_1 - \delta, y_2)}{2\delta}, \quad \left. \frac{\partial \mathcal{D}}{\partial y_2} \right|_{\mathbf{y}} \approx \frac{\mathcal{D}(y_1, y_2 + \delta) - \mathcal{D}(y_1, y_2 - \delta)}{2\delta} \quad (5.42)$$

Second-order derivative (Hessian matrix) The Hessian matrix $\mathbf{H}_{\mathcal{D}}(\mathbf{y}) \in \mathbb{R}^{2 \times 2}$ is the symmetric matrix of second-order partial derivatives of $\mathcal{D}$, as formulated in Equation (5.43), given that $\frac{\partial^2 \mathcal{D}}{\partial y_1 \partial y_2} = \frac{\partial^2 \mathcal{D}}{\partial y_2 \partial y_1}$. This Hessian matrix represents also the Jacobian matrix's transpose of the analytical gradient $\nabla \mathcal{D}(\mathbf{y})$, as formulated in Equation (5.44).

$$\mathbf{H}_{\mathcal{D}}(\mathbf{y}) = \begin{bmatrix} \frac{\partial^2 \mathcal{D}}{\partial y_1^2} & \frac{\partial^2 \mathcal{D}}{\partial y_1 \partial y_2} \\ \frac{\partial^2 \mathcal{D}}{\partial y_2 \partial y_1} & \frac{\partial^2 \mathcal{D}}{\partial y_2^2} \end{bmatrix}_{\mathbf{y}} = \begin{bmatrix} \frac{(y_2 - p_y^*(\mathbf{y}))^2}{\|\mathbf{y} - \mathbf{p}^*(\mathbf{y})\|_2^3} & -\frac{(y_1 - p_x^*(\mathbf{y}))(y_2 - p_y^*(\mathbf{y}))}{\|\mathbf{y} - \mathbf{p}^*(\mathbf{y})\|_2^3} \\ -\frac{(y_1 - p_x^*(\mathbf{y}))(y_2 - p_y^*(\mathbf{y}))}{\|\mathbf{y} - \mathbf{p}^*(\mathbf{y})\|_2^3} & \frac{(y_1 - p_x^*(\mathbf{y}))^2}{\|\mathbf{y} - \mathbf{p}^*(\mathbf{y})\|_2^3} \end{bmatrix} \quad (5.43)$$

$$\mathbf{H}_{\mathcal{D}}(\mathbf{y}) = \mathbf{J}_{\nabla \mathcal{D}(\mathbf{y})}^T \quad (5.44)$$

In practice, $\mathbf{H}_{\mathcal{D}}$ is computed numerically as the centered finite difference of the normalized neighbor gradients ($\hat{\nabla} \mathcal{D}$), as shown in Equation (5.45), where $(\cdot)_{y_j}$ denotes the j -th component of the enclosed vector.

This numerical approximation is the discretization implemented in the proposed approach for distance map-based 2D point cloud registration when requiring 2nd-order derivatives.

$$\mathbf{H}_{\mathcal{D}}(\mathbf{y}) \approx \frac{1}{2\delta} \begin{bmatrix} \left(\hat{\nabla} \mathcal{D}(y_1+\delta, y_2) - \hat{\nabla} \mathcal{D}(y_1-\delta, y_2) \right)_{y_1} & \left(\hat{\nabla} \mathcal{D}(y_1, y_2+\delta) - \hat{\nabla} \mathcal{D}(y_1, y_2-\delta) \right)_{y_1} \\ \left(\hat{\nabla} \mathcal{D}(y_1+\delta, y_2) - \hat{\nabla} \mathcal{D}(y_1-\delta, y_2) \right)_{y_2} & \left(\hat{\nabla} \mathcal{D}(y_1, y_2+\delta) - \hat{\nabla} \mathcal{D}(y_1, y_2-\delta) \right)_{y_2} \end{bmatrix} \quad (5.45)$$

5.2.1.3 Representation levels

As for the representation of the distance map required for the proposed registration process, each cell of the discrete grid may store different amounts of precomputed information about the distance field. Therefore, we consider three representation levels of increasing information and memory cost. The choice of the representation determines which quantities are available at $O(1)$ query cost during registration, and which must be recomputed on the fly from neighboring cells or analytical expressions.

Level 0 – $\mathcal{M}_0$ (NN index only) Each cell stores only the integer index $k^*(\mathbf{y})$ of the nearest reference point in $\mathcal{P}$ (see Equation (5.46)). The scalar distance $\mathcal{D}(\mathbf{y})$, the gradient vector $\nabla \mathcal{D}(\mathbf{y})$, and the Hessian matrix $\mathbf{H}_{\mathcal{D}}(\mathbf{y})$ are not precomputed nor stored explicitly. These values may be recovered at query time through their analytical expressions discussed in Section 5.2.1.2.

$$\mathcal{M}_0(\mathbf{y}) = k^*(\mathbf{y}) = \underset{i \in \{1, \dots, n\}}{\operatorname{argmin}} \|\mathbf{y} - \mathbf{p}_i\|_2 \in \mathbb{Z} \quad (5.46)$$

Furthermore, $\mathcal{M}_0$ provides the most memory-efficient data structure for data association, and may serve as a drop-in replacement for the kd-tree NN search [173] used in classical ICP methods [61, 171]. This modification reduces the per-query cost from $O(\log n)$ to $O(1)$ at the expense of the precomputation step of $\mathcal{M}_0$.

Level 1 – $\mathcal{M}_1$ (distance and first-order derivatives) Each cell stores the scalar distance $\mathcal{D}(\mathbf{y})$ together with the two components of its spatial gradient $\nabla \mathcal{D}(\mathbf{y}) \in \mathbb{R}^2$, as shown in Equation (5.47). The integer index $k^*(\mathbf{y})$ of the nearest reference point in $\mathcal{P}$ (see Equation (5.46)) is optional for $\mathcal{M}_1$. Hence, $\mathcal{M}_1$ requires three or four scalar values per cell, depending on whether the $k^*(\mathbf{y})$ index is included into $\mathcal{M}_1(\mathbf{y})$.

$$\mathcal{M}_1(\mathbf{y}) = \left(\mathcal{D}(\mathbf{y}), \frac{\partial \mathcal{D}}{\partial y_1}(\mathbf{y}), \frac{\partial \mathcal{D}}{\partial y_2}(\mathbf{y}) \right) \quad (5.47)$$

Subsequently, the scalar distance $\mathcal{D}(\mathbf{y})$ and the gradient vector $\nabla \mathcal{D}(\mathbf{y})$ are precomputed based on Equation (5.40) and using the central finite difference scheme of Equation (5.42), respectively. The Hessian matrix $\mathbf{H}_{\mathcal{D}}(\mathbf{y})$ is not precomputed nor stored explicitly. This value may be recovered at query time through its analytical expression or through the central finite difference approximation from neighbor cells of the normalized gradient discussed in Section 5.2.1.2.

Level 2 – $\mathcal{M}_2$ (distance, first-, and second-order derivatives) Although the analysis in Section 5.2.1.2 demonstrates that the analytical expression for $\mathbf{H}_{\mathcal{D}}$ leads to a symmetric matrix, its centered finite difference approximation presented in Equation (5.45) does not ensure that property. Therefore, each cell of $\mathcal{M}_2$ additionally stores all entries of the approximated Hessian matrix $\mathbf{H}_{\mathcal{D}}(\mathbf{y}) \in \mathbb{R}^{2 \times 2}$, as shown in Equation (5.48). Similarly to

$\mathcal{M}_1$, the integer index $k^*(\mathbf{y})$ (see Equation (5.46)) is optional for the $\mathcal{M}_2$ representation level.

$$\mathcal{M}_2(\mathbf{y}) = \left(\mathcal{D}(\mathbf{y}), \frac{\partial \mathcal{D}}{\partial y_1}(\mathbf{y}), \frac{\partial \mathcal{D}}{\partial y_2}(\mathbf{y}), \frac{\partial^2 \mathcal{D}}{\partial y_1^2}(\mathbf{y}), \frac{\partial^2 \mathcal{D}}{\partial y_1 \partial y_2}(\mathbf{y}), \frac{\partial^2 \mathcal{D}}{\partial y_2 \partial y_1}(\mathbf{y}), \frac{\partial^2 \mathcal{D}}{\partial y_2^2}(\mathbf{y}) \right) \quad (5.48)$$

The representation level $\mathcal{M}_2$ makes available all values throughout the registration process, thanks to the precomputation of the scalar distance $\mathcal{D}(\mathbf{y})$ stored in the discrete grid of the distance and based on Equation (5.40), while both the gradient vector $\nabla \mathcal{D}(\mathbf{y})$ and the Hessian matrix $\mathbf{H}_{\mathcal{D}}(\mathbf{y})$ are computed by centered finite differences from neighbor cells based on Equation (5.42) and Equation (5.45), respectively.

5.2.1.4 Computation and data structures

Independently of the representation level, the computation of $\mathcal{M}_l$ is based on a wavefront propagation algorithm using a best-first priority queue, analogous to Dijkstra's shortest-path algorithm [186]. All variants considered in our approach share the same high-level pipeline:

1. initialize the grid based on its resolution δ (m/px), the bounding box of $\mathcal{P}$, optional fixed map size L (m), and optional truncation distance $d_{\max}$ (m);
2. seed the priority queue with the grid cells occupied by reference points, initializing their distance to zero;
3. propagate the integer index $k^*(\mathbf{y})$ of the nearest reference point and its corresponding distance value outward via best-first wavefront expansion through the $\mathcal{N}_8$ neighborhood (8-connectivity) of each cell in the priority queue for the whole map or until $d_{\max}$ is achieved from the nearest reference point;
4. depending on the required representation level $\mathcal{M}_l$, execute post-processing passes to compute the gradient $\nabla \mathcal{D}$ and the Hessian $\mathbf{H}_{\mathcal{D}}$ through centered finite differences (see Section 5.2.1.2).

Coordinates conversion All implemented variants in our approach for distance map-based representations use the same affine mapping between metric coordinates $\mathbf{y} \in \mathbb{R}^2$ w.r.t. the sensor coordinate frame and integer grid indices $(r, c) \in \mathbb{Z}^2$, and vice-versa. Given $\mathbf{y}_{\mathcal{M}, \min} \in \mathbb{R}^2$, the metric coordinates corresponding to the cell $(r, c)_{\mathcal{M}, \min} = (0, 0)$ of the grid, determined during the initialization procedure of the representation level $\mathcal{M}_l$, Equation (5.49) defines the mapping from the sensor coordinates to the grid indices. Note that negative indices or values that exceed the grid size represent invalid cells outside the data structure.

$$(r, c) = \left\lfloor \frac{\mathbf{y} - \mathbf{y}_{\mathcal{M}, \min}}{\delta} \right\rfloor \in \mathbb{Z}^2 \quad (5.49)$$

Furthermore, Equation (5.50) formulates a sub-pixel (non-integer) version of Equation (5.49), omitting the floor operator. We use this version for bilinear interpolation at query time, enabling sub-pixel accuracy for the distance and derivatives values returned to the registration optimizer.

$$(r, c)_{\mathbb{R}} = \frac{\mathbf{y} - \mathbf{y}_{\mathcal{M}, \min}}{\delta} \in \mathbb{R} \quad (5.50)$$

Lastly, the inverse mapping from grid indices to metric coordinates is obtained by multiplying the indices by the resolution δ and offsetting by $\mathbf{y}_{\mathcal{M}, \min}$, as formulated in Equation (5.51).

$$\mathbf{y} = \delta \cdot (r, c)^T + \mathbf{y}_{\mathcal{M}, \min} \in \mathbb{R}^2 \quad (5.51)$$

Data structure storage Independently of the required $\mathcal{M}_l$, we store the discrete grid in one of two memory layouts. Under *dense storage* ($\mathcal{M}_{l,\text{dense}}$), the grid is allocated as a contiguous array of cells (generic Eigen¹ matrix in C++), where every cell in the grid is populated during precomputation. Under *sparse storage*, only cells satisfying $\mathcal{D}(\mathbf{y}) \leq d_{\max}$ are stored, since moving points projecting further than $d_{\max}$ distance from any reference point are excluded from the registration cost. This truncated representation is analogous to the **Truncated Signed Distance Field (TSDF)** used in volumetric reconstruction [187], where only a narrow band of the distance values is retained around a surface.

Moreover, we implement the sparse grid in C++ as a hash table using the container `std::unordered_map`². The hash key for a cell (r, c) is computed by encoding both integer indices (`int32`) into a single 64-bit unsigned integer, as formulated in Equation (5.52), where $\ll$ denotes a bit-wise left shift and $|$ denotes a bit-wise OR. An important property of `std::unordered_map` is that the container does not invalidate existing element pointers upon rehashing, allowing the priority queue to store raw cell pointers safely throughout the entire wavefront propagation. After propagation, a rehash call consolidates the bucket layout before the gradient and Hessian post-processing passes iterate over all stored entries.

$$\text{key}(r, c) = (\text{uint64}(r) \ll 32) | (\text{uint64}(c) \& 0xFFFFFFFF) \quad (5.52)$$

Computation of the data structures First, Algorithm 5.3 details the initialization procedure for the grid data structure. This structure is parametrized by a map size (L) and a truncation distance ($d_{\max}$). If $L > 0$, the metric bounding box **w.r.t.** the local coordinate frame of the sensor for the distance map is symmetrically defined as $[-L, L]$. Otherwise, the bounding box is computed directly from the reference point cloud $\mathcal{P}$ by iterating over all valid points and accumulating the minimum and maximum coordinates. Then, the grid dimensions in pixels are derived from the bounding box. If $d_{\max} > 0$ is set and no explicit map size is provided, the bounding box is expanded by $d_{\max}$ on both sides before converting to pixel cells, ensuring that the truncation band is fully covered around the point cloud boundary. Otherwise, the bounding box is converted to pixel cell counts using the resolution δ based on the original metric bounding box limits. Finally, the grid is resized to the computed dimensions. The subsequent initialization of cell values depends on the underlying data structure, as detailed in Algorithm 5.4, Algorithm 5.5, and Algorithm 5.6.

Next, Algorithm 5.4, Algorithm 5.5, and Algorithm 5.6 present the computation procedure for each type of distance map structures considered in the scope of this PhD Thesis: *Nano* ($\mathcal{M}_{0,\text{dense}}$), *Standard* ($\mathcal{M}_{1,\text{dense}}$ or $\mathcal{M}_{2,\text{dense}}$), and *Sparse* ($\mathcal{M}_{1,\text{sparse}}$ or $\mathcal{M}_{2,\text{sparse}}$), respectively. Although the $\mathcal{M}_0$ representation could be implemented using a sparse storage layout, the overhead of `std::unordered_map`, which may reserve more hash table buckets than the number of elements actually stored, does not justify a sparse structure for a single integer index per cell.

Furthermore, Algorithm 5.4 details the full construction of the *Nano* variant. First, the entire grid $\mathcal{M}_{0,\text{dense}}$ is filled with -1 , representing unvisited cells, and the priority queue Q is cleared. Next, the reference point cloud $\mathcal{P}$ is projected onto the grid using the coordinate conversion `global2local(.)` of Equation (5.49). Invalid points and points whose projected cell (r, c) lies outside the grid are skipped. For each remaining point $\mathbf{p}_i$, the corresponding cell is assigned the index i . A pointer to the cell is then pushed into Q with distance $d = 0$, which

¹<https://libeigen.gitlab.io/>

²https://en.cppreference.com/w/cpp/container/unordered_map.html

Algorithm 5.3: Initialization procedure to resize the distance map data structures.

Input: Reference point cloud $\mathcal{P} = \{\mathbf{p}_i\}_{i=1}^n$, grid resolution $\delta > 0$ (m/px), map size L (m), truncation distance $d_{\max}$ (m)

Output: Resized grid $\mathcal{M}_l$

```

1 if  $L > 0$  then                                     // Compute metric bounding box of the occupied area
2    $\mathbf{y}_{\mathcal{P},\min} = (-L, -L)$ 
3    $\mathbf{y}_{\mathcal{P},\max} = (L, L)$ 
4 else
5    $\mathbf{y}_{\mathcal{P},\min} = (+\infty, +\infty)$ 
6    $\mathbf{y}_{\mathcal{P},\max} = (-\infty, -\infty)$ 
7   forall  $i = 1, \dots, n$  do
8     if  $\mathbf{p}_i$  is not a valid point then
9       continue
10     $\mathbf{y}_{\mathcal{P},\min} = \min\{\mathbf{y}_{\mathcal{P},\min}, \mathbf{p}_i\}$ 
11     $\mathbf{y}_{\mathcal{P},\max} = \max\{\mathbf{y}_{\mathcal{P},\max}, \mathbf{p}_i\}$ 
12  $\mathbf{y}_{\mathcal{P},\Delta} = \mathbf{y}_{\mathcal{P},\max} - \mathbf{y}_{\mathcal{P},\min}$ 
13 if  $d_{\max} > 0$  and  $L \leq 0$  then                     // Compute size of the grid data structure
14    $\text{rows} = \lfloor \frac{\mathbf{y}_{\mathcal{P},\Delta,x} + 2 \cdot d_{\max}}{\delta} \rfloor + 1$ 
15    $\text{cols} = \lfloor \frac{\mathbf{y}_{\mathcal{P},\Delta,y} + 2 \cdot d_{\max}}{\delta} \rfloor + 1$ 
16    $\mathbf{y}_{\mathcal{M},\min} = \mathbf{y}_{\mathcal{P},\min} - d_{\max}$ 
17 else
18    $\text{rows} = \lfloor \frac{\mathbf{y}_{\mathcal{P},\Delta,x}}{\delta} \rfloor + 1$ 
19    $\text{cols} = \lfloor \frac{\mathbf{y}_{\mathcal{P},\Delta,y}}{\delta} \rfloor + 1$ 
20    $\mathbf{y}_{\mathcal{M},\min} = \mathbf{y}_{\mathcal{P},\min}$ 
21  $\mathcal{M}.\text{resize}(\text{rows}, \text{cols})$ 

```

sets the priority of the queue entry. Since $\mathcal{M}_{0,\text{dense}}$ does not store the distance value, the queue entry must carry it explicitly alongside the cell pointer, so that $\mathcal{Q}$ correctly orders cells during wavefront propagation.

At each iteration of the wavefront propagation, the element q_k with the smallest distance is popped from $\mathcal{Q}$. A negative stored NN index ($k < 0$) indicates an entry overwritten by a closer parent that surpasses the truncation distance $d_{\max}$ when set ($d_{\max} > 0$), and is therefore discarded. For each cell (r', c') in the $\mathcal{N}_8(r_k, c_k)$ neighborhood, the neighbor cell is updated to the closer parent provided the truncation condition is satisfied, i.e., $d_k < d_{\max}$ when truncation is enabled, or unconditionally true when the full map is computed. The first case is when the cell is unvisited ($o < 0$). The other case is if already visited ($o \geq 0$) but with a stored distance greater than the candidate ($d_o > d_k$). In both cases, the updated cell is pushed into $\mathcal{Q}$, allowing the wavefront propagation to correct sub-optimal assignments from earlier iterations.

The *Standard* variant detailed in Algorithm 5.5 follows a similar construction pipeline as the *Nano* variant (see Algorithm 5.4), with two key differences arising from the more complete cell data structure. Since the distance is now part of the cell itself in $\mathcal{M}_{1,\text{dense}}$ and $\mathcal{M}_{2,\text{dense}}$ representation levels, the priority queue $\mathcal{Q}$ does not need to carry the distance separately, accessing directly from the representation of the distance map. The other difference is the distinction between unvisited and already-visited neighbor cells, which required an explicit branch in Algorithm 5.4, is no longer necessary. Since all cells are initialized with $+\infty$, the single condition ($d_k < d_{\max}$ **or** $d_{\max} \leq 0$) **and** $d_o > d_k$ handles both cases uniformly, where an unvisited cell holds a $+\infty$ distance, while a previously visited cell is updated only if a closer parent is found. After the wavefront

Algorithm 5.4: Computation of the *Nano* distance map ($\mathcal{M}_{0,\text{dense}}$).

Input: Reference point cloud $\mathcal{P} = \{\mathbf{p}_i\}_{i=1}^n$, Resized grid $\mathcal{M}_{0,\text{dense}}$
Output: Dense data structure $\mathcal{M}_{0,\text{dense}}$ (see Equation (5.46))

```

1  $\mathcal{M}_{0,\text{dense}}.\text{fill}(-1)$  // Initialization
2  $Q.\text{clear}()$ 
3 for  $i = 1$  to  $n$  do // Seeding
4   if  $\mathbf{p}_i$  is not a valid point then
5     continue
6    $(r, c) = \text{global2local}(\mathbf{p}_i)$ 
7   if  $(r, c)$  is not inside the grid then
8     continue
9    $\mathcal{M}_{0,\text{dense}}(r, c) = i$ 
10   $Q.\text{push}(\&\mathcal{M}_{0,\text{dense}}(r, c), 0)$ 
11 while  $Q$  is not empty do // Wavefront propagation
12    $q_k = Q.\text{pop}()$ 
13    $k = q_k.\text{cell}()$ 
14   if  $k < 0$  then
15     continue
16    $(r_k, c_k) = \mathcal{M}_{0,\text{dense}}.\text{getCellPos}(\&q_k.\text{cell}())$ 
17    $(r_p, c_p) = \text{global2local}(\mathbf{p}_k)$ 
18   foreach  $(r', c') \in \mathcal{N}_g(r_k, c_k)$  do
19      $o = \mathcal{M}_{0,\text{dense}}(r', c')$ 
20      $d_k = \|(r', c')^T - (r_p, c_p)^T\|_2 \cdot \delta$ 
21     if  $o < 0$  then // Unvisited cell
22       if  $d_k < d_{\max}$  or  $d_{\max} \leq 0$  then
23          $\mathcal{M}_{0,\text{dense}}(r', c') = k$ 
24          $Q.\text{push}(\&\mathcal{M}_{0,\text{dense}}(r', c'), d_k)$ 
25       continue
26      $(r_o, c_o) = \text{global2local}(\mathbf{p}_o)$ 
27      $d_o = \|(r_o, c_o)^T - (r_p, c_p)^T\|_2 \cdot \delta$ 
28     if  $(d_k < d_{\max} \text{ or } d_{\max} \leq 0)$  and  $d_o > d_k$  then // Update if closer parent exists
29        $\mathcal{M}_{0,\text{dense}}(r', c') = k$ 
30        $Q.\text{push}(\&\mathcal{M}_{0,\text{dense}}(r', c'), d_k)$ 

```

terminates, the gradient $\nabla \mathcal{D}$ is computed over the entire dense grid via the central finite difference scheme of Equation (5.42). If the $\mathcal{M}_{2,\text{dense}}$ representation is requested, a second post-processing pass computes the Hessian $\mathbf{H}_{\mathcal{D}}$ via Equation (5.45). Both passes use one-sided (forward or backward) differences at grid borders.

As for the distance computation on the *Sparse* variant described in Algorithm 5.6, its construction differs from the *Standard* variant (see Algorithm 5.5) in three steps arising from the `std::unordered_map` storage layout. The `std::unordered_map` is not pre-allocated for the entire map, in order to maintain a sparse storage layout when the truncation distance is set ($d_{\max} > 0$). Instead, the implemented wrapper around `std::unordered_map` defines a default cell whenever a queried index (r, c) is not yet present in the hash table. Next, each entry pushed into the priority queue Q carries more information than in the *Standard* variant: a pointer to the cell $\mathcal{M}_{1|2,\text{sparse}}(r, c)$, the current cell index $\{(r, c)\}$, and the parent **NN** index $\{(r_p, c_p)\}$. This

Algorithm 5.5: Computation of the *Standard* distance map ($\mathcal{M}_{1|2,\text{dense}}$).

Input: Reference point cloud $\mathcal{P} = \{\mathbf{p}_i\}_{i=1}^n$, Resized grid $\mathcal{M}_{1|2,\text{dense}}$
Output: Dense data structure $\mathcal{M}_{1|2,\text{dense}}$ (see Equation (5.47) and Equation (5.48))

```

1  $\mathcal{M}_{1|2,\text{dense}}.\text{fill}(\{-1, +\infty\})$  // Initialization
2  $Q.\text{clear}()$ 
3 for  $i = 1$  to  $n$  do // Seeding
4   if  $\mathbf{p}_i$  is not a valid point then
5     continue
6    $(r, c) = \text{global2local}(\mathbf{p}_i)$ 
7   if  $(r, c)$  is not inside the grid then
8     continue
9    $\mathcal{M}_{1|2,\text{dense}}(r, c) = \{i, 0\}$ 
10   $Q.\text{push}(\&\mathcal{M}_{1|2,\text{dense}}(r, c))$ 
11 while  $Q$  is not empty do // Wavefront propagation
12    $q_k = Q.\text{pop}()$ 
13    $m = q_k.\text{cell}()$ 
14    $k = m.\text{index}()$ 
15   if  $k < 0$  then
16     continue
17    $(r_k, c_k) = \mathcal{M}_{1|2,\text{dense}}.\text{getCellPos}(\&m)$ 
18    $(r_p, c_p) = \text{global2local}(\mathbf{p}_k)$ 
19   foreach  $(r', c') \in \mathcal{N}_8(r_k, c_k)$  do
20      $d_o = m.\text{dist}()$ 
21      $d_k = \|(r', c')^T - (r_p, c_p)^T\|_2 \cdot \delta$ 
22     if  $(d_k < d_{\max} \text{ or } d_{\max} \leq 0)$  and  $d_o > d_k$  then // Update if closer parent exists
23        $\mathcal{M}_{1|2,\text{dense}}(r', c') = \{k, d_k\}$ 
24        $Q.\text{push}(\&\mathcal{M}_{1|2,\text{dense}}(r', c'))$ 

```

implementation avoids additional hash queries at every propagation step. Lastly, once the wavefront terminates and before precomputing the derivatives, a rehash call is issued to consolidate the internal bucket layout of the `std::unordered_map` to match the final number of populated cells, minimizing the load factor and reducing hash collision overhead for the subsequent read-intensive derivative passes.

Table 5.2 summarizes the three distance map variants considered in this PhD Thesis in terms of storage layout, memory occupation, query lookup complexity, and truncation requirements. The *Nano* variant stores a single integer index per cell, resulting in the lowest memory footprint of $O(N_G)$. Both the *Standard* and *Sparse* variants store between 3 and 8 scalar values per cell depending on the representation level $\mathcal{M}_{1|2}$ and whether the optional **NN** index $k^*(\mathbf{y})$ is included. While the *Standard* variant allocates the full grid as a dense array with $O(N_G)$ memory, the *Sparse* variant backed by `std::unordered_map` populates only the $N_d \leq N_G$ cells within the truncation band, making it the most memory-efficient choice when $d_{\max}$ is set, in general. All three variants provide $O(1)$ query lookup, though the *Sparse* variant achieves this only on average, with worst-case $O(N_d)$ under hash collision, mitigated by the post-propagation rehash discussed previously.

Algorithm 5.6: Computation of the *Sparse* distance map ($\mathcal{M}_{1|2,\text{sparse}}$).

Input: Reference point cloud $\mathcal{P} = \{\mathbf{p}_i\}_{i=1}^n$, Resized grid $\mathcal{M}_{1|2,\text{sparse}}$
Output: Sparse data structure $\mathcal{M}_{1|2,\text{sparse}}$ (see Equation (5.47) and Equation (5.48))

```

1  $\mathcal{M}_{1|2,\text{sparse}}.\text{fill}(\{-1, +\infty\})$  // Initialization
2  $Q.\text{clear}()$ 
3 for  $i = 1$  to  $n$  do // Seeding
4   if  $\mathbf{p}_i$  is not a valid point then
5     continue
6    $(r, c) = \text{global2local}(\mathbf{p}_i)$ 
7   if  $(r, c)$  is not inside the grid then
8     continue
9    $\mathcal{M}_{1|2,\text{sparse}}(r, c) = \{i, 0\}$ 
10   $Q.\text{push}(\&\mathcal{M}_{1|2,\text{sparse}}(r, c), \{(r, c)\}, \{(r, c)\})$ 
11 while  $Q$  is not empty do // Wavefront propagation
12    $q_k = Q.\text{pop}()$ 
13    $m = q_k.\text{cell}()$ 
14    $k = m.\text{index}()$ 
15   if  $k < 0$  then
16     continue
17    $(r_k, c_k) = q_k.\text{cell\_idx}()$ 
18    $(r_p, c_p) = q_k.\text{parent\_idx}()$ 
19   foreach  $(r', c') \in \mathcal{N}_8(r_k, c_k)$  do
20      $d_o = m.\text{dist}()$ 
21      $d_k = \|(r', c')^T - (r_p, c_p)^T\|_2 \cdot \delta$ 
22     if  $(d_k < d_{\max} \text{ or } d_{\max} \leq 0)$  and  $d_o > d_k$  then // Update if closer parent exists
23        $\mathcal{M}_{1|2,\text{sparse}}(r', c') = \{k, d_k\}$ 
24        $Q.\text{push}(\&\mathcal{M}_{1|2,\text{sparse}}(r', c'), \{(r', c')\}, \{(r_p, c_p)\})$ 
25  $\mathcal{M}_{1|2,\text{sparse}}.\text{rehash}(\mathcal{M}_{1|2,\text{sparse}}.\text{size}())$ 

```

Table 5.2: Comparison of the three distance map variants (*Nano*, *Standard*, and *Sparse*) considered in the scope this PhD Thesis. N_G denotes the total number of grid cells; $N_d \leq N_G$ represents the number of populated cells within the truncation distance $d_{\max}$.

Variant	Storage	Values/cell	Query lookup	Truncation
<i>Nano</i> ($\mathcal{M}_{0,\text{dense}}$)	Dense (array)	1 (int)	$O(1)$	optional
<i>Standard</i> ($\mathcal{M}_{1 2,\text{dense}}$)	Dense (array)	3–8 [*]	$O(1)$	optional
<i>Sparse</i> ($\mathcal{M}_{1 2,\text{sparse}}$)	Hash (<code>std::unordered_map</code>)	3–8 [*]	$O(1)$ [†]	advisable

^{*} Depends on the representation level: 3 values for $\mathcal{M}_1$, 7 for $\mathcal{M}_2$, +1 if the optional **NN** index $k^*(\mathbf{y})$ is stored.

[†] $O(1)$ average case; worst-case $O(N_d)$ under hash collision, mitigated by the post-propagation rehash.

5.2.2 Problem formulation

Let $\mathbf{z}_i \in \mathbb{R}^2$, $i = 1, \dots, M$, represent the 2D point measurements from the moving (current) point cloud. Similarly to the reference cloud $\mathcal{P}$, the moving cloud is also acquired by a 2D laser scanner and expressed in the sensor coordinate frame $\mathcal{S}$. The goal of 2D point cloud registration is to find the robot pose $\mathbf{x} \in SE(2)$ that best aligns the moving cloud to the reference cloud. The robot pose is represented as a homogeneous transformation on the $SE(2)$ manifold, as formulates in Equation (5.53), where $\mathbf{R}(\theta) \in SO(2)$ is the 2D orthonormal rotation

matrix (see Equation (5.54)) and $\mathbf{t} = (t_x, t_y)^T \in \mathbb{R}^2$ is the translation vector.

$$\mathbf{x} = [\mathbf{R} \mid \mathbf{t}] \in SE(2) \quad (5.53)$$

$$\mathbf{R}(\theta) = \begin{bmatrix} r_{1,1} & r_{1,2} \\ r_{2,1} & r_{2,2} \end{bmatrix} = \begin{bmatrix} \cos(\theta) & -\sin(\theta) \\ \sin(\theta) & \cos(\theta) \end{bmatrix} \in SO(2) \quad (5.54)$$

Furthermore, let $\mathbf{T}_S^R \in SE(2)$ denote the known rigid-body transformation of the sensor *w.r.t.* the robot frame (the extrinsic calibration of the laser scanner), as formulated in Equation (5.55). Then, Equation (5.56) formulates the inverse transformation for the extrinsic parameters of the sensor ($\mathbf{T}_R^S$). As a result, the moving point $\mathbf{z}_i$ expressed in the robot frame is formulated by Equation (5.57), while the same point transformed by the current pose estimate $\mathbf{x}$ is defined by Equation (5.58).

$$\mathbf{T}_S^R = [\mathbf{R}_S^R \mid \mathbf{t}_S^R] \in SE(2) \quad (5.55)$$

$$\mathbf{T}_R^S = [\mathbf{R}_R^S \mid \mathbf{t}_R^S] = [\mathbf{R}_S^{RT} \mid -\mathbf{R}_S^{RT} \mathbf{t}_S^R] \in SE(2) \quad (5.56)$$

$$\mathbf{z}_{R,i} = \mathbf{T}_S^R \mathbf{z}_i \quad (5.57)$$

$$\tilde{\mathbf{z}}_{R,i} = \mathbf{x} \mathbf{T}_S^R \mathbf{z}_i \quad (5.58)$$

Since the optimization variable $\mathbf{x}$ lives on the $SE(2)$ manifold, incremental updates cannot be performed by simple Euclidean addition, as discussed in Section 5.1.2.5. A perturbation $\Delta \mathbf{x} = (\Delta x, \Delta y, \Delta \theta)^T \in \mathbb{R}^3$ is defined in the tangent (Euclidean) space with the exponential and logarithmic mappings defined as Equation (5.59) and Equation (5.60), respectively. Consequently, the manifold update between Gauss-Newton iterations adopt a right-sided perturbation update defined in Equation (5.61), with the corresponding inverse (distance) operation, as formulated in Equation (5.62).

$$\exp(\Delta \mathbf{x}) = [\mathbf{R}(\Delta \theta) \mid (\Delta x, \Delta y)^T] \in SE(2) \quad (5.59)$$

$$\log(\mathbf{x})_{\mathbf{x} \in SE(2)} = (t_x, t_y, \text{atan2}(r_{2,1}, r_{1,1}))^T \in \mathbb{R}^3 \quad (5.60)$$

$$\mathbf{x}_{k+1} = \mathbf{x}_k \oplus \Delta \mathbf{x} = \mathbf{x}_k \cdot \exp(\Delta \mathbf{x}) \quad (5.61)$$

$$\Delta \mathbf{x} = \mathbf{x}_1 \ominus \mathbf{x}_2 = \log(\mathbf{x}_1^{-1} \mathbf{x}_2)^\vee \quad (5.62)$$

Lastly, the cross-product matrix (hat operator) in 3D maps a vector $\mathbf{v} \in \mathbb{R}^3$ to a 3×3 skew-symmetric matrix and gives the rotational sensitivity of a rotated vector. In 2D, an analogous operator is defined for a vector $\mathbf{v} = (v_x, v_y)^T \in \mathbb{R}^2$, as shown in Equation (5.63).

$$[\mathbf{v}] = \left. \frac{\partial \mathbf{R}(\theta) \mathbf{v}}{\partial \theta} \right|_{\theta=0} = \begin{bmatrix} 0 & -1 \\ 1 & 0 \end{bmatrix} \cdot \mathbf{v} = \begin{pmatrix} -v_y \\ v_x \end{pmatrix} \quad (5.63)$$

Intuitively, $[\mathbf{v}]$ is the 90° counter-clockwise rotation of $\mathbf{v}$, *i.e.* the instantaneous velocity of $\mathbf{v}$ under an infinitesimal positive rotation. This skew operator appears in all Jacobian matrices of the 2D point cloud registration problem derived in Section 5.2, corresponding to the column associated with the rotational degree of freedom $\Delta \theta$ of the perturbation $\Delta \mathbf{x}$.

5.2.3 Point-to-point error formulation

In terms of defining the error, first, we propose a point-to-point error defined as the scalar value returned by the distance map when the transformed moving point $\mathbf{z}_i$ is projected onto the reference sensor coordinate frame S . Let $\tilde{\mathbf{z}}_i = \mathbf{T}_R^S \mathbf{x} \mathbf{T}_S^R \mathbf{z}_i \in \mathbb{R}^2$ denote the moving point transformed by the current pose estimate $\mathbf{x}$ and expressed in the reference sensor frame S . The point-to-point scalar residual $e_i^{\text{p2p}}(\mathbf{x}) \in \mathbb{R}$ is then given by Equation (5.64). This scalar residual measures the Euclidean distance from the transformed moving point $\tilde{\mathbf{z}}_i$ to its nearest reference point in $\mathcal{P}$, as encoded by the precomputed distance map $\mathcal{D}$.

$$e_i^{\text{p2p}}(\mathbf{x}) = \mathcal{D}(\tilde{\mathbf{z}}_i) = \mathcal{D}\left(\mathbf{T}_R^S \cdot \mathbf{x} \cdot \mathbf{T}_S^R \mathbf{z}_i\right) \in \mathbb{R} \quad (5.64)$$

5.2.3.1 Jacobian matrix via right-sided perturbation

Applying a right-sided perturbation $\mathbf{x} \oplus \Delta\mathbf{x} = \mathbf{x} \cdot \exp(\Delta\mathbf{x})$ as defined in Equation (5.61), the perturbed error for the proposed point-point error formulation based on distance maps is given by Equation (5.65).

$$e_i^{\text{p2p}}(\mathbf{x} \oplus \Delta\mathbf{x}) = \mathcal{D}\left(\mathbf{T}_R^S \cdot \mathbf{x} \cdot \exp(\Delta\mathbf{x}) \cdot \mathbf{T}_S^R \mathbf{z}_i\right) \quad (5.65)$$

Then, the 1×3 Jacobian matrix $\mathbf{J}_i^{\text{p2p}}$ is obtained by differentiating Equation (5.65) w.r.t. $\Delta\mathbf{x}$ via the chain rule, evaluating at $\Delta\mathbf{x} = \mathbf{0}$, as formulated in Equation (5.66). The chain rule decomposes the derivative into two factors. The first is the spatial gradient $\nabla \mathcal{D}(\tilde{\mathbf{z}}_i)^T \in \mathbb{R}^{1 \times 2}$ over the distance map, evaluated at the projected point $\tilde{\mathbf{z}}_i$ (see Equation (5.64)). The second is the Jacobian of $\tilde{\mathbf{z}}_i$ w.r.t. $\Delta\mathbf{x}$ at $\Delta\mathbf{x} = \mathbf{0}$, denoted $\mathbf{J}_i^{\text{p2p,ICP}} \in \mathbb{R}^{2 \times 3}$ and defined by the underbrace in Equation (5.66). This second factor coincides with the Jacobian of the classical point-to-point **ICP** error, where $\mathbf{z}_{R,i} = \mathbf{T}_S^R \mathbf{z}_i$ is the moving point expressed in the robot frame and $[\cdot]$ is the 2D skew operator defined in Equation (5.63).

$$\mathbf{J}_i^{\text{p2p}} = \left. \frac{\partial e_i^{\text{p2p}}(\mathbf{x} \oplus \Delta\mathbf{x})}{\partial \Delta\mathbf{x}} \right|_{\Delta\mathbf{x}=\mathbf{0}} = \nabla \mathcal{D}(\tilde{\mathbf{z}}_i)^T \cdot \underbrace{\mathbf{R}_R^S \mathbf{R}(\theta) \cdot (\mathbf{I}_{2 \times 2} \mid [\mathbf{z}_{R,i}])}_{\mathbf{J}_i^{\text{p2p,ICP}}} \in \mathbb{R}^{1 \times 3} \quad (5.66)$$

The translational columns of $\mathbf{J}_i^{\text{p2p}}$, associated with Δx and Δy , arise from the identity block $\mathbf{I}_{2 \times 2}$, while the rotational column, associated with $\Delta\theta$, arises from $[\mathbf{z}_{R,i}]$. The spatial gradient $\nabla \mathcal{D}(\tilde{\mathbf{z}}_i)$ is retrieved at $O(1)$ cost from the precomputed distance map at representation level $\mathcal{M}_1$ or $\mathcal{M}_2$, as discussed in Section 5.2.1.3.

5.2.4 Point-to-plane error formulation

Furthermore, we propose a point-to-plane variant for the distance map-based formulation of the 2D point cloud registration problem. This variant incorporates surface normal information to improve convergence in environments dominated by planar or locally smooth structures, analogously to the classical formulation of Chen and Medioni [171]. Here, normals $\mathbf{n}_{z_i} \in \mathbb{R}^2$ are estimated on the moving point cloud in the sensor frame, as detailed in Section 5.2.4.2.

Let $\hat{\nabla} \mathcal{D}(\tilde{\mathbf{z}}_i) = \nabla \mathcal{D}(\tilde{\mathbf{z}}_i) / \|\nabla \mathcal{D}(\tilde{\mathbf{z}}_i)\|_2$ denote the normalized spatial gradient of the distance map at the projected point $\tilde{\mathbf{z}}_i$, obtained by normalizing the finite-difference gradient stored in $\mathcal{M}_1$ or $\mathcal{M}_2$ at query time or from the gradient's analytical expression when using the $\mathcal{M}_0$ representation level (see Section 5.2.1.2). Then,

the product $\hat{\nabla} \mathcal{D}(\tilde{\mathbf{z}}_i) \cdot \mathcal{D}(\tilde{\mathbf{z}}_i)$ recovers the displacement vector from the projected moving point $\tilde{\mathbf{z}}_i$ to its nearest reference point in $\mathcal{P}$. The point-to-plane error $e_i^{\text{p2pl}}(\mathbf{x})$ projects this displacement onto the surface normal rotated to the reference sensor frame, as formulated in Equation (5.67).

$$e_i^{\text{p2pl}}(\mathbf{x}) = (\mathbf{R}(\theta) \mathbf{n}_{z_i})^T \cdot \hat{\nabla} \mathcal{D}(\tilde{\mathbf{z}}_i) \cdot \mathcal{D}(\tilde{\mathbf{z}}_i) \quad (5.67)$$

The normal $\mathbf{n}_{z_i}$ is expressed in the sensor frame of the moving cloud and must be rotated to the reference sensor frame before the projection. Since all rotations in $SO(2)$ are planar and therefore commutative, the extrinsic calibration cancels itself, $\mathbf{R}_R^S \mathbf{R}(\theta) \mathbf{R}_S^R = \mathbf{R}_S^{R^T} \mathbf{R}_S^R \mathbf{R}(\theta) = \mathbf{R}(\theta)$, and Equation (5.67) follows directly.

5.2.4.1 Jacobian matrix via right-sided perturbation

Applying the right-sided perturbation $\mathbf{x} \oplus \Delta \mathbf{x} = \mathbf{x} \cdot \exp(\Delta \mathbf{x})$ as defined in Equation (5.61), the perturbed error on the proposed point-to-plane formulation is given by Equation (5.68), where $\tilde{\mathbf{z}}'_i = \mathbf{T}_R^S \mathbf{x} \exp(\Delta \mathbf{x}) \mathbf{T}_S^R \mathbf{z}_i$ is the perturbed projected point.

$$e_i^{\text{p2pl}}(\mathbf{x} \oplus \Delta \mathbf{x}) = (\mathbf{R}(\theta) \mathbf{R}(\Delta \theta) \mathbf{n}_{z_i})^T \cdot \hat{\nabla} \mathcal{D}(\tilde{\mathbf{z}}'_i) \cdot \mathcal{D}(\tilde{\mathbf{z}}'_i) \quad (5.68)$$

When differentiating *w.r.t.* $\Delta \mathbf{x}$ at $\Delta \mathbf{x} = \mathbf{0}$, the product rule yields three additive contributions to the 1×3 Jacobian $\mathbf{J}_i^{\text{p2pl}}$, as formulated in Equation (5.69). The first term accounts for the rotation of the normal under the perturbation $\Delta \theta$. The second and third terms account for the change in the gradient direction and the change in the distance value, respectively, as the projected point moves under $\Delta \mathbf{x}$. Both terms share the factor $\mathbf{J}_i^{\text{p2p,ICP}} \in \mathbb{R}^{2 \times 3}$ defined in Equation (5.66).

$$\begin{aligned} \mathbf{J}_i^{\text{p2pl}} = \frac{\partial e_i^{\text{p2pl}}(\mathbf{x} \oplus \Delta \mathbf{x})}{\partial \Delta \mathbf{x}} \bigg|_{\Delta \mathbf{x}=\mathbf{0}} &= \mathbf{n}_{z_i}^T \left(\underbrace{\mathbf{0}_{2 \times 2} \mid \left[\mathbf{R}(\theta)^T \cdot \hat{\nabla} \mathcal{D}(\tilde{\mathbf{z}}_i) \cdot \mathcal{D}(\tilde{\mathbf{z}}_i) \right]}_{\text{rotation of the normal}} \right) \\ &\quad + \underbrace{(\mathbf{R}(\theta) \mathbf{n}_{z_i})^T \cdot \mathbf{H}_{\mathcal{D}}(\tilde{\mathbf{z}}_i) \cdot \mathcal{D}(\tilde{\mathbf{z}}_i) \cdot \mathbf{J}_i^{\text{p2p,ICP}}}_{\text{change in gradient direction}} \\ &\quad + \underbrace{(\mathbf{R}(\theta) \mathbf{n}_{z_i})^T \cdot \hat{\nabla} \mathcal{D}(\tilde{\mathbf{z}}_i) \cdot \nabla \mathcal{D}(\tilde{\mathbf{z}}_i)^T \cdot \mathbf{J}_i^{\text{p2p,ICP}}}_{\text{change in distance value}} \end{aligned} \quad (5.69)$$

The first term in Equation (5.69) is non-zero only in the rotational column ($\Delta \theta$), since the translational perturbation does not rotate the normal. The second and third terms contribute to all three columns and require the precomputed Hessian $\mathbf{H}_{\mathcal{D}}$ and gradient $\nabla \mathcal{D}$, both retrieved at $O(1)$ cost from the $\mathcal{M}_2$ representation discussed in Section 5.2.1.3.

5.2.4.2 Normals estimation

Surface normals $\mathbf{n}_{z_i} \in \mathbb{R}^2$ are estimated on the moving point cloud prior to the registration optimization in the scope of this PhD Thesis. This precomputation is performed once per scan, independently of the iterative Gauss-Newton solver. Since a 2D laser scanner produces an angular-ordered sequence of range measurements, the moving cloud $\{\mathbf{z}_i\}_{i=1}^M$ constitutes a 1D-organized array of range observations in the sensor frame, where the

i -th range encodes spatial adjacency along the scan arc. The normal estimation exploits this ordering to collect a spatially coherent local neighborhood for each point, as described below.

Radius neighborhood search For each point $\mathbf{z}_i$, the neighborhood $\mathcal{N}_r(\mathbf{z}_i)$ is formed by collecting all points $\mathbf{z}_j$ whose Euclidean distance to $\mathbf{z}_i$ does not exceed the radius $r > 0$ (m), as formulated in Equation (5.70). Rather than a kd-tree search [173] in the 2D point cloud, the 1D organization of the scan is exploited: indices are traversed in both directions from i along the ordered scan sequence, wrapping circularly at the scan boundaries to handle 360° sensors, and halted when the Euclidean distance to $\mathbf{z}_i$ exceeds r .

$$\mathcal{N}_r(\mathbf{z}_i) = \{\mathbf{z}_j \mid \|\mathbf{z}_j - \mathbf{z}_i\|_2 \leq r\} \quad (5.70)$$

If the cardinality $|\mathcal{N}_r(\mathbf{z}_i)|$ falls below a minimum threshold $N_{\min} \geq 3$, the neighborhood is deemed insufficient for a reliable normal estimate. With fewer than three points, the 2×2 covariance matrix $\mathbf{C}_i$ is at most rank-1, and its eigen-decomposition does not yield a geometrically valid surface normal direction. In such cases, the normal $\mathbf{n}_{z_i}$ is set to zero, marking the point as invalid for the point-to-plane error term.

Covariance-based normal estimation via Principle Component Analysis (PCA) For each valid neighborhood $\mathcal{N}_r(\mathbf{z}_i)$, the 2×2 covariance matrix $\mathbf{C}_i$ is computed as formulated in Equation (5.71), where $N_i = |\mathcal{N}_r(\mathbf{z}_i)|$ and $\bar{\mathbf{z}}_i$ is the sample mean. This formulation admits a single-pass accumulation of $\sum \mathbf{z}_j \mathbf{z}_j^T$ and $\sum \mathbf{z}_j$ without storing the neighborhood explicitly, and is equivalent to the standard population covariance.

$$\bar{\mathbf{z}}_i = \frac{1}{N_i} \sum_{\mathbf{z}_j \in \mathcal{N}_r(\mathbf{z}_i)} \mathbf{z}_j, \quad \mathbf{C}_i = \frac{1}{N_i} \sum_{\mathbf{z}_j \in \mathcal{N}_r(\mathbf{z}_i)} \mathbf{z}_j \mathbf{z}_j^T - \bar{\mathbf{z}}_i \bar{\mathbf{z}}_i^T \in \mathbb{R}^{2 \times 2} \quad (5.71)$$

Since $\mathbf{C}_i$ is real symmetric and positive semi-definite, its eigen-decomposition yields two orthogonal eigenvectors. The eigenvector $\mathbf{v}_{\min}$ associated with the smallest eigenvalue $\lambda_{\min}$ corresponds to the direction of least variance within $\mathcal{N}_r(\mathbf{z}_i)$, which is the local surface normal direction, as given by Equation (5.72).

$$\mathbf{n}_{z_i} = \pm \mathbf{v}_{\min}, \quad \lambda_{\min} = \min(\lambda_1, \lambda_2), \quad \mathbf{C}_i \mathbf{v}_{\min} = \lambda_{\min} \mathbf{v}_{\min} \quad (5.72)$$

Normal orientation The eigenvector $\mathbf{v}_{\min}$ is defined only up to a sign ambiguity. To ensure a consistent orientation convention, the normal is flipped so that it points toward the sensor origin, *i.e.* so that $\mathbf{n}_{z_i}^T \mathbf{z}_i \leq 0$, as formulated in Equation (5.73). This convention reflects the fact that, for a 2D laser scanner, surface normals on obstacle boundaries consistently face the sensor.

$$\mathbf{n}_{z_i} \leftarrow \begin{cases} -\mathbf{v}_{\min} & \text{if } \mathbf{v}_{\min}^T \mathbf{z}_i > 0, \\ \mathbf{v}_{\min} & \text{otherwise.} \end{cases} \quad (5.73)$$

5.3 Experimental Results and Discussion

5.3.1 Synthetic benchmarks

5.3.2 Real-world laser odometry evaluation

5.3.3 Computational performance

5.3.4 Discussion

5.4 Summary and Conclusions

TBC

Chapter 6

Simultaneous Localization and Mapping (SLAM) in Dynamic Environments

TBC

6.1 Background

TBC

6.2 Online SLAM

TBC

6.3 Dynamics Removal

TBC

6.4 Offline Global Bundle Adjustment

TBC

6.5 Experimental Results and Discussion

TBC

6.6 Summary and Conclusions

TBC

Chapter 7

Conclusions and Future Work

TBC

7.1 Conclusions and Contributions

TBC

7.2 Final Remarks and Future Work

7.2.1 Future Work

7.2.2 Involvement in the Scientific Community

TBC

7.2.2.1 Collaborations

Giorgio Sapienza, Rome (PhD Internship, 6 months, February - July 2024).

5dpo Robotics Team:

[188]

R.B. Sousa, P.G. Costa, and A.P. Moreira. “A pose control algorithm for omnidirectional robots”. In: *2021 IEEE International Conference on Autonomous Robot Systems and Competitions (ICARSC)*. Santa Maria da Feira, Portugal, 2021, pp. 91–96. doi: [10.1109/ICARSC52212.2021.9429803](https://doi.org/10.1109/ICARSC52212.2021.9429803)

[189]

R.B. Sousa, C.D. Rocha, J.G. Martins, J.P. Costa, J.T. Padrão, J.M. Sarmiento, J.P. Carvalho, M.S. Lopes, P.G. Costa, and A.P. Moreira. “A robotic framework for the robot@factory 4.0 competition”. In: *2024 IEEE International Conference on Autonomous Robot Systems and Competitions (ICARSC)*. Paredes de Coura, Portugal, 2024, pp. 66–73. doi: [10.1109/ICARSC61747.2024.10535935](https://doi.org/10.1109/ICARSC61747.2024.10535935)

[190]

M.S. Lopes, J.D. Ribeiro, A.P. Moreira, C.D. Rocha, J.G. Martins, J.M. Sarmiento, J.P. Carvalho, P.G. Costa, and R.B. Sousa. “From competition to classroom: a hands-on approach to robotics learning”. In: *2025 IEEE International Conference on Autonomous Robot Systems and Competitions (ICARSC)*. Funchal, Portugal, 2025, pp. 170–176. doi: [10.1109/ICARSC65809.2025.10970153](https://doi.org/10.1109/ICARSC65809.2025.10970153)

7.2.2.2 Scientific projects

APPLICATION: SLAM4all (SLAM em tempo real para multi-domínio com múltiplos robôs sob transmissão de dados limitada).

PROJECTS: VINCI7D (Gerber parser); GreenAuto; PRODUTECH R3.

- PRODUTECH R3:

[191]

T.B. Levin, J.M. Oliveira, R.B. Sousa, M.F. Silva, B.S. Parreira, H.M. Sobreira, and H.S. Mendonça. “Image and command transmission over the 5G network for teleoperation of mobile robots”. In: *2024 7th Iberian Robotics Conference (ROBOT)*. Madrid, Spain, 2024, pp. 1–8. doi: [10.1109/ROBOT61475.2024.10797434](https://doi.org/10.1109/ROBOT61475.2024.10797434)

[192]

D. Caldana, A. Cordeiro, J.P. Souza, R.B. Sousa, P.M. Rebelo, A.J. Silva, and M.F. Silva. “Pallet and pocket detection based on deep learning techniques”. In: *2024 7th Iberian Robotics Conference (ROBOT)*. Madrid, Spain, 2024, pp. 1–8. doi: [10.1109/ROBOT61475.2024.10797410](https://doi.org/10.1109/ROBOT61475.2024.10797410)

[193]

F.D. Pacheco, P.M. Rebelo, R.B. Sousa, M.F. Silva, and H.S. Mendonça. “Integrated RFID system for intralogistics operations with industrial mobile robots”. In: *2025 IEEE International Conference on Autonomous Robot Systems and Competitions (ICARSC)*. Funchal, Portugal, 2025, pp. 124–131. doi: [10.1109/ICARSC65809.2025.10970172](https://doi.org/10.1109/ICARSC65809.2025.10970172)

- VINCI 7D:

[194]

R.B. Sousa, C. Rocha, H.S. Mendonça, A.P. Moreira, and M.F. Silva. “Gerber file parsing for conversion to bitmap image—the VINCI7D case study”. In: *IEEE Access* 10 (2022), pp. 69659–69679. doi: [10.1109/ACCESS.2022.3187042](https://doi.org/10.1109/ACCESS.2022.3187042)

7.2.2.3 Supervision

B.Sc. supervision (not official) at FEUP (February - June 2023; 2nd semester 2022/23): Rodrigo de Vasconcelos e Miguel (Compensação do Movimento do Robô nos Dados de Sensores Exteroceptivos); José Miguel Guedes (Controlo de Trajetórias de Robôs Móveis com o Hangfa Discovery Q2); André Luís Pires (Robótica e Sistemas Autónomos em Localização e Mapeamento de Robôs Móveis com o 3-omni 5DPO MSL); Ana Luísa Mota (Robótica e Sistemas Autónomos em Localização e Mapeamento de Robôs Móveis com o Hangfa Discovery Q2).

B.Sc. supervision (not official) at ISEP (March - July 2023; 2nd semester 2022/23): Paulo Ricardo Barbosa (Teste de algoritmos de mapeamento open source recorrendo ao sensor RS-Helios-5515).

BII at INESC TEC: Daniele da Mota Caldana (Aplicação para deteção de local para manipulação de paletes e paletes danificadas); Duarte Filipe Monteiro (Caixas metálicas GreenAuto); David José Castro (Caixas metálicas GreenAuto); João Miguel Oliveira (tele-controlo AMRs); Thiago Baldassarri Levin (comunicações 5G).

References

- [1] ABI Research. *Commercial and industrial robotics addressable market, segmentation, and technologies market data overview: 3Q 2025*. Tech. rep. ABI Research, July 2025. URL: <https://www.abiresearch.com/market-research/product/7786487-commercial-and-industrial-robotics-address>.
- [2] R. Siegwart, I.R. Nourbakhsh, and D. Scaramuzza. *Introduction to autonomous mobile robots*. Ed. by R.C. Arkin. Second. The MIT Press, 2004. ISBN: 978-0-262-01535-6.
- [3] C. Cadena, L. Carlone, H. Carrillo, Y. Latif, D. Scaramuzza, J. Neira, I. Reid, and J.J. Leonard. “Past, present, and future of simultaneous localization and mapping: toward the robust-perception age”. In: *IEEE Transactions on Robotics* 32.6 (2016), pp. 1309–1332. DOI: [10.1109/TRO.2016.2624754](https://doi.org/10.1109/TRO.2016.2624754).
- [4] P. Biber and T. Duckett. “Experimental analysis of sample-based maps for long-term SLAM”. In: *International Journal of Robotics Research* 28.1 (2009), pp. 20–33. DOI: [10.1177/0278364908096286](https://doi.org/10.1177/0278364908096286).
- [5] H. Kretzschmar and C. Stachniss. “Information-theoretic compression of pose graphs for laser-based SLAM”. In: *International Journal of Robotics Research* 31.11 (2012), pp. 1219–1230. DOI: [10.1177/0278364912455072](https://doi.org/10.1177/0278364912455072).
- [6] F. Dayoub, G. Cielniak, and T. Duckett. “Long-term experiments with an adaptive spherical view representation for navigation in changing environments”. In: *Robotics and Autonomous Systems* 59.5 (2011), pp. 285–295. DOI: [10.1016/j.robot.2011.02.013](https://doi.org/10.1016/j.robot.2011.02.013).
- [7] B. Bacca, J. Salví, and X. Cufí. “Long-term mapping and localization using feature stability histograms”. In: *Robotics and Autonomous Systems* 61.12 (2013), pp. 1539–1558. DOI: [10.1016/j.robot.2013.07.003](https://doi.org/10.1016/j.robot.2013.07.003).
- [8] Z. Xing, X. Zhu, and D. Dong. “DE-SLAM: SLAM for highly dynamic environment”. In: *Journal of Field Robotics* 39.5 (2022), pp. 528–542. DOI: [10.1002/rob.22062](https://doi.org/10.1002/rob.22062).
- [9] J. Biswas and M.M. Veloso. “Episodic non-Markov localization”. In: *Robotics and Autonomous Systems* 87 (2017), pp. 162–176. DOI: [10.1016/j.robot.2016.09.005](https://doi.org/10.1016/j.robot.2016.09.005).
- [10] A. Walcott-Bryant, M. Kaess, H. Johannsson, and J.J. Leonard. “Dynamic pose graph SLAM: long-term mapping in low dynamic environments”. In: *2012 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. Vilamoura-Algarve, Portugal, 2012, pp. 1871–1878. DOI: [10.1109/IROS.2012.6385561](https://doi.org/10.1109/IROS.2012.6385561).
- [11] M. Zhang, Y. Chen, and M. Li. “SDF-Loc: signed distance field based 2D relocalization and map update in dynamic environments”. In: *2019 American Control Conference (ACC)*. Philadelphia, PA, USA, 2019, pp. 1997–2004. DOI: [10.23919/ACC.2019.8814347](https://doi.org/10.23919/ACC.2019.8814347).

- [12] P. Ganti and S.L. Waslander. “Network uncertainty informed semantic feature selection for visual SLAM”. In: *2019 16th Conference on Computer and Robot Vision (CRV)*. Kingston, QC, Canada, 2019, pp. 121–128. DOI: [10.1109/CRV.2019.00024](https://doi.org/10.1109/CRV.2019.00024).
- [13] S. Zhu, X. Zhang, S. Guo, J. Li, and H. Liu. “Lifelong localization in semi-dynamic environment”. In: *2021 IEEE International Conference on Robotics and Automation (ICRA)*. Xi’an, China, 2021, pp. 14389–14395. DOI: [10.1109/ICRA48506.2021.9561584](https://doi.org/10.1109/ICRA48506.2021.9561584).
- [14] R. Guo and X. Liu. “Ground enhanced RGB-D SLAM for dynamic environments”. In: *2021 IEEE International Conference on Robotics and Biomimetics (ROBIO)*. Sanya, China, 2021, pp. 1171–1177. DOI: [10.1109/ROBIO54168.2021.9739362](https://doi.org/10.1109/ROBIO54168.2021.9739362).
- [15] T. Deng, H. Xie, J. Wang, and W. Chen. “Long-term visual simultaneous localization and mapping: using a Bayesian persistence filter-based global map prediction”. In: *IEEE Robotics & Automation Magazine* 30.1 (2023), pp. 36–49. DOI: [10.1109/MRA.2022.3228492](https://doi.org/10.1109/MRA.2022.3228492).
- [16] Y. Wang, Y. Fan, J. Wang, and W. Chen. “Long-term navigation for autonomous robots based on spatio-temporal map prediction”. In: *Robotics and Autonomous Systems* 179 (2024), p. 104724. DOI: [10.1016/j.robot.2024.104724](https://doi.org/10.1016/j.robot.2024.104724).
- [17] X. Li, S. Yuan, H. Cai, S. Lu, W. Wang, and J. Liu. “LL-localizer: a lifelong localization system based on dynamic i-octree”. In: *IEEE Transactions on Instrumentation and Measurement* 74 (2025), pp. 1–11. DOI: [10.1109/TIM.2025.3598401](https://doi.org/10.1109/TIM.2025.3598401).
- [18] B. Bescos, J.M. Fácil, J. Civera, and J. Neira. “DynaSLAM: tracking, mapping, and inpainting in dynamic scenes”. In: *IEEE Robotics and Automation Letters* 3.4 (2018), pp. 4076–4083. DOI: [10.1109/LRA.2018.2860039](https://doi.org/10.1109/LRA.2018.2860039).
- [19] K. Wang, Y. Lin, L. Wang, L. Han, M. Hua, X. Wang, S. Lian, and B. Huang. “A unified framework for mutual improvement of SLAM and semantic segmentation”. In: *2019 International Conference on Robotics and Automation (ICRA)*. Montreal, QC, Canada, 2019, pp. 5224–5230. DOI: [10.1109/ICRA.2019.8793499](https://doi.org/10.1109/ICRA.2019.8793499).
- [20] J.C.V. Soares, V.S. Medeiros, G.F. Abati, M. Becker, G. Caurin, M. Gattass, and M.A. Meggiolaro. “Visual localization and mapping in dynamic and changing environments”. In: *Journal of Intelligent & Robotic Systems* 109.4 (2023), p. 95. DOI: [10.1007/s10846-023-02019-6](https://doi.org/10.1007/s10846-023-02019-6).
- [21] J.P. Saarinen, H. Andreasson, T. Stoyanov, and A.J. Lilienthal. “3D normal distributions transform occupancy maps: an efficient representation for mapping in dynamic environments”. In: *International Journal of Robotics Research* 32.14 (2013), pp. 1627–1644. DOI: [10.1177/0278364913499415](https://doi.org/10.1177/0278364913499415).
- [22] M.T. Lázaro, R. Capobianco, and G. Grisetti. “Efficient long-term mapping in dynamic environments”. In: *2018 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. Madrid, Spain, 2018, pp. 153–160. DOI: [10.1109/IROS.2018.8594310](https://doi.org/10.1109/IROS.2018.8594310).
- [23] F. Boniardi, T. Caselitz, R. Kümmerle, and W. Burgard. “A pose graph-based localization system for long-term navigation in CAD floor plans”. In: *Robotics and Autonomous Systems* 112 (2019), pp. 84–97. DOI: [10.1016/j.robot.2018.11.003](https://doi.org/10.1016/j.robot.2018.11.003).

- [24] G. Kim and A. Kim. “LT-mapper: a modular framework for LiDAR-based lifelong mapping”. In: *2022 International Conference on Robotics and Automation (ICRA)*. Philadelphia, PA, USA, 2022, pp. 7995–8002. doi: [10.1109/ICRA46639.2022.9811916](https://doi.org/10.1109/ICRA46639.2022.9811916).
- [25] E. Stefanini, E. Ciancolini, A. Settimi, and L. Pallottino. “Safe and robust map updating for long-term operations in dynamic environments”. In: *Sensors* 23.13 (2023), p. 6066. doi: [10.3390/s23136066](https://doi.org/10.3390/s23136066).
- [26] I. Hroob, S. Molina, R. Polvara, G. Cielniak, and M. Hanheide. “Adaptive robot localization in dynamic environments through self-learned long-term 3D stable points segmentation”. In: *Robotics and Autonomous Systems* 181 (2024), p. 104786. doi: [10.1016/j.robot.2024.104786](https://doi.org/10.1016/j.robot.2024.104786).
- [27] I. Hroob, B. Mersch, C. Stachniss, and M. Hanheide. “Generalizable stable points segmentation for 3D LiDAR scan-to-map long-term localization”. In: *IEEE Robotics and Automation Letters* 9.4 (2024), pp. 3546–3553. doi: [10.1109/LRA.2024.3368236](https://doi.org/10.1109/LRA.2024.3368236).
- [28] J. Zhu, W. Mi, H. Li, T. Xue, J. Zhao, S. Wang, and T. Zhang. “SPLIN: a structured plane-based LiDAR-inertial SLAM with dynamic object removal”. In: *IEEE Transactions on Instrumentation and Measurement* 74 (2025), pp. 1–15. doi: [10.1109/TIM.2025.3643083](https://doi.org/10.1109/TIM.2025.3643083).
- [29] L. Yang, S.M. Prakhya, S. Zhu, and Z. Liu. “Lifelong 3D mapping framework for hand-held & robot-mounted LiDAR mapping systems”. In: *IEEE Robotics and Automation Letters* 9.11 (2024), pp. 9446–9453. doi: [10.1109/LRA.2024.3417113](https://doi.org/10.1109/LRA.2024.3417113).
- [30] S. Wu, R. Zhang, and Y. Wang. “Dynamic visual SLAM based on efficient two-stage statistical model”. In: *2022 6th CAA International Conference on Vehicular Control and Intelligence (CVCI)*. Nanjing, China, 2022, pp. 1–6. doi: [10.1109/CVCI56766.2022.9964605](https://doi.org/10.1109/CVCI56766.2022.9964605).
- [31] C. Liu, J. Qin, S. Wang, L. Yu, and Y. Wang. “Accurate RGB-D SLAM in dynamic environments based on dynamic visual feature removal”. In: *Science China Information Sciences* 65.10 (2022), p. 202206. doi: [10.1007/s11432-021-3425-8](https://doi.org/10.1007/s11432-021-3425-8).
- [32] Z.-J. Du, S.-S. Huang, T.-J. Mu, Q. Zhao, R.R. Martin, and K. Xu. “Accurate dynamic SLAM using CRF-based long-term consistency”. In: *IEEE Transactions on Visualization and Computer Graphics* 28.4 (2022), pp. 1745–1757. doi: [10.1109/TVCG.2020.3028218](https://doi.org/10.1109/TVCG.2020.3028218).
- [33] Y. Wang, Y. Tian, J. Chen, C. Chen, K. Xu, and X. Ding. “MSSD-SLAM: multifeature semantic RGB-D inertial SLAM with structural regularity for dynamic environments”. In: *IEEE Transactions on Instrumentation and Measurement* 74 (2025), pp. 1–17. doi: [10.1109/TIM.2024.3509541](https://doi.org/10.1109/TIM.2024.3509541).
- [34] S. Cheng, S. He, F. Duan, and N. An. “TLS-SLAM: Gaussian splatting SLAM tailored for large-scale scenes”. In: *IEEE Robotics and Automation Letters* 10.3 (2025), pp. 2814–2821. doi: [10.1109/LRA.2025.3536876](https://doi.org/10.1109/LRA.2025.3536876).
- [35] X. Hu, Y. Wu, M. Zhao, Z. Cao, X. Zhang, and X. Ji. “DYO-SLAM: visual localization and object mapping in dynamic scenes”. In: *IEEE Transactions on Circuits and Systems for Video Technology* 35.11 (2025), pp. 11200–11215. doi: [10.1109/TCSVT.2025.3571049](https://doi.org/10.1109/TCSVT.2025.3571049).
- [36] Z. Jia, Q. Li, L. Zhan, and Z. Wang. “DGS-SLAM: robust visual SLAM with 3D Gaussian splatting in dynamic environments”. In: *IEEE Transactions on Circuits and Systems for Video Technology* (2025), pp. 1–1. doi: [10.1109/TCSVT.2025.3645351](https://doi.org/10.1109/TCSVT.2025.3645351).

- [37] F. Pomerleau, P. Krüsi, F. Colas, P. Furgale, and R. Siegwart. “Long-term 3D map maintenance in dynamic environments”. In: *2014 IEEE International Conference on Robotics and Automation (ICRA)*. Hong Kong, China, 2014, pp. 3712–3719. doi: [10.1109/ICRA.2014.6907397](https://doi.org/10.1109/ICRA.2014.6907397).
- [38] X. Ding, Y. Wang, R. Xiong, D. Li, L. Tang, H. Yin, and L. Zhao. “Persistent stereo visual localization on cross-modal invariant map”. In: *IEEE Transactions on Intelligent Transportation Systems* 21.11 (2020), pp. 4646–4658. doi: [10.1109/TITS.2019.2942760](https://doi.org/10.1109/TITS.2019.2942760).
- [39] H. Thomas, M.G. de Saint Aurin, J. Zhang, and T.D. Barfoot. “Learning spatiotemporal occupancy grid maps for lifelong navigation in dynamic scenes”. In: *2022 International Conference on Robotics and Automation (ICRA)*. Philadelphia, PA, USA, 2022, pp. 484–490. doi: [10.1109/ICRA46639.2022.9812297](https://doi.org/10.1109/ICRA46639.2022.9812297).
- [40] E. Einhorn and H.-M. Gross. “Generic NDT mapping in dynamic environments and its application for lifelong SLAM”. In: *Robotics and Autonomous Systems* 69 (2015), pp. 28–39. doi: [10.1016/j.robot.2014.08.008](https://doi.org/10.1016/j.robot.2014.08.008).
- [41] Z. Pan, H. Chen, S. Li, and Y. Liu. “ClusterMap building and relocalization in urban environments for unmanned vehicles”. In: *Sensors* 19.19 (2019), p. 4252. doi: [10.3390/s19194252](https://doi.org/10.3390/s19194252).
- [42] S. Yang, G. Fan, L. Bai, C. Zhao, and D. Li. “SGC-VSLAM: a semantic and geometric constraints VSLAM for dynamic indoor environments”. In: *Sensors* 20.8 (2020), p. 2432. doi: [10.3390/s20082432](https://doi.org/10.3390/s20082432).
- [43] Y. Yue, C. Yang, J. Zhang, M. Wen, Z. Wu, H. Zhang, and D. Wang. “Day and night collaborative dynamic mapping in unstructured environment based on multimodal sensors”. In: *2020 IEEE International Conference on Robotics and Automation (ICRA)*. Paris, France, 2020, pp. 2981–2987. doi: [10.1109/ICRA40945.2020.9197072](https://doi.org/10.1109/ICRA40945.2020.9197072).
- [44] Y. Chen, S. Sun, H. Yin, and M.H. Ang. “Exploring the effect of 3D object removal using deep learning for LiDAR-based mapping and long-term vehicular localization”. In: *2022 IEEE 25th International Conference on Intelligent Transportation Systems (ITSC)*. Macau, China, 2022, pp. 1730–1735. doi: [10.1109/ITSC55140.2022.9921969](https://doi.org/10.1109/ITSC55140.2022.9921969).
- [45] Y. Song, D. Zhu, J. Li, Y. Tian, and M. Li. “Learning local feature descriptor with motion attribute for vision-based localization”. In: *2019 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. Macau, China, 2019, pp. 3794–3801. doi: [10.1109/IROS40897.2019.8967749](https://doi.org/10.1109/IROS40897.2019.8967749).
- [46] G. Singh, M. Wu, S.-K. Lam, and D.V. Minh. “Hierarchical loop closure detection for long-term visual SLAM with semantic-geometric descriptors”. In: *2021 IEEE International Intelligent Transportation Systems Conference (ITSC)*. Indianapolis, IN, USA, 2021, pp. 2909–2916. doi: [10.1109/ITSC48978.2021.9564866](https://doi.org/10.1109/ITSC48978.2021.9564866).
- [47] N. Yu, Y. Wang, M. Gan, N. Ma, and G. Huang. “Lifelong SLAM based on environmental semantic and depth information in dynamic scenes”. In: *2023 7th Asian Conference on Artificial Intelligence Technology (ACAIT)*. Jiaxing, China, 2023, pp. 764–770. doi: [10.1109/ACAIT60137.2023.10528652](https://doi.org/10.1109/ACAIT60137.2023.10528652).
- [48] G.D. Tipaldi, D. Meyer-Delius, and W. Burgard. “Lifelong localization in changing environments”. In: *International Journal of Robotics Research* 32.14 (2013), pp. 1662–1678. doi: [10.1177/0278364913502830](https://doi.org/10.1177/0278364913502830).

- [49] M. Rapp, M. Hahn, M. Thom, J. Dickmann, and K. Dietmayer. “Semi-Markov process based localization using radar in dynamic environments”. In: *2015 IEEE 18th International Conference on Intelligent Transportation Systems (ITSC)*. Gran Canaria, Spain, 2015, pp. 423–429. doi: [10.1109/ITSC.2015.77](https://doi.org/10.1109/ITSC.2015.77).
- [50] L. Wang, W. Chen, and J. Wang. “Long-term localization with time series map prediction for mobile robots in dynamic environments”. In: *2020 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. Las Vegas, NV, USA, 2020, pp. 1–7. doi: [10.1109/IROS45743.2020.9468884](https://doi.org/10.1109/IROS45743.2020.9468884).
- [51] J.M. Santos, T. Krajník, J.P. Fentanes, and T. Duckett. “Lifelong information-driven exploration to complete and refine 4-D spatio-temporal maps”. In: *IEEE Robotics and Automation Letters* 1.2 (2016), pp. 684–691. doi: [10.1109/LRA.2016.2516594](https://doi.org/10.1109/LRA.2016.2516594).
- [52] T. Krajník, J.P. Fentanes, J.M. Santos, and T. Duckett. “FreMEEn: frequency map enhancement for long-term mobile robot autonomy in changing environments”. In: *IEEE Transactions on Robotics* 33.4 (2017), pp. 964–977. doi: [10.1109/TRO.2017.2665664](https://doi.org/10.1109/TRO.2017.2665664).
- [53] F. Heintz, N. Belbachir, and E. Curry. *Strategic Research, Innovation, and Deployment Agenda 2025–2027 (Version no. 1.3)*. Tech. rep. Brussels, Belgium: AI, Data and Robotics Association (ADRA), Feb. 2024. URL: https://adr-association.eu/sites/default/files/2024-09/Adra%20Strategic%20Research%20Jul24_v2-2_0.pdf.
- [54] A. Pastor. *14 types of Autonomous Mobile Robots (AMR) you need to know for your business*. Last accessed on Dec. 4, 2025. Autonomous Mobile Robot (AGV) Network. URL: <https://www.agvnetwork.com/types-of-autonomous-mobile-robots>.
- [55] B. Cisneros. *The power of dynamic slotting: optimizing warehouse efficiency for modern businesses*. Last accessed on Dec. 4, 2025. Quality Material Handling Inc., Jan. 2025. URL: <https://www.qmhinc.com/dynamic-slotting/>.
- [56] Agenda GreenAuto. *GreenAuto: green innovation for the automotive industry*. Last accessed on Jan. 14, 2025. Peugeot Citroën Automóveis Portugal, S.A. (Stellantis Mangualde). URL: <https://www.agendagreenauto.pt/>.
- [57] R.B. Sousa, M.R. Petry, P.G. Costa, and A.P. Moreira. “OptiOdom: a generic approach for odometry calibration of wheeled mobile robots”. In: *Journal of Intelligent & Robotic Systems* 105.39 (2022). doi: [10.1007/s10846-022-01630-3](https://doi.org/10.1007/s10846-022-01630-3).
- [58] *ISO 10218-1:2025: robotics – safety requirements. Part 1: industrial robots*. Standard. International Organization for Standardization (ISO), 2025.
- [59] *ANSI/A3 R15.06-2025: American national standard for industrial robots and robot systems – safety requirements (parts 1, 2, and 3)*. Standard. American National Standards Institute (ANSI). Association for Advancing Automation (A3), 2025.
- [60] *ISO 3691-4:2023: industrial trucks – safety requirements. Part 1: driverless industrial trucks and their systems*. Standard. International Organization for Standardization (ISO), 2023.
- [61] P.J. Besl and N.D. McKay. “A method for registration of 3-D shapes”. In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 14.2 (1992), pp. 239–256. doi: [10.1109/34.121791](https://doi.org/10.1109/34.121791).
- [62] G. Grisetti, T. Guadagnino, I. Aloise, M. Colosi, B. Della Corte, and D. Schlegel. “Least squares optimization: from theory to practice”. In: *Robotics* 9.3 (2020), p. 51. doi: [10.3390/robotics9030051](https://doi.org/10.3390/robotics9030051).

- [63] Flowbotic Mobile Systems. *Flowbotic*. Last accessed on Feb. 9, 2026. URL: <https://www.flowbotic.eu/>.
- [64] United Nations. *Transforming our world: the 2030 agenda for sustainable development*. Last accessed on Jan. 15, 2025. 2015. URL: <https://sdgs.un.org/publications/transforming-our-world-2030-agenda-sustainable-development-17981>.
- [65] R.B. Sousa, H.M. Sobreira, and A.P. Moreira. “A systematic literature review on long-term localization and mapping for mobile robots”. In: *Journal of Field Robotics* 40.5 (2023), pp. 1245–1322. DOI: [10.1002/rob.22170](https://doi.org/10.1002/rob.22170).
- [66] R.B. Sousa, H.M. Sobreira, M.F. Silva, and A.P. Moreira. “Line fitting-based corner-like detector for 2D laser scanners data”. In: *2024 10th International Conference on Automation, Robotics and Applications (ICARA)*. Athens, Greece, 2024, pp. 210–216. DOI: [10.1109/ICARA60736.2024.10553006](https://doi.org/10.1109/ICARA60736.2024.10553006).
- [67] R.B. Sousa, H.M. Sobreira, J.G. Martins, P.G. Costa, M.F. Silva, and A.P. Moreira. “Integrating multimodal perception into ground mobile robots”. In: *2025 IEEE International Conference on Autonomous Robot Systems and Competitions (ICARSC)*. Funchal, Portugal, 2025, pp. 104–111. DOI: [10.1109/ICARSC65809.2025.10970176](https://doi.org/10.1109/ICARSC65809.2025.10970176).
- [68] J.D. Ribeiro, R.B. Sousa, J.G. Martins, A.S. Aguiar, F.N. dos Santos, and H.M. Sobreira. “Indoor benchmark of 3-D LiDAR SLAM at iilab – industry and innovation laboratory”. In: *IEEE Access* 13 (2025), pp. 212421–212442. DOI: [10.1109/ACCESS.2025.3643753](https://doi.org/10.1109/ACCESS.2025.3643753).
- [69] M.S. Lopes, A. Cordeiro, R.B. Sousa, J.A. Beça, P. Costa, J.P.C. Souza, and M.F. Silva. “CARGO: container-unloading autonomous robot with generalized operation”. In: *2026 12th International Conference on Automation, Robotics and Applications (ICARA)*. Istanbul, Turkey, 2026, pp. 1–6.
- [70] J.D. Ribeiro, R.B. Sousa, J.G. Martins, A.S. Aguiar, F.N. dos Santos, and H.M. Sobreira. *IILABS 3D: iilab indoor LiDAR-based SLAM dataset*. iilab – industry and innovation lab: INESC TEC – Institute for Systems, Computer Engineering, Technology, and Science, 2025. DOI: [10.25747/VHNJ-WM80](https://doi.org/10.25747/VHNJ-WM80).
- [71] Chengdu Hangfa Robotics Co. *Discovery Q2 four-wheeled omni-directional platform*. Last accessed on Feb. 9, 2026. URL: <http://en.hangfa.com/products/research-platform/discoveryQ2-7.html>.
- [72] S. Kohlbrecher, O. von Stryk, J. Meyer, and U. Klingauf. “A flexible and scalable SLAM system with full 3D motion estimation”. In: *2011 IEEE International Symposium on Safety, Security, and Rescue Robotics (SSRR)*. Kyoto, Japan, 2011, pp. 155–160. DOI: [10.1109/SSRR.2011.6106777](https://doi.org/10.1109/SSRR.2011.6106777).
- [73] J. Sturm, N. Engelhard, F. Endres, W. Burgard, and D. Cremers. “A benchmark for the evaluation of RGB-D SLAM systems”. In: *2012 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. Vilamoura-Algarve, Portugal, 2012, pp. 573–580. DOI: [10.1109/IROS.2012.6385773](https://doi.org/10.1109/IROS.2012.6385773).
- [74] G. Grisetti, C. Stachniss, and W. Burgard. “Improved techniques for grid mapping with Rao-Blackwellized particle filters”. In: *IEEE Transactions on Robotics* 23.1 (2007), pp. 34–46. DOI: [10.1109/TRO.2006.889486](https://doi.org/10.1109/TRO.2006.889486).
- [75] B.P. Gerkey, R.T. Vaughan, K. Stoy, A. Howard, G.S. Sukhatme, and M.J. Mataric. “Most valuable player: a robot device server for distributed control”. In: *2001 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. Vol. 3. Maui, HI, USA, 2001, pp. 1226–1231. DOI: [10.1109/IROS.2001.977150](https://doi.org/10.1109/IROS.2001.977150).

- [76] R. Vaughan. “Massively multi-robot simulation in stage”. In: *Swarm Intelligence 2.2* (2008), pp. 189–208. DOI: [10.1007/s11721-008-0014-4](https://doi.org/10.1007/s11721-008-0014-4).
- [77] The Player Project. *Stage*. 2017. URL: <https://playerproject.github.io/stage/>.
- [78] S. Thrun, W. Burgard, and D. Fox. *Probabilistic robotics*. The MIT Press, 2005. ISBN: 978-0-262-20162-9.
- [79] B.D. Lucas and T. Kanade. “An iterative image registration techniques with an application to stereo vision”. In: *Proceedings of the 7th International Joint Conference on Artificial Intelligence (IJCAI)*. San Francisco, CA, USA, 1981, pp. 121–130. URL: https://www.ri.cmu.edu/pub_files/pub3/lucas_bruce_d_1981_2/lucas_bruce_d_1981_2.pdf.
- [80] M. Lauer, S. Lange, and M. Riedmiller. “Calculating the perfect match: an efficient and accurate approach for robot self-localization”. In: *RoboCup 2005: Robot Soccer World Cup IX*. Ed. by A. Bredenfeld, A. Jacoff, I. Noda, and Y. Takahashi. Berlin, Heidelberg: Springer-Verlag, 2006, pp. 142–153. DOI: [10.1007/11780519_13](https://doi.org/10.1007/11780519_13).
- [81] J. Serafin and G. Grisetti. “NICP: dense normal based point cloud registration”. In: *2015 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. Hamburg, Germany, 2015, pp. 742–749. DOI: [10.1109/IROS.2015.7353455](https://doi.org/10.1109/IROS.2015.7353455).
- [82] P. Biber and W. Strasser. “The normal distributions transform: a new approach to laser scan matching”. In: *2003 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. Vol. 3. Las Vegas, NV, USA, 2003, pp. 2743–2748. DOI: [10.1109/IROS.2003.1249285](https://doi.org/10.1109/IROS.2003.1249285).
- [83] A. Doucet, N. Freitas, K. Murphy, and S. Russell. “Rao-Blackwellised filtering for dynamic Bayesian networks”. In: *Proceedings of the 16th Conference on Uncertainty in Artificial Intelligence (UAI)*. San Francisco, CA, USA, 2000, pp. 176–183. URL: <https://dl.acm.org/doi/10.5555/2073946.2073968>.
- [84] G. Grisetti, R. Kümmerle, C. Stachniss, and W. Burgard. “A tutorial on graph-based SLAM”. In: *IEEE Intelligent Transportation Systems Magazine 2.4* (2010), pp. 31–43. DOI: [10.1109/MITS.2010.939925](https://doi.org/10.1109/MITS.2010.939925).
- [85] M. Montemerlo, S. Thrun, D. Koller, and B. Wegbreit. “FastSLAM: a factored solution to the simultaneous localization and mapping problem”. In: *Proceedings of the Eighteenth National Conference on Artificial Intelligence (AAAI)*. Edmonton, Alberta, Canada, 2002, pp. 593–598. URL: <https://aaai.org/Papers/AAAI/2002/AAAI02-089.pdf>.
- [86] M. Montemerlo, S. Thrun, D. Koller, and B. Wegbreit. “FastSLAM 2.0: an improved particle filtering algorithm for simultaneous localization and mapping that provably converges”. In: *Proceedings of the Eighteenth International Joint Conference on Artificial Intelligence (IJCAI)*. San Francisco, CA, USA, 2003, pp. 1151–1156. URL: <https://dl.acm.org/doi/10.5555/1630659.1630824>.
- [87] G. Bresson, Z. Alsayed, L. Yu, and S. Glaser. “Simultaneous localization and mapping: a survey of current trends in autonomous driving”. In: *IEEE Transactions on Intelligent Vehicles 2.3* (2017), pp. 194–220. DOI: [10.1109/TIV.2017.2749181](https://doi.org/10.1109/TIV.2017.2749181).
- [88] G. Grisetti, C. Stachniss, and W. Burgard. “Nonlinear constraint network optimization for efficient map learning”. In: *IEEE Transactions on Intelligent Transportation Systems 10.3* (2009), pp. 428–439. DOI: [10.1109/TITS.2009.2026444](https://doi.org/10.1109/TITS.2009.2026444).

- [89] M. Kaess, A. Ranganathan, and F. Dellaert. “iSAM: incremental smoothing and mapping”. In: *IEEE Transactions on Robotics* 24.6 (2008), pp. 1365–1378. doi: [10.1109/TRO.2008.2006706](https://doi.org/10.1109/TRO.2008.2006706).
- [90] M. Kaess, H. Johannsson, R. Roberts, V. Ila, J.J. Leonard, and F. Dellaert. “iSAM2: incremental smoothing and mapping with fluid relinearization and incremental variable reordering”. In: *2011 IEEE International Conference on Robotics and Automation (ICRA)*. Shanghai, China, 2011, pp. 3281–3288. doi: [10.1109/ICRA.2011.5979641](https://doi.org/10.1109/ICRA.2011.5979641).
- [91] F. Dellaert. *Factor graphs and GTSAM: a hands-on introduction*. Tech. rep. Georgia Institute of Technology, Sept. 2012. URL: <https://repository.gatech.edu/server/api/core/bitstreams/b3606eb4-ce55-4c16-8495-767bd46f0351/content>.
- [92] F. Dellaert and M. Kaess. “Factor graphs for robot perception”. In: *Foundations and Trends in Robotics* 6.1-2 (2017), pp. 1–139. doi: [10.1561/23000000043](https://doi.org/10.1561/23000000043).
- [93] R. Kümmerle, G. Grisetti, H. Strasdat, K. Konolige, and W. Burgard. “g2o: a general framework for graph optimization”. In: *2011 IEEE International Conference on Robotics and Automation (ICRA)*. Shanghai, China, 2011, pp. 3607–3613. doi: [10.1109/ICRA.2011.5979949](https://doi.org/10.1109/ICRA.2011.5979949).
- [94] R.J. Rossi. *Mathematical statistics: an introduction to likelihood based inference*. First. John Wiley & Sons, 2018. ISBN: 978-1-118-77104-4.
- [95] S. Macenski and I. Jambrecic. “SLAM toolbox: SLAM for the dynamic world”. In: *Journal of Open Source Software* 6.61 (2021), p. 2783. doi: [10.21105/joss.02783](https://doi.org/10.21105/joss.02783).
- [96] Open Robotics. *ROS: robot operating system*. Last accessed on Feb. 25, 2026. 2025. URL: <https://www.ros.org/>.
- [97] S. Agarwal, K. Mierle, and The Ceres Solver Team. *Ceres solver*. Version 2.2. Oct. 2023. URL: <https://github.com/ceres-solver/ceres-solver>.
- [98] M.J. Page, D. Moher, P.M. Bossuyt, I. Boutron, T.C. Hoffmann, C.D. Mulrow, L. Shamseer, J.M. Tetzlaff, E.A. Akl, S.E. Brennan, R. Chou, J. Glanville, J.M. Grimshaw, A. Hróbjartsson, M.M. Lalu, T. Li, E.W. Loder, E. Mayo-Wilson, S. McDonald, L.A. McGuinness, L.A. Stewart, J. Thomas, A.C. Tricco, V.A. Welch, P. Whiting, and J.E. McKenzie. “PRISMA 2020 explanation and elaboration: updated guidance and exemplars for reporting systematic reviews”. In: *BMJ* 372.n160 (2021). doi: [10.1136/bmj.n160](https://doi.org/10.1136/bmj.n160).
- [99] H. Durrant-Whyte and T. Bailey. “Simultaneous localization and mapping: part I”. In: *IEEE Robotics & Automation Magazine* 13.2 (2006), pp. 99–110. doi: [10.1109/MRA.2006.1638022](https://doi.org/10.1109/MRA.2006.1638022).
- [100] T. Bailey and H. Durrant-Whyte. “Simultaneous localization and mapping (SLAM): part II”. In: *IEEE Robotics & Automation Magazine* 13.3 (2006), pp. 108–117. doi: [10.1109/MRA.2006.1678144](https://doi.org/10.1109/MRA.2006.1678144).
- [101] S. Thrun. “Simultaneous localization and mapping”. In: *Robotics and Cognitive Approaches to Spatial Mapping*. Ed. by M.E. Jefferies and W.-K. Yeap. 2008, pp. 13–41. doi: [10.1007/978-3-540-75388-9_3](https://doi.org/10.1007/978-3-540-75388-9_3).
- [102] D. Scaramuzza and F. Fraundorfer. “Visual odometry [tutorial]”. In: *IEEE Robotics & Automation Magazine* 18.4 (2011), pp. 80–92. doi: [10.1109/MRA.2011.943233](https://doi.org/10.1109/MRA.2011.943233).
- [103] F. Fraundorfer and D. Scaramuzza. “Visual odometry: part II: matching, robustness, optimization, and applications”. In: *IEEE Robotics & Automation Magazine* 19.2 (2012), pp. 78–90. doi: [10.1109/MRA.2012.2182810](https://doi.org/10.1109/MRA.2012.2182810).

- [104] K. Yousif, A. Bab-Hadiashar, and R. Hoseinnezhad. “An overview to visual odometry and visual SLAM: applications to mobile robotics”. In: *Intelligent Industrial Systems* 1.4 (2015), pp. 289–311. DOI: [10.1007/s40903-015-0032-7](https://doi.org/10.1007/s40903-015-0032-7).
- [105] G. Dissanayake, S. Huang, Z. Wang, and R. Ranasinghe. “A review of recent developments in simultaneous localization and mapping”. In: *2011 6th International Conference on Industrial and Information Systems*. Kandy, Sri Lanka, 2011, pp. 477–482. DOI: [10.1109/ICIINFS.2011.6038117](https://doi.org/10.1109/ICIINFS.2011.6038117).
- [106] L. Paull, S. Saeedi, M. Seto, and H. Li. “AUV navigation and localization: a review”. In: *IEEE Journal of Oceanic Engineering* 39.1 (2014), pp. 131–149. DOI: [10.1109/JOE.2013.2278891](https://doi.org/10.1109/JOE.2013.2278891).
- [107] S. Huang and G. Dissanayake. “A critique of current developments in simultaneous localization and mapping”. In: *International Journal of Advanced Robotic Systems* 13.5 (2016). DOI: [10.1177/1729881416669482](https://doi.org/10.1177/1729881416669482).
- [108] S. Saeedi, M. Trentini, M. Seto, and H. Li. “Multiple-robot simultaneous localization and mapping: a review”. In: *Journal of Field Robotics* 33.1 (2016), pp. 3–46. DOI: [10.1002/rob.21620](https://doi.org/10.1002/rob.21620).
- [109] S. Lowry, N. Sünderhauf, P. Newman, J.J. Leonard, D. Cox, P. Corke, and M.J. Milford. “Visual place recognition: a survey”. In: *IEEE Transactions on Robotics* 32.1 (2016), pp. 1–19. DOI: [10.1109/TRO.2015.2496823](https://doi.org/10.1109/TRO.2015.2496823).
- [110] S. Kuutti, S. Fallah, K. Katsaros, M. Dianati, F. Mccullough, and A. Mouzakitis. “A survey of the state-of-the-art localization techniques and their potentials for autonomous vehicle applications”. In: *IEEE Internet of Things Journal* 5.2 (2018), pp. 829–846. DOI: [10.1109/JIOT.2018.2812300](https://doi.org/10.1109/JIOT.2018.2812300).
- [111] L. Kunze, N. Hawes, T. Duckett, M. Hanheide, and T. Krajník. “Artificial intelligence for long-term robot autonomy: a survey”. In: *IEEE Robotics and Automation Letters* 3.4 (2018), pp. 4023–4030. DOI: [10.1109/LRA.2018.2860628](https://doi.org/10.1109/LRA.2018.2860628).
- [112] M. Zaffar, S. Ehsan, R. Stolkin, and K.M. Maier. “Sensors, SLAM and long-term autonomy: a review”. In: *2018 NASA/ESA Conference on Adaptive Hardware and Systems (AHS)*. Edinburgh, UK, 2018, pp. 285–290. DOI: [10.1109/AHS.2018.8541483](https://doi.org/10.1109/AHS.2018.8541483).
- [113] J. Fayyad, M.A. Jaradat, D. Gruyer, and H. Najjaran. “Deep learning sensor fusion for autonomous vehicle perception and localization: a review”. In: *Sensors* 20.15 (2020), p. 4220. DOI: [10.3390/s20154220](https://doi.org/10.3390/s20154220).
- [114] J.P. Queralta, J. Taipalmaa, B.C. Pullinen, V.K. Sarker, T.N. Gia, H. Tenhunen, M. Gabbouj, J. Raitoharju, and T. Westerlund. “Collaborative multi-robot search and rescue: planning, coordination, perception, and active vision”. In: *IEEE Access* 8 (2020), pp. 191617–191643. DOI: [10.1109/ACCESS.2020.3030190](https://doi.org/10.1109/ACCESS.2020.3030190).
- [115] C. Badue, R. Guidolini, R.V. Carneiro, P. Azevedo, V.B. Cardoso, A. Forechi, L. Jesus, R. Berriel, T.M. Paixão, F. Mutz, L. de Paula Veronese, T. Oliveira-Santos, and A.F. de Souza. “Self-driving cars: a survey”. In: *Expert Systems with Applications* 165 (2021), p. 113816. DOI: [10.1016/j.eswa.2020.113816](https://doi.org/10.1016/j.eswa.2020.113816).
- [116] K. Ebadi, L. Bernreiter, H. Biggie, G. Catt, Y. Chang, A. Chatterjee, C.E. Denniston, S.-P. Deschênes, K. Harlow, S. Khattak, L. Nogueira, M. Palieri, P. Petráček, M. Petrлік, A. Reinke, V. Krátký, S. Zhao, A.-A. Agha-mohammadi, K. Alexis, C. Heckman, K. Khosoussi, N. Kottege, B. Morrell, M. Hutter, F. Pauling, F. Pomerleau, M. Saska, S. Scherer, R. Siegwart, J.L. Williams, and L. Carlone. “Present

- and future of SLAM in extreme environments: the DARPA SubT challenge”. In: *IEEE Transactions on Robotics* 40 (2024), pp. 936–959. DOI: [10.1109/TRO.2023.3323938](https://doi.org/10.1109/TRO.2023.3323938).
- [117] L. Benos, A.C. Tagarakis, G. Dolias, R. Berruto, D. Kateris, and D. Bochtis. “Machine learning in agriculture: a comprehensive updated review”. In: *Sensors* 21.11 (2021), p. 3758. DOI: [10.3390/s21113758](https://doi.org/10.3390/s21113758).
- [118] E. Chukwu and L. Garg. “A systematic review of blockchain in healthcare: frameworks, prototypes, and implementations”. In: *IEEE Access* 8 (2020), pp. 21196–21214. DOI: [10.1109/ACCESS.2020.2969881](https://doi.org/10.1109/ACCESS.2020.2969881).
- [119] H. Moayed, M. Mosallanezhad, A.S.A. Rashid, W.A.W. Jusoh, and M.A. Muazu. “A systematic review and meta-analysis of artificial neural network application in geotechnical engineering: theory and applications”. In: *Neural Computing and Applications* 32.2 (2020), pp. 495–518. DOI: [10.1007/s00521-019-04109-9](https://doi.org/10.1007/s00521-019-04109-9).
- [120] M.I. Qureshi, N. Khan, S.M. Gillani, and H. Raza. “A systematic review of past decade of mobile learning: what we learned and where to go”. In: *International Journal of Interactive Mobile Technologies* 14.6 (2020), pp. 67–81. DOI: [10.3991/ijim.v14i06.13479](https://doi.org/10.3991/ijim.v14i06.13479).
- [121] Association for Computing Machinery. *ACM digital library*. Last accessed on Feb. 11, 2026. URL: <https://dl.acm.org/>.
- [122] Digital Science & Research Solutions Inc. *Dimensions AI*. Last accessed on Feb. 11, 2026. URL: <https://www.dimensions.ai/>.
- [123] IEEE. *IEEE Xplore*. Last accessed on Feb. 11, 2026. URL: <https://ieeexplore.ieee.org/Xplore/home.jsp>.
- [124] Elsevier B.V. *Engineering Village*. Last accessed on Feb. 11, 2026. URL: <https://www.engineeringvillage.com/app/search/quick/>.
- [125] Elsevier B.V. *Scopus*. Last accessed on Feb. 11, 2026. URL: <https://www.scopus.com/search/form.uri>.
- [126] Clarivate. *Web of Science core collection*. Last accessed on Feb. 11, 2026. URL: <https://www.webofscience.com/wos/woscc/basic-search>.
- [127] V. Freitas. *Parsifal: perform systematic literature reviews*. Last accessed on Feb. 11, 2026. URL: <https://parsif.al/>.
- [128] E. Einhorn and H.-M. Gross. “Generic 2D/3D SLAM with NDT maps for lifelong application”. In: *2013 European Conference on Mobile Robots*. Barcelona, Spain, 2013, pp. 240–247. DOI: [10.1109/ECMR.2013.6698849](https://doi.org/10.1109/ECMR.2013.6698849).
- [129] S.-Y. An, L.-K. Lee, and S.-Y. Oh. “Ceiling vision-based active SLAM framework for dynamic and wide-open environments”. In: *Autonomous Robots* 40.2 (2016), pp. 291–324. DOI: [10.1007/s10514-015-9453-0](https://doi.org/10.1007/s10514-015-9453-0).
- [130] L. Sun, Z. Yan, A. Zaganidis, C. Zhao, and T. Duckett. “Recurrent-OctoMap: learning state-based map refinement for long-term semantic mapping with 3-D-Lidar data”. In: *IEEE Robotics and Automation Letters* 3.4 (2018), pp. 3749–3756. DOI: [10.1109/LRA.2018.2856268](https://doi.org/10.1109/LRA.2018.2856268).
- [131] J. Fu, C. Lin, Y. Taguchi, A. Cohen, Y. Zhang, S. Mylabathula, and J.J. Leonard. “PlaneSDF-based change detection for long-term dense mapping”. In: *IEEE Robotics and Automation Letters* 7.4 (2022), pp. 9667–9674. DOI: [10.1109/LRA.2022.3191794](https://doi.org/10.1109/LRA.2022.3191794).

- [132] K. Nielsen and G. Hendeby. “Multi-hypothesis SLAM for non-static environments with reoccurring landmarks”. In: *IEEE Transactions on Intelligent Vehicles* 8.4 (2023), pp. 3191–3203. doi: [10.1109/TIV.2022.3214978](https://doi.org/10.1109/TIV.2022.3214978).
- [133] R.C. Atkinson and R.M. Shiffrin. “Human memory: a proposed system and its control processes”. In: ed. by K.W. Spence and J.T. Spence. Vol. 2. *Psychology of Learning and Motivation*. Academic Press, 1968, pp. 89–195. doi: [10.1016/S0079-7421\(08\)60422-3](https://doi.org/10.1016/S0079-7421(08)60422-3).
- [134] R. Mur-Artal and J.D. Tardós. “ORB-SLAM2: an open-source SLAM system for monocular, stereo, and RGB-D cameras”. In: *IEEE Transactions on Robotics* 33.5 (2017), pp. 1255–1262. doi: [10.1109/TRO.2017.2705103](https://doi.org/10.1109/TRO.2017.2705103).
- [135] J. Zhu, H. Li, Z. Wang, S. Wang, and T. Zhang. “i-Octree: a fast, lightweight, and dynamic octree for proximity search”. In: *2024 IEEE International Conference on Robotics and Automation (ICRA)*. Yokohama, Japan, 2024, pp. 12290–12296. doi: [10.1109/ICRA57147.2024.10611019](https://doi.org/10.1109/ICRA57147.2024.10611019).
- [136] R. Mur-Artal, J.M.M. Montiel, and J.D. Tardós. “ORB-SLAM: a versatile and accurate monocular SLAM system”. In: *IEEE Transactions on Robotics* 31.5 (2015), pp. 1147–1163. doi: [10.1109/TRO.2015.2463671](https://doi.org/10.1109/TRO.2015.2463671).
- [137] J. Zhang and S. Singh. “LOAM: Lidar odometry and mapping in real-time”. In: *Robotics: Science and Systems (RSS X)*. Vol. 2. 9. Berkeley, USA, 2014, pp. 1–9. doi: [10.15607/RSS.2014.X.007](https://doi.org/10.15607/RSS.2014.X.007).
- [138] M.J. Milford and G.F. Wyeth. “SeqSLAM: visual route-based navigation for sunny summer days and stormy winter nights”. In: *2012 IEEE International Conference on Robotics and Automation (ICRA)*. Saint Paul, MN, USA, 2012, pp. 1643–1649. doi: [10.1109/ICRA.2012.6224623](https://doi.org/10.1109/ICRA.2012.6224623).
- [139] M. Cummins and P. Newman. “FAB-MAP: probabilistic localization and mapping in the space of appearance”. In: *International Journal of Robotics Research* 27.6 (2008), pp. 647–665. doi: [10.1177/0278364908090961](https://doi.org/10.1177/0278364908090961).
- [140] J. Sivic and A. Zisserman. “Video Google: a text retrieval approach to object matching in videos”. In: *Proceedings Ninth IEEE International Conference on Computer Vision*. Vol. 2. Nice, France, 2003, pp. 1470–1477. doi: [10.1109/ICCV.2003.1238663](https://doi.org/10.1109/ICCV.2003.1238663).
- [141] C. Campos, R. Elvira, J.J.G. Rodríguez, J.M.M. Montiel, and J.D. Tardós. “ORB-SLAM3: an accurate open-source library for visual, visual-inertial, and multimap SLAM”. In: *IEEE Transactions on Robotics* 37.6 (2021), pp. 1874–1890. doi: [10.1109/TRO.2021.3075644](https://doi.org/10.1109/TRO.2021.3075644).
- [142] U. Hassler. “Autoregressive moving average processes (ARMA)”. In: *Stochastic Processes and Calculus: An Elementary Introduction with Applications*. Cham: Springer International Publishing, 2016, pp. 45–75. doi: [10.1007/978-3-319-23428-1_3](https://doi.org/10.1007/978-3-319-23428-1_3).
- [143] M. Magnusson, A. Lilienthal, and T. Duckett. “Scan registration for autonomous mining vehicles using 3D-NDT”. In: *Journal of Field Robotics* 24.10 (2007), pp. 803–827. doi: [10.1002/rob.20204](https://doi.org/10.1002/rob.20204).
- [144] K.M. Wurm, C. Stachniss, and G. Grisetti. “Bridging the gap between feature- and grid-based SLAM”. In: *Robotics and Autonomous Systems* 58.2 (2010), pp. 140–148. doi: [10.1016/j.robot.2009.09.009](https://doi.org/10.1016/j.robot.2009.09.009).
- [145] H. Sobreira, A.P. Moreira, P.G. Costa, and J. Lima. “Robust mobile robot localization based on security laser scanner”. In: *2015 IEEE International Conference on Autonomous Robot Systems and Competitions (ICARSC)*. Vila Real, Portugal, 2015, pp. 162–167. doi: [10.1109/ICARSC.2015.28](https://doi.org/10.1109/ICARSC.2015.28).

- [146] J.A. Castellanos, J. Neira, O. Strauss, and J.D. Tardos. “Detecting high level features for mobile robot localization”. In: *1996 IEEE/SICE/RSJ International Conference on Multisensor Fusion and Integration for Intelligent Systems*. Washington, DC, USA, 1996, pp. 611–618. doi: [10.1109/MFI.1996.572237](https://doi.org/10.1109/MFI.1996.572237).
- [147] S.I. Roumeliotis and G.A. Bekey. “SEGMENTS: a layered, dual-Kalman filter algorithm for indoor feature extraction”. In: *2000 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. Vol. 1. Takamatsu, Japan, 2000, pp. 454–461. doi: [10.1109/IROS.2000.894646](https://doi.org/10.1109/IROS.2000.894646).
- [148] H.A. Mohamed, A.M. Moussa, M.M. Elhabiby, N. El-Sheimy, and A.B. Sesay. “Improved real-time scan matching using corner features”. In: *The International Archives of the Photogrammetry, Remote Sensing and Spatial Information Sciences XLI-B5* (2016), pp. 533–539. doi: [10.5194/isprs-archives-XLI-B5-533-2016](https://doi.org/10.5194/isprs-archives-XLI-B5-533-2016).
- [149] Y. Liu, L. Sui, P. Li, L. Zhang, Q. Wu, J. Du, Y. Liu, and H. Yu. “A radar linear feature fitting algorithm combining adaptive clustering and corner detection operator”. In: *Journal of Sensors* 2023.1 (2023), p. 6991467. doi: [10.1155/2023/6991467](https://doi.org/10.1155/2023/6991467).
- [150] P. Nunez, R. Vazquez-Martin, J.C. del Toro, A. Bandera, and F. Sandoval. “Feature extraction from laser scan data based on curvature estimation for mobile robotics”. In: *2006 IEEE International Conference on Robotics and Automation (ICRA)*. Orlando, FL, USA, 2006, pp. 1167–1172. doi: [10.1109/ROBOT.2006.1641867](https://doi.org/10.1109/ROBOT.2006.1641867).
- [151] R. Vázquez-Martín, P. Núñez, A. Bandera, and F. Sandoval. “Curvature-based environment description for robot navigation using laser range sensors”. In: *Sensors* 9.8 (2009), pp. 5894–5918. doi: [10.3390/s90805894](https://doi.org/10.3390/s90805894).
- [152] X. Feng, S. Guo, X. Li, and Y. He. “Robust mobile robot localization by tracking natural landmarks”. In: *Artificial Intelligence and Computational Intelligence (AICI 2009)*. Ed. by H. Deng, L. Wang, F.L. Wang, and J. Lei. Shanghai, China, 2009, pp. 278–287. doi: [10.1007/978-3-642-05253-8_31](https://doi.org/10.1007/978-3-642-05253-8_31).
- [153] S.T. Pfister, S.I. Roumeliotis, and J.W. Burdick. “Weighted line fitting algorithms for mobile robot map building and efficient data representation”. In: *2003 IEEE International Conference on Robotics and Automation (ICRA)*. Vol. 1. Taipei, Taiwan, 2003, pp. 1304–1311. doi: [10.1109/ROBOT.2003.1241772](https://doi.org/10.1109/ROBOT.2003.1241772).
- [154] H. Gao, X. Zhang, Y. Fang, and J. Yuan. “A line segment extraction algorithm using laser data based on seeded region growing”. In: *International Journal of Advanced Robotic Systems* 15.1 (2018), p. 1729881418755245. doi: [10.1177/1729881418755245](https://doi.org/10.1177/1729881418755245).
- [155] Z. Falomir, L. Museros, V. Castelló, and L. Gonzalez-Abril. “Qualitative distances and qualitative image descriptions for representing indoor scenes in robotics”. In: *Pattern Recognition Letters* 34.7 (2013), pp. 731–743. doi: [10.1016/j.patrec.2012.08.012](https://doi.org/10.1016/j.patrec.2012.08.012).
- [156] H.G. Jung, Y.H. Cho, P.J. Yoon, and J. Kim. “Scanning laser radar-based target position designation for parking aid system”. In: *IEEE Transactions on Intelligent Transportation Systems* 9.3 (2008), pp. 406–424. doi: [10.1109/TITS.2008.922980](https://doi.org/10.1109/TITS.2008.922980).
- [157] W. Lin, J. Hu, H. Xu, C. Ye, X. Ye, and Z. Li. “Graph-based SLAM in indoor environment using corner feature from laser sensor”. In: *2017 32nd Youth Academic Annual Conference of Chinese Association of Automation (YAC)*. Hefei, China, 2017, pp. 1211–1216. doi: [10.1109/YAC.2017.7967597](https://doi.org/10.1109/YAC.2017.7967597).

- [158] Y. Liu, L. Zhang, K. Qian, L. Sui, Y. Lu, F. Qian, T. Yan, H. Yu, and F. Gao. “An adaptive threshold line segment feature extraction algorithm for laser radar scanning environments”. In: *Electronics* 11.11 (2022), p. 1759. doi: [10.3390/electronics11111759](https://doi.org/10.3390/electronics11111759).
- [159] N. Certad, R. Acuna, A. Terrones, D. Ralev, J. Cappelletto, and J.C. Grieco. “Study and improvements in landmarks extraction in 2D range images based on an adaptive curvature estimation”. In: *2012 VI Andean Region International Conference*. Cuenca, Ecuador, 2012, pp. 95–98. doi: [10.1109/Andescon.2012.31](https://doi.org/10.1109/Andescon.2012.31).
- [160] S. Qu, G. Chen, C. Ye, F. Lu, F. Wang, Z. Xu, and Y. Gel. “An efficient L-shape fitting method for vehicle pose detection with 2D LiDAR”. In: *2018 IEEE International Conference on Robotics and Biomimetics (ROBIO)*. Kuala Lumpur, Malaysia, 2018, pp. 1159–1164. doi: [10.1109/ROBIO.2018.8665265](https://doi.org/10.1109/ROBIO.2018.8665265).
- [161] E. Fernández-Moral, J. González-Jiménez, and V. Arévalo. “Extrinsic calibration of 2D laser rangefinders from perpendicular plane observations”. In: *International Journal of Robotics Research* 34.11 (2015), pp. 1401–1417. doi: [10.1177/0278364915580683](https://doi.org/10.1177/0278364915580683).
- [162] Y. Zhao, F. Liu, and R. Wang. “Location technology of indoor robot based on laser sensor”. In: *2016 7th IEEE International Conference on Software Engineering and Service Science (ICSESS)*. Beijing, China, 2016, pp. 683–686. doi: [10.1109/ICSESS.2016.7883160](https://doi.org/10.1109/ICSESS.2016.7883160).
- [163] J. Xavier, M. Pacheco, D. Castro, A. Ruano, and U. Nunes. “Fast line, arc/circle and leg detection from laser scan data in a player driver”. In: *2005 IEEE International Conference on Robotics and Automation (ICRA)*. Barcelona, Spain, 2005, pp. 3930–3935. doi: [10.1109/ROBOT.2005.1570721](https://doi.org/10.1109/ROBOT.2005.1570721).
- [164] M.A. Fischler and R.C. Bolles. “Random sample consensus: a paradigm for model fitting with applications to image analysis and automated cartography”. In: *Communications of the ACM* 24.6 (1981), pp. 381–395. doi: [10.1145/358669.358692](https://doi.org/10.1145/358669.358692).
- [165] L.R. Muñoz, M.G. Villanueva, and C.A.L. Álvarez. “An extended line tracking algorithm”. In: *2013 IEEE International Autumn Meeting on Power Electronics and Computing (ROPEC)*. Morelia, Mexico, 2013, pp. 1–5. doi: [10.1109/ROPEC.2013.6702752](https://doi.org/10.1109/ROPEC.2013.6702752).
- [166] M. Ester, H.-P. Kriegel, J. Sander, and X. Xu. “A density-based algorithm for discovering clusters in large spatial databases with noise”. In: *Proceedings of the Second International Conference on Knowledge Discovery and Data Mining (KDD’96)*. 1996, pp. 226–231. doi: [Portland, Oregon, USA](https://doi.org/10.1145/358669.358692).
- [167] G.A. Borges and M.-J. Aldon. “Line extraction in 2D range images for mobile robotics”. In: *Journal of Intelligent and Robotic Systems* 40.3 (2004), pp. 267–297. doi: [10.1023/B:JINT.0000038945.55712.65](https://doi.org/10.1023/B:JINT.0000038945.55712.65).
- [168] J.F. Kenney and E.S. Keeping. “Linear regression and correlation”. In: 3rd ed. *Mathematics of Statistics*. Princeton: Van Nostrand, 1962. Chap. 15, pp. 252–285.
- [169] 5dpo. *5dpo robotics team. Faculty of Engineering, University of Porto (FEUP)*. Last accessed on Mar. 17, 2026. URL: <https://5dpo.github.io/>.
- [170] H. Bay, T. Tuytelaars, and L. Van Gool. “SURF: speeded up robust features”. In: *Computer Vision – European Conference on Computer Vision (ECCV) 2006*. Ed. by A. Leonardis, H. Bischof, and A. Pinz. Graz, Austria, 2006, pp. 404–417. doi: [10.1007/11744023_32](https://doi.org/10.1007/11744023_32).
- [171] Y. Chen and G. Medioni. “Object modeling by registration of multiple range images”. In: *1991 IEEE International Conference on Robotics and Automation (ICRA)*. Vol. 3. Sacramento, CA, USA, 1991, pp. 2724–2729. doi: [10.1109/ROBOT.1991.132043](https://doi.org/10.1109/ROBOT.1991.132043).

- [172] s. Rusinkiewicz and M. Levoy. “Efficient variants of the ICP algorithm”. In: *Proceedings Third International Conference on 3-D Digital Imaging and Modeling*. Quebec City, QC, Canada, 2001, pp. 145–152. DOI: [10.1109/IM.2001.924423](https://doi.org/10.1109/IM.2001.924423).
- [173] J.L. Bentley. “Multidimensional binary search trees used for associative searching”. In: *Communications of the ACM* 18.9 (1975), pp. 509–517. DOI: [10.1145/361002.361007](https://doi.org/10.1145/361002.361007).
- [174] A. Segal, D. Haehnel, and S. Thrun. “Generalized-ICP”. In: *Proceedings of the Robotics: Science and Systems (RSS V)*. Seattle, USA, 2009. DOI: [10.15607/RSS.2009.V.021](https://doi.org/10.15607/RSS.2009.V.021).
- [175] L. Montesano, J. Minguez, and L. Montano. “Probabilistic scan matching for motion estimation in unstructured environments”. In: *2005 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. Edmonton, AB, Canada, 2005, pp. 3499–3504. DOI: [10.1109/IROS.2005.1545182](https://doi.org/10.1109/IROS.2005.1545182).
- [176] A. Censi. “Scan matching in a probabilistic framework”. In: *2006 IEEE International Conference on Robotics and Automation (ICRA)*. Orlando, FL, USA, 2006, pp. 2291–2296. DOI: [10.1109/ROBOT.2006.1642044](https://doi.org/10.1109/ROBOT.2006.1642044).
- [177] W. Hess, D. Kohler, H. Rapp, and D. Andor. “Real-time loop closure in 2D LIDAR SLAM”. In: *2016 IEEE International Conference on Robotics and Automation (ICRA)*. Stockholm, Sweden, 2016, pp. 1271–1278. DOI: [10.1109/ICRA.2016.7487258](https://doi.org/10.1109/ICRA.2016.7487258).
- [178] H. Sobreira, C.M. Costa, I. Sousa, L. Rocha, J. Lima, P.C.M.A. Farias, P. Costa, and A.P. Moreira. “Map-matching algorithms for robot self-localization: a comparison between perfect match, iterative closest point and normal distributions transform”. In: *Journal of Intelligent & Robotic Systems* 93.3 (2019), pp. 533–546. DOI: [10.1007/s10846-017-0765-5](https://doi.org/10.1007/s10846-017-0765-5).
- [179] M. Riedmiller and H. Braun. “A direct adaptive method for faster backpropagation learning: the RPROP algorithm”. In: *IEEE International Conference on Neural Networks (ICNN)*. Vol. 1. San Francisco, CA, USA, 1993, pp. 586–591. DOI: [10.1109/ICNN.1993.298623](https://doi.org/10.1109/ICNN.1993.298623).
- [180] T.D. Barfoot. *State estimation for robotics*. First. Cambridge: Cambridge University Press, 2017. ISBN: 978-1-107-15939-6.
- [181] J. Nocedal and S.J. Wright. *Numerical optimization*. 2nd ed. Springer Series in Operations Research and Financial Engineering. New York: Springer, 2006. ISBN: 978-0-387-30303-1.
- [182] T.B. Schön and F. Lindsten. *Manipulating the multivariate Gaussian density*. Tech. rep. Linköping, Sweden: Division of Automatic Control. Linköping University, Jan. 2011. URL: <http://ais.informatik.uni-freiburg.de/teaching/ss23/robotics/etc/schonl2011.pdf>.
- [183] P.J. Huber. *Robust statistics*. First. New York: Wiley, 1981. ISBN: 978-0-471-41805-4. DOI: [10.1002/0471725250](https://doi.org/10.1002/0471725250).
- [184] P. Babin, P. Giguère, and F. Pomerleau. “Analysis of robust functions for registration algorithms”. In: *2019 International Conference on Robotics and Automation (ICRA)*. Motreal, QC, Canada, 2019, pp. 1451–1457. DOI: [10.1109/ICRA.2019.8793791](https://doi.org/10.1109/ICRA.2019.8793791).
- [185] M.W. Jones, J.A. Baerentzen, and M. Sramek. “3D distance fields: a survey of techniques and applications”. In: *IEEE Transactions on Visualization and Computer Graphics* 12.4 (2006), pp. 581–599. DOI: [10.1109/TVCG.2006.56](https://doi.org/10.1109/TVCG.2006.56).

- [186] E.W. Dijkstra. “A note on two problems in connexion with graphs”. In: *Numerische Mathematik* 1.1 (1959), pp. 269–271. doi: [10.1007/BF01386390](https://doi.org/10.1007/BF01386390).
- [187] B. Curless and M. Levoy. “A volumetric method for building complex models from range images”. In: *Proceedings of the 23rd Annual Conference on Computer Graphics and Interactive Techniques (SIGGRAPH '96)*. New Orleans, LA, USA, 1996, pp. 303–312. doi: [10.1145/237170.237269](https://doi.org/10.1145/237170.237269).
- [188] R.B. Sousa, P.G. Costa, and A.P. Moreira. “A pose control algorithm for omnidirectional robots”. In: *2021 IEEE International Conference on Autonomous Robot Systems and Competitions (ICARSC)*. Santa Maria da Feira, Portugal, 2021, pp. 91–96. doi: [10.1109/ICARSC52212.2021.9429803](https://doi.org/10.1109/ICARSC52212.2021.9429803).
- [189] R.B. Sousa, C.D. Rocha, J.G. Martins, J.P. Costa, J.T. Padrão, J.M. Sarmiento, J.P. Carvalho, M.S. Lopes, P.G. Costa, and A.P. Moreira. “A robotic framework for the robot@factory 4.0 competition”. In: *2024 IEEE International Conference on Autonomous Robot Systems and Competitions (ICARSC)*. Paredes de Coura, Portugal, 2024, pp. 66–73. doi: [10.1109/ICARSC61747.2024.10535935](https://doi.org/10.1109/ICARSC61747.2024.10535935).
- [190] M.S. Lopes, J.D. Ribeiro, A.P. Moreira, C.D. Rocha, J.G. Martins, J.M. Sarmiento, J.P. Carvalho, P.G. Costa, and R.B. Sousa. “From competition to classroom: a hands-on approach to robotics learning”. In: *2025 IEEE International Conference on Autonomous Robot Systems and Competitions (ICARSC)*. Funchal, Portugal, 2025, pp. 170–176. doi: [10.1109/ICARSC65809.2025.10970153](https://doi.org/10.1109/ICARSC65809.2025.10970153).
- [191] T.B. Levin, J.M. Oliveira, R.B. Sousa, M.F. Silva, B.S. Parreira, H.M. Sobreira, and H.S. Mendonça. “Image and command transmission over the 5G network for teleoperation of mobile robots”. In: *2024 7th Iberian Robotics Conference (ROBOT)*. Madrid, Spain, 2024, pp. 1–8. doi: [10.1109/ROBOT61475.2024.10797434](https://doi.org/10.1109/ROBOT61475.2024.10797434).
- [192] D. Caldana, A. Cordeiro, J.P. Souza, R.B. Sousa, P.M. Rebelo, A.J. Silva, and M.F. Silva. “Pallet and pocket detection based on deep learning techniques”. In: *2024 7th Iberian Robotics Conference (ROBOT)*. Madrid, Spain, 2024, pp. 1–8. doi: [10.1109/ROBOT61475.2024.10797410](https://doi.org/10.1109/ROBOT61475.2024.10797410).
- [193] F.D. Pacheco, P.M. Rebelo, R.B. Sousa, M.F. Silva, and H.S. Mendonça. “Integrated RFID system for intralogistics operations with industrial mobile robots”. In: *2025 IEEE International Conference on Autonomous Robot Systems and Competitions (ICARSC)*. Funchal, Portugal, 2025, pp. 124–131. doi: [10.1109/ICARSC65809.2025.10970172](https://doi.org/10.1109/ICARSC65809.2025.10970172).
- [194] R.B. Sousa, C. Rocha, H.S. Mendonça, A.P. Moreira, and M.F. Silva. “Gerber file parsing for conversion to bitmap image—the VINCI7D case study”. In: *IEEE Access* 10 (2022), pp. 69659–69679. doi: [10.1109/ACCESS.2022.3187042](https://doi.org/10.1109/ACCESS.2022.3187042).

Appendix A

A Systematic Literature Review on Long-Term Localization and Mapping for Mobile Robots

- [65] R.B. Sousa, H.M. Sobreira, and A.P. Moreira. “A systematic literature review on long-term localization and mapping for mobile robots”. In: *Journal of Field Robotics* 40.5 (2023), pp. 1245–1322. DOI: [10.1002/rob.22170](https://doi.org/10.1002/rob.22170).

Appendix B

Integrating Multimodal Perception into Ground Mobile Robots

- [67] R.B. Sousa, H.M. Sobreira, J.G. Martins, P.G. Costa, M.F. Silva, and A.P. Moreira. “Integrating multimodal perception into ground mobile robots”. In: *2025 IEEE International Conference on Autonomous Robot Systems and Competitions (ICARSC)*. Funchal, Portugal, 2025, pp. 104–111. doi: [10.1109/ICARSC65809.2025.10970176](https://doi.org/10.1109/ICARSC65809.2025.10970176).

Appendix C

Line Fitting-Based Corner-Like Detector for 2D Laser Scanners Data

- [66] R.B. Sousa, H.M. Sobreira, M.F. Silva, and A.P. Moreira. “Line fitting-based corner-like detector for 2D laser scanners data”. In: *2024 10th International Conference on Automation, Robotics and Applications (ICARA)*. Athens, Greece, 2024, pp. 210–216. doi: [10.1109/ICARA60736.2024.10553006](https://doi.org/10.1109/ICARA60736.2024.10553006).

